

# SIMILARITY AND DISTANCE METRIC LEARNING

MDI 341, MS BIG DATA, TÉLÉCOM PARISTECH

---

**Aurélien Bellet**

Inria Lille

March 1, 2017


1. Introduction
2. Linear metric learning
3. Nonlinear extensions
4. Large-scale metric learning
5. Metric learning for structured data

## INTRODUCTION

---


- Similarity / distance judgments are essential components of many human cognitive processes
  - Compare perceptual or conceptual representations
  - Perform recognition, categorization...
- Underlie most machine learning and data mining techniques

## Nearest neighbor classification



# MOTIVATION

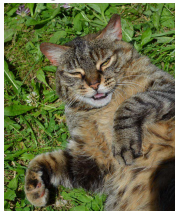
## Clustering



# MOTIVATION

## Information retrieval

**Query document**




**Most similar documents**



# MOTIVATION

## Data visualization

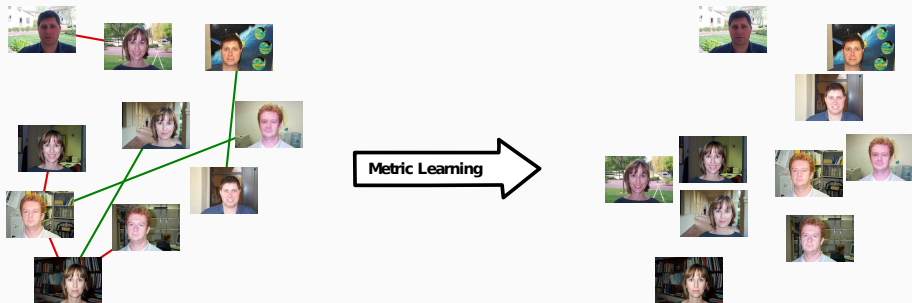


(image taken from [van der Maaten and Hinton, 2008])



- Choice of similarity is crucial to the performance
- Humans weight features differently depending on context
  - Facial recognition vs. determining facial expression
- Fundamental question: **how to appropriately measure similarity or distance** for a given task?
- Metric learning → infer this automatically from data
- Note: we will refer to *distance* or *similarity* indistinctly as *metric*

# METRIC LEARNING IN A NUTSHELL



## Basic recipe

1. Pick a **parametric distance or similarity function**
  - Say, a distance  $D_M(x, x')$  function parameterized by a matrix  $M$
2. Collect **similarity judgments** on data pairs/triplets
  - $\mathcal{S} = \{(x_i, x_j) : x_i \text{ and } x_j \text{ are similar}\}$
  - $\mathcal{D} = \{(x_i, x_j) : x_i \text{ and } x_j \text{ are dissimilar}\}$
  - $\mathcal{R} = \{(x_i, x_j, x_k) : x_i \text{ is more similar to } x_j \text{ than to } x_k\}$
3. **Estimate parameters** s.t. metric best agrees with judgments
  - Solve an optimization problem of the form

$$M^* = \arg \min_M \left[ \underbrace{\ell(M, \mathcal{S}, \mathcal{D}, \mathcal{R})}_{\text{loss function}} + \underbrace{\lambda \text{reg}(M)}_{\text{regularization}} \right]$$

# LINEAR METRIC LEARNING

---

## Definition (Distance function)

A distance over a set  $\mathcal{X}$  is a pairwise function  $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  which satisfies the following properties  $\forall x, x', x'' \in \mathcal{X}$ :

- (1)  $d(x, x') \geq 0$  (nonnegativity)
- (2)  $d(x, x') = 0$  if and only if  $x = x'$  (identity of indiscernibles)
- (3)  $d(x, x') = d(x', x)$  (symmetry)
- (4)  $d(x, x'') \leq d(x, x') + d(x', x'')$  (triangle inequality)

- Note: a **pseudo-distance** satisfies the above except (2)

## Minkowski distances

- A family of distances induced by  $L_p$  norms ( $p \geq 1$ )

$$d_p(\mathbf{x}, \mathbf{x}') = \|\mathbf{x} - \mathbf{x}'\|_p = \left( \sum_{i=1}^d |x_i - x'_i|^p \right)^{1/p}$$

- When  $p = 2$ : “ordinary” Euclidean distance

$$d_{euc}(\mathbf{x}, \mathbf{x}') = \left( \sum_{i=1}^d |x_i - x'_i|^2 \right)^{1/2} = \sqrt{(\mathbf{x} - \mathbf{x}')^T (\mathbf{x} - \mathbf{x}')}$$

- When  $p = 1$ : Manhattan distance  $d_{man}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^d |x_i - x'_i|$
- When  $p \rightarrow \infty$ : Chebyshev distance  $d_{che}(\mathbf{x}, \mathbf{x}') = \max_i |x_i - x'_i|$


# MAHALANOBIS DISTANCE

- Mahalanobis (pseudo) distance:

$$D_M(x, x') = \sqrt{(x - x')^T M (x - x')}$$

where  $M \in \mathbb{R}^{d \times d}$  is symmetric positive semi-definite (PSD)

- Denote by  $\mathbb{S}_+^d$  the cone of symmetric PSD  $d \times d$  matrices



- A symmetric matrix  $\mathbf{M}$  is in  $\mathbb{S}_+^d$  (also denoted  $\mathbf{M} \succeq 0$ ) iff:
  - Its eigenvalues are all nonnegative
  - $\mathbf{x}^T \mathbf{M} \mathbf{x} \geq 0, \forall \mathbf{x} \in \mathbb{R}^d$
  - $\mathbf{M} = \mathbf{L}^T \mathbf{L}$  for some  $\mathbf{L} \in \mathbb{R}^{k \times d}, k \leq d$
- Equivalent to Euclidean distance after linear transformation:

$$D_{\mathbf{M}}(\mathbf{x}, \mathbf{x}') = \sqrt{(\mathbf{x} - \mathbf{x}')^T \mathbf{L}^T \mathbf{L} (\mathbf{x} - \mathbf{x}')} = \sqrt{(\mathbf{L}\mathbf{x} - \mathbf{L}\mathbf{x}')^T (\mathbf{L}\mathbf{x} - \mathbf{L}\mathbf{x}')}$$

- If  $\text{rank}(\mathbf{M}) = k \leq d$ , then  $\mathbf{L} \in \mathbb{R}^{k \times d}$  does dimensionality reduction
- For convenience, we often work with the squared distance



A first approach with pairwise constraints [Xing et al., 2002]



- Targeted task: clustering with side information

## Formulation

$$\begin{aligned} \max_{\mathbf{M} \in \mathbb{S}_+^d} \quad & \sum_{(x_i, x_j) \in \mathcal{D}} D_{\mathbf{M}}(x_i, x_j) \\ \text{s.t.} \quad & \sum_{(x_i, x_j) \in \mathcal{S}} D_{\mathbf{M}}^2(x_i, x_j) \leq 1 \end{aligned}$$

- Problem is convex in  $\mathbf{M}$  and always feasible (take  $\mathbf{M} = \mathbf{0}$ )
- Solved with projected gradient descent
  - Project onto distance constraint:  $O(d^2)$  time
  - Project onto  $\mathbb{S}_+^d$ :  $O(d^3)$  time
- Only look at sums of distances

A first approach with pairwise constraints [Xing et al., 2002]



A first approach with triplet constraints [Schultz and Joachims, 2003]


## Formulation

$$\begin{aligned} \min_{\mathbf{M} \in \mathbb{S}_+^d, \boldsymbol{\xi} \geq 0} \quad & \|\mathbf{M}\|_{\mathcal{F}}^2 + \lambda \sum_{i,j,k} \xi_{ijk} \\ \text{s.t.} \quad & D_{\mathbf{M}}^2(\mathbf{x}_i, \mathbf{x}_k) - D_{\mathbf{M}}^2(\mathbf{x}_i, \mathbf{x}_j) \geq 1 - \xi_{ijk} \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \end{aligned}$$

- Regularization by Frobenius norm  $\|\mathbf{M}\|_{\mathcal{F}}^2 = \sum_{i,j=1}^d M_{ij}^2$
- Formulation is **convex**
- One **large margin soft constraint** per triplet
- Can be solved with similar techniques as SVM

## Large Margin Nearest Neighbor [Weinberger et al., 2005]

- Targeted task:  $k$ -NN classification
- Constraints derived from labeled data
  - $\mathcal{S} = \{(\mathbf{x}_i, \mathbf{x}_j) : y_i = y_j, \mathbf{x}_j \text{ belongs to } k\text{-neighborhood of } \mathbf{x}_i\}$
  - $\mathcal{R} = \{(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) : (\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}, y_i \neq y_k\}$



## Large Margin Nearest Neighbor [Weinberger et al., 2005]

### Formulation



$$\begin{aligned} \min_{M \in \mathbb{S}_+^d, \xi \geq 0} \quad & (1 - \mu) \sum_{(x_i, x_j) \in \mathcal{S}} D_M^2(x_i, x_j) \quad + \quad \mu \sum_{i,j,k} \xi_{ijk} \\ \text{s.t.} \quad & D_M^2(x_i, x_k) - D_M^2(x_i, x_j) \geq 1 - \xi_{ijk} \quad \forall (x_i, x_j, x_k) \in \mathcal{R} \end{aligned}$$

$\mu \in [0, 1]$  trade-off parameter

- **Convex** formulation, unlike NCA [Goldberger et al., 2004]
- Number of constraints in the order of  $kn^2$ 
  - Solver based on projected gradient descent with working set
  - Simple alternative: only consider closest “impostors”
- Chicken and egg situation: which metric to build constraints?

# MAHALANOBIS DISTANCE LEARNING

## Large Margin Nearest Neighbor [Weinberger et al., 2005]



## Pointers to metric learning algorithms for other tasks

- Learning to rank [McFee and Lanckriet, 2010]
- Multi-task learning [Parameswaran and Weinberger, 2010]
- Transfer learning [Zhang and Yeung, 2010]
- Semi-supervised learning [Hoi et al., 2008]


## Interesting regularizers

- Add regularization term to **prevent overfitting**
- We have already seen the Frobenius norm  $\|\mathbf{M}\|_{\mathcal{F}}^2 = \sum_{i,j=1}^d M_{ij}^2$ 
  - Convex, smooth  $\rightarrow$  easy to optimize
- Mixed  $L_{2,1}$  norm:  $\|\mathbf{M}\|_{2,1} = \sum_{i=1}^d \|\mathbf{M}_i\|_2$ 
  - Tends to zero-out entire columns  $\rightarrow$  feature selection
  - Convex but nonsmooth
  - Efficient proximal gradient algorithms
- Trace (or nuclear) norm:  $\|\mathbf{M}\|_* = \sum_{i=1}^d \sigma_i(\mathbf{M})$ 
  - Favors low-rank matrices  $\rightarrow$  dimensionality reduction
  - Convex but nonsmooth
  - Efficient Frank-Wolfe algorithms



# MAHALANOBIS DISTANCE LEARNING


## $L_{2,1}$ norm illustration



(image taken from [Ying et al, 2009])

# LINEAR SIMILARITY LEARNING

- Mahalanobis distance satisfies some distance properties
  - Nonnegativity, symmetry, triangle inequality
  - Natural regularization, required by some applications
- In practice, these properties may not be satisfied
  - By human similarity judgments



- By some good visual recognition systems
- Alternative: learn **bilinear similarity** function  $S_M(x, x') = x^T M x'$ 
  - Example: OASIS algorithm (presented later)
  - No PSD constraint on  $M \rightarrow$  computationally easier

## NONLINEAR EXTENSIONS

---

## Definition (Kernel function)

A symmetric function  $K$  is a kernel if there exists a mapping function  $\phi : \mathcal{X} \rightarrow \mathbb{H}$  from the instance space  $\mathcal{X}$  to a Hilbert space  $\mathbb{H}$  such that  $K$  can be written as an inner product in  $\mathbb{H}$ :

$$K(x, x') = \langle \phi(x), \phi(x') \rangle.$$

Equivalently,  $K$  is a kernel if it is positive semi-definite (PSD), i.e.,

$$\sum_{i=1}^n \sum_{j=1}^n c_i c_j K(x_i, x_j) \geq 0$$

for all finite sequences of  $x_1, \dots, x_n \in \mathcal{X}$  and  $c_1, \dots, c_n \in \mathbb{R}$ .

## Kernel trick for metric learning

- Notations

- Kernel  $K(\mathbf{x}, \mathbf{x}') = \langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle$ , training data  $\{\mathbf{x}_i\}_{i=1}^n$
- $\phi_i \stackrel{\text{def}}{=} \phi(\mathbf{x}_i) \in \mathbb{R}^D$ ,  $\Phi \stackrel{\text{def}}{=} [\phi_1, \dots, \phi_n] \in \mathbb{R}^{n \times D}$

- Mahalanobis distance in kernel space

$$D_M^2(\phi_i, \phi_j) = (\phi_i - \phi_j)^T M (\phi_i - \phi_j) = (\phi_i - \phi_j)^T L^T L (\phi_i - \phi_j)$$

- Setting  $L^T = \Phi U^T$ , where  $U \in \mathbb{R}^{D \times n}$ , we get

$$D_M^2(\phi(\mathbf{x}), \phi(\mathbf{x}')) = (\mathbf{k} - \mathbf{k}')^T M (\mathbf{k} - \mathbf{k}')$$


- $M = U^T U \in \mathbb{R}^{n \times n}$ ,  $\mathbf{k} = \Phi^T \phi(\mathbf{x}) = [K(\mathbf{x}_1, \mathbf{x}), \dots, K(\mathbf{x}_n, \mathbf{x})]^T$

- Theoretically justified (representer theorem)

## Kernel trick for metric learning

- Similar trick as kernel SVM
  - Use a nonlinear kernel (e.g., Gaussian RBF)
  - Inexpensive computations through the kernel
  - Nonlinear metric learning while retaining convexity
- Need to learn  $O(n^2)$  parameters
- Linear metric learning algorithm must be **kernelized**
  - Interface to data limited to inner products only
  - Several algorithms shown to be kernelizable
- General trick (simple and works well in practice):
  1. Kernel PCA: nonlinear mapping to low-dimensional space
  2. Apply linear metric learning algorithm to transformed data

# LEARNING A NONLINEAR METRIC




- More flexible approach: learn **nonlinear mapping**  $\phi$  to optimize

$$D_{\phi}(\mathbf{x}, \mathbf{x}') = \|\phi(\mathbf{x}) - \phi(\mathbf{x}')\|_2$$

- Possible parameterizations for  $\phi$ :
  - Regression trees
  - Deep neural nets
  - ...
- Typically nonconvex formulations

# LEARNING MULTIPLE LOCAL METRICS

- Simple linear metrics perform well locally
- Idea: different metrics for different parts of the space





## Multiple Metric LMNN [Weinberger and Saul, 2009]


- Group data into  $C$  clusters
- Learn a metric for each cluster in a coupled fashion

### Formulation

$$\begin{aligned} \min_{\substack{M_1, \dots, M_C \\ \xi \geq 0}} \quad & (1 - \mu) \sum_{(x_i, x_j) \in \mathcal{S}} D_{M_{C(x_j)}}^2(x_i, x_j) + \mu \sum_{i,j,k} \xi_{ijk} \\ \text{s.t.} \quad & D_{M_{C(x_k)}}^2(x_i, x_k) - D_{M_{C(x_j)}}^2(x_i, x_j) \geq 1 - \xi_{ijk} \quad \forall (x_i, x_j, x_k) \in \mathcal{R} \end{aligned}$$

- Remains convex
- Computationally more expensive than standard LMNN
- Subject to overfitting (many parameters)

## Multiple Metric LMNN [Weinberger and Saul, 2009]



# LARGE-SCALE METRIC LEARNING

---

- How to deal with large datasets?
  - Number of similarity judgments can grow as  $O(n^2)$  or  $O(n^3)$
- How to deal with high-dimensional data?
  - Cannot store  $d \times d$  matrix
  - Cannot afford computational complexity in  $O(d^2)$  or  $O(d^3)$

### Online learning

- Online algorithm
  - Receive *one* similarity judgment
  - Suffer loss based on current metric
  - Update metric and iterate
- Goal: minimize **regret**

$$\sum_{t=1}^T \ell_t(\mathbf{M}_t) - \sum_{t=1}^T \ell_t(\mathbf{M}^*) \leq f(T),$$

- $\ell_t$ : loss suffered at time  $t$
- $\mathbf{M}_t$ : metric learned at time  $t$
- $\mathbf{M}^*$ : best metric in hindsight

## OASIS [Chechik et al., 2010]

### Formulation

- Set  $M^0 = I$
- At step  $t$ , receive  $(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R}$  and update by solving



$$\begin{aligned} M^t = \arg \min_{M, \xi} \quad & \frac{1}{2} \|M - M^{t-1}\|_{\mathcal{F}}^2 + C\xi \\ \text{s.t.} \quad & 1 - S_M(\mathbf{x}_i, \mathbf{x}_j) + S_M(\mathbf{x}_i, \mathbf{x}_k) \leq \xi \\ & \xi \geq 0 \end{aligned}$$

- $S_M(\mathbf{x}, \mathbf{x}') = \mathbf{x}^T M \mathbf{x}'$ ,  $C$  trade-off parameter

- Simple closed-form solution at each iteration
- Trained with 160M triplets in 3 days on 1 CPU

# CASE OF LARGE $n$

## OASIS [Chechik et al., 2010]



### Stochastic and distributed optimization

- Assume metric learning problem of the form

$$\min_M \frac{1}{|\mathcal{R}|} \sum_{(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R}} \ell(M, \mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k)$$

- Can use **Stochastic Gradient Descent**
  - Use a random sample (mini-batch) to estimate gradient
  - Better than full gradient descent when  $n$  is large
- Can be combined with **distributed optimization**
  - Distribute triplets on workers
  - Each worker use a mini-batch to estimate gradient
  - Coordinator averages estimates and updates



### Simple workarounds

- Learn a **diagonal matrix**
  - Learn  $d$  parameters
  - Only a weighting of features!
- Learn metric after dimensionality reduction (e.g., PCA)
  - Used in many papers
  - Potential loss of information
  - Learned metric difficult to interpret

### Matrix decompositions



- **Low-rank decomposition**  $M = L^T L$  with  $L \in \mathbb{R}^{r \times d}$ 
  - Learn  $r \times d$  parameters
  - Generally nonconvex, must tune  $r$
- **Rank-1 decomposition**  $M = \sum_{k=1}^K w_k \mathbf{b}_k \mathbf{b}_k^T$ 
  - Learn  $K$  parameters
  - Must choose good basis set
- Special case: sparse data [Liu et al., 2015]
  - Decomposition as rank-1 4-sparse matrices
  - Greedy algorithm incorporating a single basis at each iteration
  - Computational cost independent of  $d$

# METRIC LEARNING FOR STRUCTURED DATA

---

- Each data instance is a **structured object**
  - Strings: words, DNA sequences
  - Trees: XML documents
  - Graphs: social network, molecules

ACGGCTT



- Metrics on structured data are convenient
  - Act as proxy to manipulate complex objects
  - Can use any metric-based algorithm

- Could represent each object by a feature vector
  - Idea behind many kernels for structured data
  - Could then apply standard metric learning techniques
  - Potential loss of structural information
- Instead, focus on **edit distances**
  - Directly operate on structured object
  - Variants for strings, trees, graphs
  - Natural parameterization by cost matrix

- Notations
  - Alphabet  $\Sigma$ : finite set of symbols
  - String  $x$ : finite sequence of symbols from  $\Sigma$
  - $|x|$ : length of string  $x$
  - $\$$ : empty string / symbol

## Definition (Levenshtein distance)

The Levenshtein string edit distance between  $x$  and  $x'$  is the length of the shortest sequence of operations (called an *edit script*) turning  $x$  into  $x'$ . Possible operations are insertion, deletion and substitution of symbols.

- Computed in  $O(|x| \cdot |x'|)$  time by Dynamic Programming (DP)

# STRING EDIT DISTANCE

## Parameterized version

- Use a nonnegative  $(|\Sigma| + 1) \times (|\Sigma| + 1)$  matrix  $C$ 
  - $C_{ij}$ : cost of substituting symbol  $i$  with symbol  $j$

### Example 1: Levenshtein distance

C		\$	a	b
\$		0	1	1
a		1	0	1
b		1	1	0

$\Rightarrow$  edit distance between **abb** and **aa** is 2 (needs at least two operations)


### Example 2: specific costs

C		\$	a	b
\$		0	2	10
a		2	0	4
b		10	4	0

$\Rightarrow$  edit distance between **abb** and **aa** is 10 ( $a \rightarrow \$$ ,  $b \rightarrow a$ ,  $b \rightarrow a$ )

# EDIT PROBABILITY LEARNING

- Interdependence issue
  - The optimal edit script depends on the costs
  - Updating the costs may change the optimal edit script
- Consider **edit probability**  $p(x'|x)$  [Oncina and Sebban, 2006]
  - Cost matrix: probability distribution over operations
  - Corresponds to summing over all possible scripts
- Represent process by a stochastic memoryless transducer
- Maximize expected log-likelihood of positive pairs







## Iterative **Expectation-Maximization** algorithm [Oncina and Sebban, 2006]

- Expectation step
  - Given edit probabilities, compute frequency of each operation
  - Probabilistic version of the DP algorithm
- Maximization step
  - Given frequencies, update edit probabilities
  - Done by likelihood maximization under constraints

$$\forall u \in \Sigma, \sum_{v \in \Sigma \cup \{\$ \}} c_{v|u} + \sum_{v \in \Sigma} c_{v|\$} = 1, \quad \text{with} \quad \sum_{v \in \Sigma} c_{v|\$} + \underbrace{c(\#)}_{\text{exit prob.}} = 1,$$

## Application to handwritten digit recognition [Oncina and Sebban, 2006]



## Some remarks

- Advantages
  - Elegant probabilistic framework
  - Enables data generation
  - Generalization to trees [Bernard et al., 2008]
- Drawbacks
  - Convergence to local minimum
  - Costly: DP algorithm for each pair at each iteration
  - Cannot use negative pairs

## GESL [Bellet et al., 2012]

- Inspired from successful algorithms for non-structured data
  - Large-margin constraints
  - Convex optimization
- Requires key simplification: **fix the edit script**

$$e_c(x, x') = \sum_{u,v \in \Sigma \cup \{\$ \}} c_{uv} \cdot \#_{uv}(x, x')$$

- $\#_{uv}(x, x')$ : nb of times  $u \rightarrow v$  appears in Levenshtein script
- $e_c$  is a linear function of the costs

GESL [Bellet et al., 2012]


## Formulation

$$\begin{aligned}
 \min_{\mathbf{C} \geq 0, \boldsymbol{\xi} \geq 0, B_1 \geq 0, B_2 \geq 0} \quad & \sum_{i,j} \xi_{ij} + \lambda \|\mathbf{C}\|_{\mathcal{F}}^2 \\
 \text{s.t.} \quad & e_{\mathbf{C}}(\mathbf{x}, \mathbf{x}') \geq B_1 - \xi_{ij} \quad \forall (\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{D} \\
 & e_{\mathbf{C}}(\mathbf{x}, \mathbf{x}') \leq B_2 + \xi_{ij} \quad \forall (\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S} \\
 & B_1 - B_2 = \gamma
 \end{aligned}$$

$\gamma$  margin parameter

- **Convex**, less costly and use of negative pairs
- Straightforward adaptation to trees and graphs
- Less general than proper edit distance
  - Chicken and egg situation similar to LMNN

Application to word classification [Bellet et al., 2012]



- Distance / similarity: key component of machine learning
- Metric learning often requires only weak supervision
- Many algorithms:
  - For classification, clustering, ranking...
  - Linear, nonlinear, local metrics
  - Scalable methods
- Very successful in practical applications
- More details: can refer to [\[Bellet et al., 2015\]](#)

# REFERENCES I

- [Bellet et al., 2012] Bellet, A., Habrard, A., and Sebban, M. (2012).  
**Good edit similarity learning by loss minimization.**  
Machine Learning Journal, 89(1):5–35.
- [Bellet et al., 2015] Bellet, A., Habrard, A., and Sebban, M. (2015).  
***Metric Learning.***  
Morgan & Claypool Publishers.
- [Bernard et al., 2008] Bernard, M., Boyer, L., Habrard, A., and Sebban, M. (2008).  
**Learning probabilistic models of tree edit distance.**  
Pattern Recognition, 41(8):2611–2629.
- [Chechik et al., 2010] Chechik, G., Sharma, V., Shalit, U., and Bengio, S. (2010).  
**Large Scale Online Learning of Image Similarity Through Ranking.**  
Journal of Machine Learning Research, 11:1109–1135.
- [Goldberger et al., 2004] Goldberger, J., Roweis, S., Hinton, G., and Salakhutdinov, R. (2004).  
**Neighbourhood Components Analysis.**  
In NIPS, pages 513–520.
- [Hoi et al., 2008] Hoi, S. C., Liu, W., and Chang, S.-F. (2008).  
**Semi-supervised distance metric learning for Collaborative Image Retrieval.**  
In CVPR.



## REFERENCES II

- [Liu et al., 2015] Liu, K., Bellet, A., and Sha, F. (2015).  
**Similarity Learning for High-Dimensional Sparse Data.**  
In AISTATS, pages 653–662.
- [McFee and Lanckriet, 2010] McFee, B. and Lanckriet, G. R. G. (2010).  
**Metric Learning to Rank.**  
In ICML, pages 775–782.
- [Oncina and Sebban, 2006] Oncina, J. and Sebban, M. (2006).  
**Learning Stochastic Edit Distance: application in handwritten character recognition.**  
Pattern Recognition, 39(9):1575–1587.
- [Parameswaran and Weinberger, 2010] Parameswaran, S. and Weinberger, K. Q. (2010).  
**Large Margin Multi-Task Metric Learning.**  
In NIPS, pages 1867–1875.
- [Schultz and Joachims, 2003] Schultz, M. and Joachims, T. (2003).  
**Learning a Distance Metric from Relative Comparisons.**  
In NIPS.
- [van der Maaten and Hinton, 2008] van der Maaten, L. and Hinton, G. (2008).  
**Visualizing Data using t-SNE.**  
Journal of Machine Learning Research, 9:2579–2605.

## REFERENCES III

- [Weinberger et al., 2005] Weinberger, K. Q., Blitzer, J., and Saul, L. K. (2005).  
**Distance Metric Learning for Large Margin Nearest Neighbor Classification.**  
In NIPS, pages 1473–1480.
- [Weinberger and Saul, 2009] Weinberger, K. Q. and Saul, L. K. (2009).  
**Distance Metric Learning for Large Margin Nearest Neighbor Classification.**  
Journal of Machine Learning Research, 10:207–244.
- [Xing et al., 2002] Xing, E. P., Ng, A. Y., Jordan, M. I., and Russell, S. J. (2002).  
**Distance Metric Learning with Application to Clustering with Side-Information.**  
In NIPS, pages 505–512.
- [Ying et al., 2009] Ying, Y., Huang, K., and Campbell, C. (2009).  
**Sparse Metric Learning via Smooth Optimization.**  
In NIPS, pages 2214–2222.
- [Zhang and Yeung, 2010] Zhang, Y. and Yeung, D.-Y. (2010).  
**Transfer metric learning by learning task relationships.**  
In KDD, pages 1199–1208.