

Summary

- i. Stochastic gradient descent
- ii. Anecdotal evidence
- iii. Analysis for a simple case
- iv. The tradeoffs of large scale learning
- v. Simple benchmarks for SGD.
- vi. General convergence results.
- vii. Learning with a single epoch.
- viii. SGD for Neyman Pearson classification.

I. Learning with Stochastic Gradient Descent

Example

Binary classification

- Patterns x .
- Classes $y = \pm 1$
- Examples $z = (x, y)$

Linear model

- Choose features: $\Phi(x) \in \mathbb{R}^d$
- Linear discriminant function: $f_w(x) = \text{sign} \left(w^\top \Phi(x) \right)$

SVM training

- Choose loss function

$$Q(z, w) = \ell(y, f(x, w)) = \text{(e.g.) } \log \left(1 + e^{-y w^\top \Phi(x)} \right)$$

- Cannot minimize the expected risk $E(w) = \int Q(z, w) dP(z)$.
- Can compute the empirical risk $E_n(w) = \frac{1}{n} \sum_{i=1}^n Q(z_i, w)$.



Minimize L_2 regularized empirical risk


$$\min_w \frac{\lambda}{2} \|w\|^2 + \frac{1}{n} \sum_{i=1}^n Q(z_i, w)$$

Choosing λ is the same setting a constraint $\|w\|^2 < B$.

Batch Gradient Descent

Batch: process all examples together

$$\text{Repeat: } w \leftarrow w - \gamma \left(\lambda w + \frac{1}{n} \sum_{i=1}^n \frac{\partial Q}{\partial w}(z_i, w) \right)$$

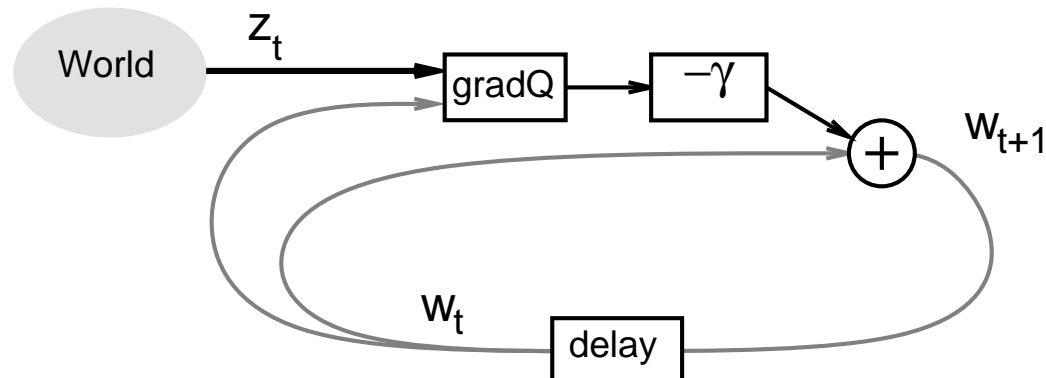


Stochastic Gradient Descent

Online: process examples one by one

Repeat: (a) Pick random example $z_t = (x_t, y_t)$

$$(b) \ w \leftarrow w - \gamma_t \left(\lambda w + \frac{\partial Q}{\partial w}(z_t, w) \right)$$



Stochastic versus Online

Stochastic

- Examples drawn randomly from a finite training set.
- In practice one often perform “epochs”.

Online

- Examples drawn on-the-fly from the real world.
- Adaptive systems.

In fact the same mathematics apply to both cases.

Stochastic/Online versus Generalization

Stochastic gradient descent optimizes

$$E(w) = \mathbb{E} [Q(z, w)]$$

for whatever distribution $dP(z)$ the examples are drawn from.

- If the examples are drawn from a finite training set,
→ SGD optimizes the empirical error $E_n(w)$.
- If the examples are drawn from Nature,
→ SGD optimizes the expected error $E(w)$.

SGD convergence speed statements are generalization results.

SGD Algorithms for everything...

Adaline (Widrow and Hoff, 1960)

$$Q_{\text{adaline}} = \frac{1}{2} (y - w^\top \Phi(x))^2$$
$$\Phi(x) \in \mathbb{R}^d, \quad y = \pm 1$$

$$w \leftarrow w + \gamma_t (y_t - w^\top \Phi(x_t)) \Phi(x_t)$$

Perceptron (Rosenblatt, 1957)

$$Q_{\text{perceptron}} = \max\{0, -y w^\top \Phi(x)\}$$
$$\Phi(x) \in \mathbb{R}^d, \quad y = \pm 1$$

$$w \leftarrow w + \gamma_t \begin{cases} y_t \Phi(x_t) & \text{if } y_t w^\top \Phi(x_t) \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

Multilayer perceptrons (Rumelhart et al., 1986) ...

SVM (Cortes and Vapnik, 1995) ...

Lasso (Tibshirani, 1996)

$$Q_{\text{lasso}} = \lambda |w|_1 + \frac{1}{2} (y - w^\top \Phi(x))^2$$
$$w = (u_1 - v_1, \dots, u_d - v_d)$$
$$\Phi(x) \in \mathbb{R}^d, \quad y \in \mathbb{R}, \quad \lambda > 0$$

$$u_i \leftarrow [u_i - \gamma_t (\lambda - (y_t - w_t^\top \Phi(x_t)) \Phi_i(x_t))]_+$$
$$v_i \leftarrow [v_i - \gamma_t (\lambda + (y_t - w_t^\top \Phi(x_t)) \Phi_i(x_t))]_+$$

with notation $[x]_+ = \max\{0, x\}$.

K-Means (MacQueen, 1967)


$$Q_{\text{kmeans}} = \min_k \frac{1}{2} (z - w_k)^2$$
$$z \in \mathbb{R}^d, \quad w_1 \dots w_k \in \mathbb{R}^d$$
$$n_1 \dots n_k \in \mathbb{N}, \quad \text{initially } 0$$

$$k^* = \arg \min_k (z_t - w_k)^2$$
$$n_{k^*} \leftarrow n_{k^*} + 1$$
$$w_{k^*} \leftarrow w_{k^*} + \frac{1}{n_{k^*}} (z_t - w_{k^*})$$

Adaline

Model: $f(x, w) = \text{StepFunction}(w^\top \Phi(x))$


Loss function: $Q_{\text{adaline}} = \frac{1}{2}(y - w^\top \Phi(x))^2$



Update rule: $w \leftarrow w + \gamma_t(y_t - w^\top \Phi(x_t)) \Phi(x_t)$

(Widrow and Hoff, 1960)

Multilayer Networks



Model:

Acyclic combinations of sigmoid units.

E.g.: $f(x, w) = \sum_j v_j \text{Sigmoid}(u_j^\top x)$ and many variations...

Loss: $Q_{\text{mlp}} = \frac{1}{2}(y - f(x, w))^2$

Update rule:

- Compute $\partial f(x_t, w) / \partial w$ using the chain rule.
- $w \leftarrow w + \gamma_t (y_t - f(x_t, w)) \frac{\partial f(x_t, w)}{\partial w}$

Non differentiability

What if $Q(z, w)$ has a few non differentiable points?

Just pick another example!


$$\frac{\partial E(w)}{\partial w} = \frac{\partial}{\partial w} \int Q(z, w) dP(z) \stackrel{?}{=} \int \frac{\partial Q(z, w)}{\partial w} dP(z)$$

Mild sufficient condition: (bounded convergence theorem)

$$\exists \Phi, \forall z, \forall v \in \mathcal{V}(w), \quad |Q(z, v) - Q(z, w)| \leq |w - v| \Phi(z, w)$$

Rosenblatt's Perceptron

Model: $\text{StepFunction}(w^\top \Phi(x))$




Loss: $Q_{\text{perceptron}} = \max\{0, -y w^\top \Phi(x)\}$

Update rule: $w \leftarrow w + \gamma_t \begin{cases} y_t \Phi(x_t) & \text{if } y_t w^\top \Phi(x_t) \leq 0 \\ 0 & \text{otherwise} \end{cases}$ (Rosenblatt, 1957)

K-Means

Loss: $Q_{\text{kmeans}} = \min_k \frac{1}{2}(z - w_k)^2$




Update rule:
$$\begin{cases} k^* = \arg \min_k (z_t - w_k)^2 \\ n_{k^*} \leftarrow n_{k^*} + 1 \\ w_{k^*} \leftarrow w_{k^*} + \frac{1}{n_{k^*}}(z_t - w_{k^*}) \end{cases} \quad (\text{MacQueen, 1967})$$

Note the “optimal” grain $\gamma_t = 1/n_{k^*}$.

Learning Vector Quantization

Loss:

$$Q_{lvq} = \begin{cases} 0 & \text{if correct} \\ \frac{(x-w^+)^2 - (x-w^-)^2}{\delta(x-w^-)^2} & \text{if } (x-w^+)^2 < (1+\delta)(x-w^-)^2 \\ 1 & \text{otherwise} \end{cases}$$



w^- : closest centroid.

w^+ : closest centroid w/correct class.

Update:

$$\text{if } \begin{cases} x \text{ is misclassified} \\ \text{and } (x-w^+)^2 < (1+\delta)(x-w^-)^2 \end{cases}$$

$$\text{then } \begin{cases} w_{t+1}^- = w_t^- - \gamma_t k_1 (x - w_t^-) \\ w_{t+1}^+ = w_t^+ + \gamma_t k_2 (x - w_t^+) \end{cases}$$

(Kohonen et al. 1982, 1988)

Lasso

L1 Regularization:

$$\min_w \lambda |w|_1 + \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (y_i - w^\top \Phi(x_i))^2$$


Trick: $w = (u_1 - v_1, \dots, u_d - v_d)$ $u_i, v_i \in \mathbb{R}^+$

$$\text{Update Rule: } \begin{cases} u_i \leftarrow \left[u_i - \gamma_t (\lambda - (y_t - w^\top \Phi(x_t)) \Phi_i(x_t)) \right]_+ \\ v_i \leftarrow \left[v_i - \gamma_t (\lambda + (y_t - w_t^\top \Phi(x_t)) \Phi_i(x_t)) \right]_+ \end{cases}$$

using notation $[x]_+ = \max\{0, x\}$.

II. Anecdotal Evidence

Stochastic gradient is slow



Does stochastic gradient converge to the optimum?

Residual noise proportional to learning rate γ_t .

Learning rate cannot decrease too fast.

$$|w_t - w^*|^2 \equiv \frac{1}{t} \quad (\dots \text{at best} \dots)$$

Stochastic gradient is not a good optimization algorithm.


Stochastic gradient is fast

Assume training set contains
10 copies of the 100 same examples.

- **Batch** Blindly computes redundant gradients.
1 epoch on large set \equiv 1 epochs on small set.
- **Online** Take advantage of redundancy.
1 epoch on large set \equiv 10 epochs on small set.

Stochastic gradient **learns** much faster?

1957 - Perceptron



(Frank Rosenblatt, 1957)


Computing

- Relays, potentiometers, electrical motors, ...

Why did Rosenblatt use a stochastic algorithm?

- Because he did not know better? (unlikely)
- Because he did not have computing resources to do anything else?
- Does this mean that a stochastic algorithm does more with less?

1988 - Convolutional networks for OCR



(Le Cun et al., 1989)

Computing

- Sun3 (\approx first generation of Palm Pilot.)
- 9000 training examples, 2000 test examples, three weeks of training.

Why did they use stochastic gradient?

- Because he did not know better? **no** (Becker & Le Cun, 1989)
- Because he did not have computing resources to do anything else?
- Does this means that a stochastic algorithm does more with less?

1995 - Check reading

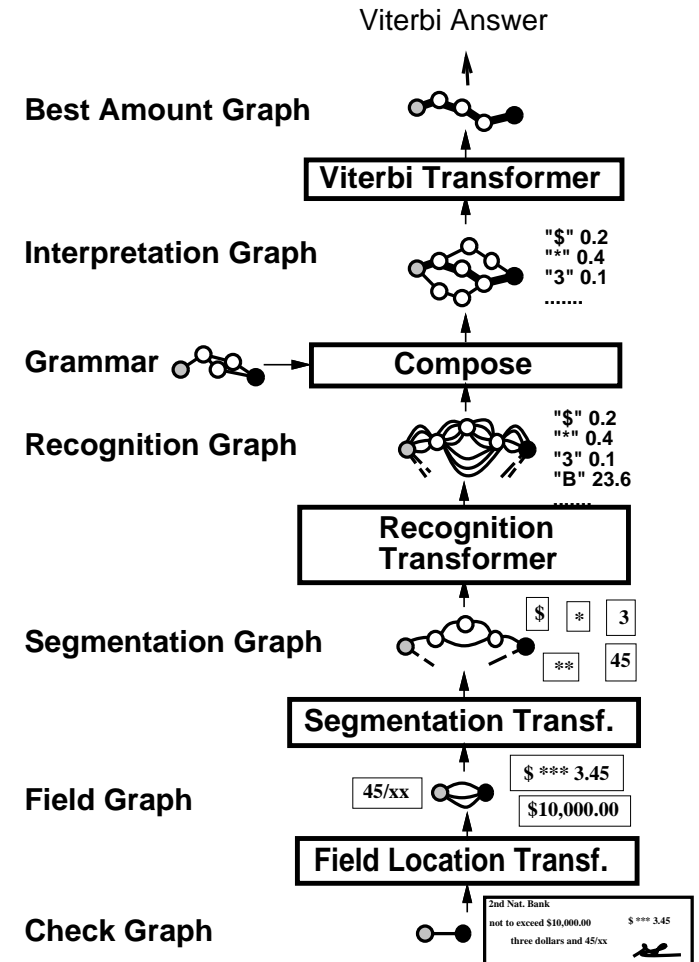
Computing

- Sun4 (\approx your average cell phone.)
- 200K segmented digits.
- 250K unsegmented check images.
- Three weeks of CRF-like training.


Why did we use stochastic gradient?

- Because we did not know better?
- Because he did not have computing resources to do anything else?
- Does this means that a stochastic algorithm does more with less?

(Bottou, LeCun, et al., CVPR 1997)



2010 – Multiple Natural Language Tasks



Computing

- $\approx 1B$ words for unsupervised training set. Six weeks of training.
- $\approx 1M$ words for task dependent training.

Why did they use stochastic gradient?

- Because they did not have computing resources to do anything else?


(Collobert, Weston et al., 2009–2011)

III. Analysis for a simple case

One-dimensional optimization

Simple example

- One dimension $x \in \mathbb{R}$
- $f(w) = \frac{1}{2}Hw^2 \quad w^* = 0$
- $g(w) = \frac{\partial f}{\partial w} + \xi = Hw + \xi$
- $\mathbb{E}[\xi] = 0 \quad \mathbb{E}[\xi^2] = G$



Without stochastic noise ($\xi = 0$)

Gradient descent with constant gain $\gamma_t = \eta$

$$w_t = w_{t-1} - \eta g(w_{t-1})$$

Then

$$\begin{aligned} w_t &= w_{t-1} - \eta H w_{t-1} \\ &= (1 - \eta H) w_{t-1} \\ &= (1 - \eta H)^t w_0 \end{aligned}$$

Error w_t^2 decreases exponentially when constant gain $\eta < 2/H$.

Remark

- H is a positive matrix in the multidimensional case.
The convergence condition is then $\eta < 2/\lambda_{\max}$.

With stochastic noise ($\mathbb{E}[\xi] = 0$, $\mathbb{E}[\xi^2] = G$)

Gradient descent with decreasing gain $\gamma_t = \eta t^{-\alpha}$

$$\begin{aligned} w_t &= w_{t-1} - \eta t^{-\alpha} g(w_{t-1}) \\ &= w_{t-1} - \eta t^{-\alpha} (H w_{t-1} + \xi) \end{aligned}$$

Error recursion

$$\begin{aligned} w_t^2 &= w_{t-1}^2 - 2\eta t^{-\alpha} w_{t-1} (H w_{t-1} + \xi) + \eta^2 t^{-2\alpha} (H w_{t-1} + \xi)^2 \\ \mathbb{E}[w_t^2 | \mathcal{P}_t] &= w_{t-1}^2 - 2\eta t^{-\alpha} w_{t-1} H w_{t-1} + \eta^2 t^{-2\alpha} (H^2 w_{t-1}^2 + G) \\ \mathbb{E}[w_t^2] &= \left(1 - 2\eta H t^{-\alpha} + \eta^2 H^2 t^{-2\alpha}\right) \mathbb{E}[w_{t-1}^2] + \eta^2 G t^{-2\alpha} \end{aligned}$$

Decomposition

$$\begin{aligned} U_t &= \left(1 - 2\eta H t^{-\alpha} + \eta^2 H^2 t^{-2\alpha}\right) U_{t-1} + \eta^2 G t^{-2\alpha} \\ &= U_S \prod_{j=S+1}^t \left(1 - 2\eta H j^{-\alpha} + \eta^2 H^2 j^{-2\alpha}\right) \\ &\quad + \sum_{i=S+1}^t \eta^2 G i^{-2\alpha} \prod_{j=i+1}^t \left(1 - 2\eta H j^{-\alpha} + \eta^2 H^2 j^{-2\alpha}\right) \end{aligned}$$

When $\alpha < 0$

– Everything diverges.

When $\alpha = 0$

- When $2\eta H - \eta^2 H^2 \geq 1$, the error U_t diverges.
- When $2\eta H - \eta^2 H^2 < 1$, the error U_t converges, but not to zero.

Otherwise we pick S large enough to make all the product terms positive.

The Green Term

Case $1/2 < \alpha < 1$

$$\begin{aligned} A_{St} &= \sum_{j=S+1}^t \log \left(1 - 2\eta H j^{-\alpha} + \eta^2 H^2 j^{-2\alpha} \right) \\ &= \sum_{j=S+1}^t -2\eta H j^{-\alpha} + \mathcal{O}(j^{-2\alpha}) \\ &= -2\eta H \int_S^t x^{-\alpha} dx + \mathcal{O}(1) \\ &= -\kappa t^{1-\alpha} + \mathcal{O}(1) \quad \text{where } \kappa \triangleq \frac{2\eta H}{1-\alpha} \end{aligned}$$

Therefore $U_S e^{A_{St}} \equiv e^{-\kappa t^{1-\alpha}} \longrightarrow 0$

The green term converges exponentially to zero.

The case $0 < \alpha \leq 1/2$ is slightly more complicated and not that interesting.

The Green Term

Case $\alpha = 1$

$$\begin{aligned} A_{St} &= \sum_{j=S+1}^t \log \left(1 - 2\eta H j^{-1} + \eta^2 H^2 j^{-2} \right) \\ &= \sum_{j=S+1}^t -2\eta H j^{-1} + \mathcal{O}(j^{-2}) \\ &= -2\eta H \int_S^t x^{-1} dx + \mathcal{O}(1) \\ &= -2\eta H \log(t) + \mathcal{O}(1) \end{aligned}$$

Therefore $U_S e^{A_{St}} \equiv t^{-2\eta H} \longrightarrow 0$

The green term converges polynomially to zero.

The choice of η impacts the degree of the polynomial convergence.

The Green Term

Case $\alpha > 1$

$$\begin{aligned} A_{St} &= \sum_{j=S+1}^t \log \left(1 - 2\eta H j^{-\alpha} + \eta^2 H^2 j^{-2\alpha} \right) \\ &= \sum_{j=S+1}^t -2\eta H j^{-\alpha} + \mathcal{O}(j^{-2\alpha}) \longrightarrow K > -\infty \end{aligned}$$

Therefore $U_S e^{A_{St}} \longrightarrow U_S e^K > 0$

The green term does not converges to zero.

Therefore U_t does not converge to zero either...

The Red Term

Case $\alpha = 1$

$$\begin{aligned} A_{it} &= \sum_{j=i+1}^t \log \left(1 - 2\eta H j^{-1} + \eta^2 H^2 j^{-2} \right) \\ &= -2\eta H \sum_{j=i+1}^t j^{-1} + \mathcal{O}(j^{-2}) = -2\eta H \int_i^t x^{-1} dx + \mathcal{O}(1) \\ &= -2\eta H [\log t - \log i] + \mathcal{O}(1) \end{aligned}$$

Therefore we can write the red term as:

$$\begin{aligned} \sum_{i=S+1}^t \eta^2 G i^{-2} e^{A_{it}} &\equiv t^{-2\eta H} \sum_{i=S+1}^t i^{-2} i^{2\eta H} \\ &\equiv t^{-2\eta H} \left[t^{2\eta H - 1} - S^{2\eta H - 1} \right] \equiv t^{-1} \longrightarrow 0 \end{aligned}$$

The red term converges like t^{-1} .

The Red Term

Case $1/2 < \alpha < 1$

$$\begin{aligned} A_{it} &= \sum_{j=i+1}^t \log \left(1 - 2\eta H j^{-\alpha} + \eta^2 H^2 j^{-2\alpha} \right) \\ &= -2\eta H \sum_{j=i+1}^t j^{-\alpha} + \mathcal{O}(j^{-2\alpha}) = -2\eta H \int_i^t x^{-\alpha} dx + \mathcal{O}(1) \\ &= -2\eta H \left[\frac{x^{1-\alpha}}{1-\alpha} \right]_i^t + \mathcal{O}(1) = -\kappa \left(t^{1-\alpha} - i^{1-\alpha} \right) + \mathcal{O}(1) \end{aligned}$$

Then
$$\sum_{i=S+1}^t \eta^2 G i^{-2\alpha} e^{A_{it}} \equiv e^{-\kappa t^{1-\alpha}} \sum_{i=S+1}^t i^{-2\alpha} e^{\kappa i^{1-\alpha}}$$

We would like to approximate with an integral.

Unfortunately the primitive of $x^{-2\alpha} e^{\kappa x^{1-\alpha}}$ is not obvious.

The Red Term

Case $1/2 < \alpha < 1$ (continued)

Therefore we insert terms that do not affect the asymptotic rate and transform the expression into a known derivative.

$$\begin{aligned} \dots &\equiv e^{-\kappa t^{1-\alpha}} \sum_{i=S+1}^t i^{-2\alpha} e^{\kappa i^{1-\alpha}} \\ &\equiv e^{-\kappa t^{1-\alpha}} \sum_{i=S+1}^t \left(\kappa (1-\alpha) i^{-2\alpha} - (1+\alpha) i^{-\alpha-1} \right) e^{\kappa i^{1-\alpha}} \\ &\equiv e^{-\kappa t^{1-\alpha}} \left[x^{-\alpha} e^{\kappa x^{1-\alpha}} \right]_S^t \\ &\equiv t^{-\alpha} \end{aligned}$$

The red term therefore converges like $t^{-\alpha}$.

Summary

When $\alpha \leq 0$ or $\alpha > 1$

$\mathbb{E} \left[w_t^2 - w^* \right]$ does not converge to zero.

When $1/2 < \alpha < 1$

(and also when $0 < \alpha \leq 1/2$.)

$$\mathbb{E} \left[w_t^2 - w^* \right] \equiv t^{-\alpha}$$

When $\alpha = 1$

$$\mathbb{E} \left[w_t^2 - w^* \right] \equiv t^{-\min\{1, 2\eta H\}}$$

Conclusions from this simple case

Best convergence speed is $\mathbb{E} \left[(w_t - w^*)^2 \right] \equiv t^{-1}$

- This is achieved with gain $\gamma_t = \eta t^{-1}$.
- One should ensure $\eta > 1/2H$ otherwise convergence is much slower.
- Alternatively one can decrease gains a little bit slower, i.e., $\gamma_t = \eta t^{-1+\varepsilon}$ $\varepsilon > 0$, and converge almost as fast.

How good or bad is this convergence speed?

- For an **offline optimization algorithm**, this is very slow.
Batch gradient descent has $\mathbb{E} \left[(w_t - w^*)^2 \right] \equiv e^{-t}$
- For an **online optimization algorithm**, this is expected.
How well can we generalize after seeing only t examples?

IV. The tradeoffs of large scale learning

The Machine Learning Problem

Example: Character Recognition

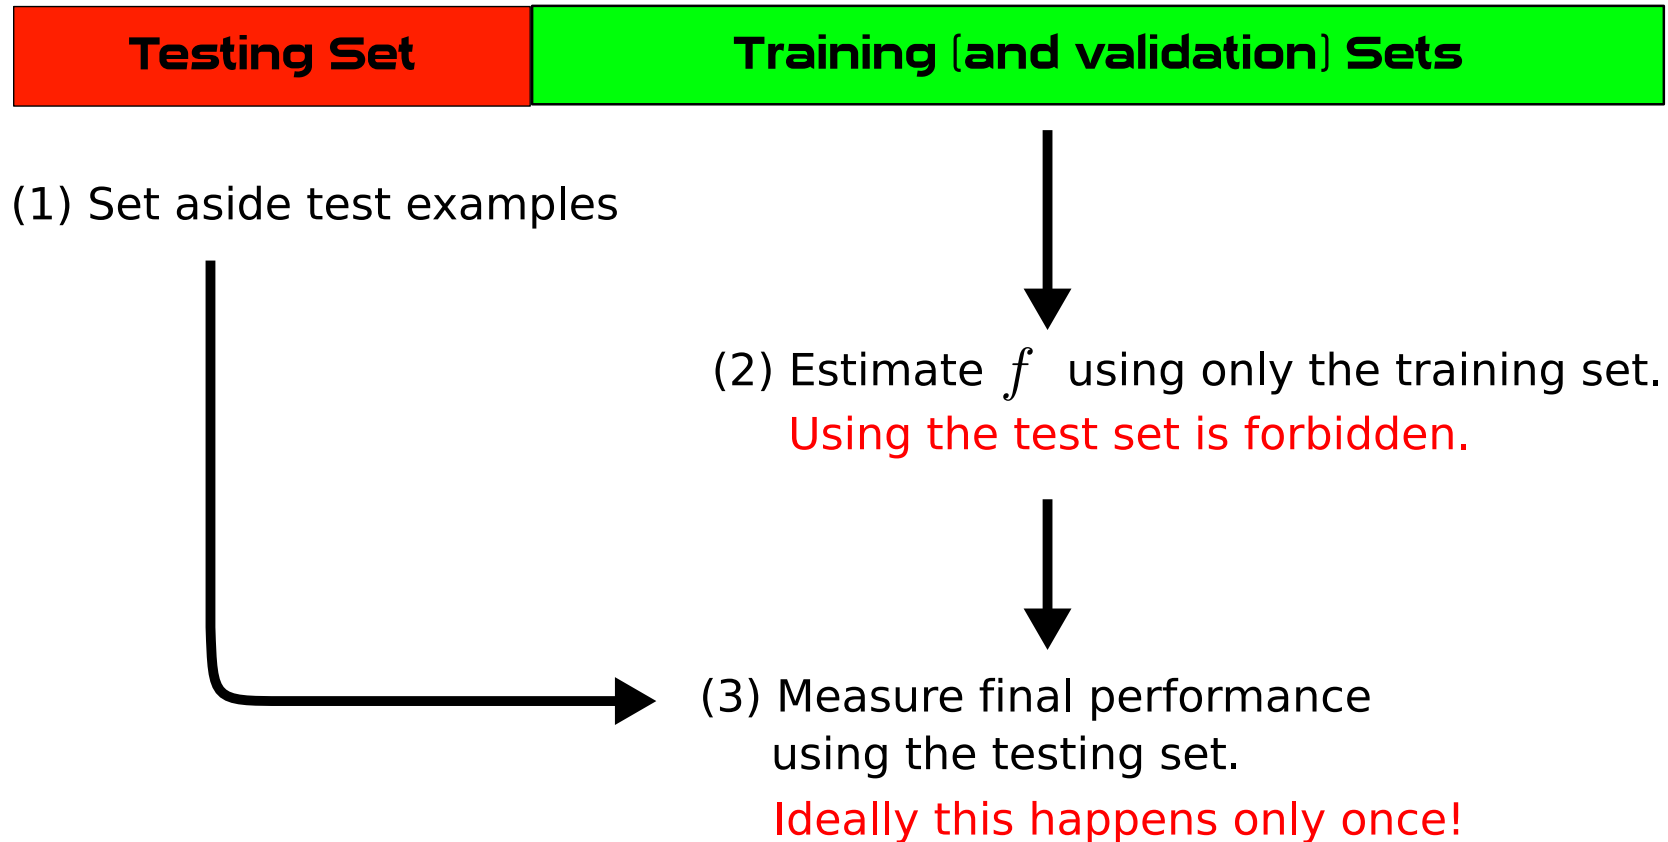


- Find recognition function f on the basis of training examples.
- Function f must work for **all** character variants, not just the training examples.

Issues

- Approximation: How to represent f ?
- Statistics: How many examples do we need to estimate f ?
- Statistics: Generalization \neq Learning by rote.
- Computation: How to compute f efficiently?

The Main Experimental Paradigm



Variations: k -fold cross-validation, etc.

This is the main driver for progress in machine learning.

Mathematical Statement (i)

- **Assumption**

Examples are drawn independently from an unknown probability distribution $P(x, y)$ that represents the laws of Nature.

- **Loss Function**

Function $\ell(\hat{y}, y)$ measures the cost of answering \hat{y} when the true answer is y .

- **Expected Risk**

We seek to find the function f^* that minimizes:

$$\min_f E(f) = \int \ell(f(x), y) dP(x, y)$$

Note: The test set error is an approximation of the expected risk.

Mathematical Statement (ii)

- **Approximation**

Not feasible to search f^* among all functions.

Instead, we search $f_{\mathcal{F}}^*$ that minimizes the Expected Risk $E(f)$ within some richly parametrized family of functions \mathcal{F} .

- **Estimation**

Not feasible to minimize the expectation $E(f)$ because $P(x, y)$ is unknown.


Instead, we search f_n that minimizes the Empirical Risk $E_n(f)$, that is, the average loss over the training set examples.

$$\min_{f \in \mathcal{F}} E_n(f) = \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i)$$

In other words, we optimize a surrogate problem!

Approximation-Estimation Tradeoff

$$\begin{aligned} E(f_n) - E(f^*) &= (E(f_F^*) - E(f^*)) && \text{Approximation Error} \\ &+ (E(f_n) - E(f_F^*)) && \text{Estimation Error} \end{aligned}$$



(e.g. Vapnik, *Statistical Learning Theory*, 1998).

Penalized Empirical Risk

Alternate Formulation

Minimize the Penalized Empirical Risk

$$\min_{f \in \mathcal{F}} \lambda \Omega(f) + E_n(f) = \lambda \Omega(f) + \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i)$$


This can be viewed as minimizing f within

$$\mathcal{F}_C = \{ f \in \mathcal{F} \mid \Omega(f) \leq C \}$$

where C is determined by the choice of the Lagrange coefficient λ .

Typical Example

Support Vector Machine search f with a L_2 penalty term inside a Hilbert space represented using kernel functions.



The Computational Problem

- **Statistical Perspective:**

“It is good to optimize an objective function than ensures a fast estimation rate when the number of examples increases.”

- **Computer Science Perspective:**

“To efficiently solve large problems, it is preferable to choose an optimization algorithm with strong asymptotic properties, e.g. superlinear.”

- **Incorrect Conclusion:**

“To address large-scale learning problems, use a superlinear algorithm to optimize an objective function with fast estimation rate.

- **A finer analysis leads to a dramatically different conclusion.**

The Computational Problem

- Baseline large-scale learning algorithm



Randomly discarding data is the simplest way to handle large datasets.

- What is the statistical benefit of processing more data?
 - What is the computational cost of processing more data?
-
- We need a theory that links Statistics and Computation!
 - 1967: Vapnik's theory does not discuss computation.
 - 1981: Valiant's learnability excludes exponential time algorithms, but (i) polynomial time already too slow, (ii) few actual results.

Learning with Approximate Optimization

Computing $f_n = \arg \min_{f \in \mathcal{F}} E_n(f)$ is often costly.

Since we already optimize a **surrogate** function
why should we compute its optimum f_n exactly?

Let's assume our optimizer returns \tilde{f}_n
such that $E_n(\tilde{f}_n) < E_n(f_n) + \rho$.

For instance, one could stop an iterative
optimization algorithm long before its convergence.

Decomposition of the Error

$$E(\tilde{f}_n) - E(f^*) = E(f_{\mathcal{F}}^*) - E(f^*)$$

Approximation error (\mathcal{E}_{app})

$$+ E(f_n) - E(f_{\mathcal{F}}^*)$$

Estimation error (\mathcal{E}_{est})

$$+ E(\tilde{f}_n) - E(f_n)$$

Optimization error (\mathcal{E}_{opt})

Problem:

Choose \mathcal{F} , n , and ρ to make this as small as possible,

subject to budget constraints $\left\{ \begin{array}{l} \text{max number of examples } n \\ \text{max computing time } T \end{array} \right.$

Note: choosing λ is the same as choosing \mathcal{F} .

Small-scale vs. Large-scale Learning

We can give *formal definitions*.

- **Definition 1:**

We have a **small-scale learning** problem when the **active budget constraint is the number of examples n** .


- **Definition 2:**

We have a **large-scale learning** problem when the **active budget constraint is the computing time T** .

Small-scale Learning

“The active budget constraint is the number of examples.”

- To reduce the estimation error, take n as large as the budget allows.
- To reduce the optimization error to zero, take $\rho = 0$.
- We need to adjust the size of \mathcal{F} .



See Structural Risk Minimization (Vapnik 74) and later works.

Large-scale Learning

“The active budget constraint is the computing time.”

- More complicated tradeoffs.


The computing time depends on the three variables: \mathcal{F} , n , and ρ .

- Example.


If we choose ρ small, we decrease the optimization error. But we must also decrease \mathcal{F} and/or n with adverse effects on the estimation and approximation errors.

- The exact tradeoff depends on the optimization algorithm.
- We can compare optimization algorithms rigorously.

Test Error versus Learning Time




Test Error versus Learning Time




Vary the number of examples. . .

Test Error versus Learning Time



Vary the number of examples, the statistical models, the algorithms,...

Test Error versus Learning Time



Not all combinations are equal.

Let's compare the red curve for different optimization algorithms.

Asymptotic Analysis

$$E(\tilde{f}_n) - E(f^*) = \mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}}$$

Asymptotic Analysis

All three errors must decrease with comparable rates.

Forcing one of the errors to decrease much faster

- would require additional computing efforts,
- but would not significantly improve the test error.

Statistics

Asymptotics of the statistical components of the error

– Thanks to refined uniform convergence arguments

$$\mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}} \sim \mathcal{E}_{\text{app}} + \left(\frac{\log n}{n} \right)^\alpha + \rho$$

with exponent $\frac{1}{2} \leq \alpha \leq 1$.

There are in fact three (four?) types of bounds to consider:

- Classical V-C bounds (pessimistic): $\mathcal{O}\left(\sqrt{\frac{h}{n}}\right)$
- Relative V-C bounds in the realizable case: $\mathcal{O}\left(\frac{h}{n} \log \frac{n}{h}\right)$
- Localized bounds (variance, Tsybakov): $\mathcal{O}\left(\left[\frac{h}{n} \log \frac{n}{h}\right]^\alpha\right)$

Value h describes the *capacity* of our system.

The simplest capacity measure is the *Vapnik-Chervonenkis* dimension of \mathcal{F} .

(Bousquet, 2002; Tsybakov, 2004; Bartlett et al., 2005; ...)

Statistics

Asymptotics of the statistical components of the error

– Thanks to refined uniform convergence arguments

$$\mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}} \sim \mathcal{E}_{\text{app}} + \left(\frac{\log n}{n} \right)^\alpha + \rho$$

with exponent $\frac{1}{2} \leq \alpha \leq 1$.

Asymptotically effective large scale learning

– Must choose \mathcal{F} , n , and ρ such that

$$\mathcal{E} \sim \mathcal{E}_{\text{app}} \sim \mathcal{E}_{\text{est}} \sim \mathcal{E}_{\text{opt}} \sim \left(\frac{\log n}{n} \right)^\alpha \sim \rho .$$

What about optimization times?

First order algorithms

Batch: process all examples together (GD)

– Example: minimization by gradient descent

$$\text{Repeat: } w \leftarrow w - \gamma \left(\lambda w + \frac{1}{n} \sum_{i=1}^n \frac{\partial Q}{\partial w}(x_i, y_i, w) \right)$$

Stochastic: process examples one by one (SGD)

– Example: minimization by stochastic gradient descent

Repeat: (a) Pick random example x_t, y_t

$$(b) \ w \leftarrow w - \gamma_t \left(\lambda w + \frac{\partial Q}{\partial w}(x_t, y_t, w) \right)$$

Second order algorithms

Batch: (2GD)

– Example: Newton's algorithm

$$\text{Repeat: } w \leftarrow w - \mathbf{H}^{-1} \left(\lambda w + \frac{1}{n} \sum_{i=1}^n \frac{\partial Q}{\partial w}(x_i, y_i, w) \right)$$

Stochastic: (2SGD)

– Example: Second order stochastic gradient descent

Repeat: (a) Pick random example x_t, y_t

$$(b) \ w \leftarrow w - \gamma_t \mathbf{H}^{-1} \left(\lambda w + \frac{\partial Q}{\partial w}(x_t, y_t, w) \right)$$

Statistics and Computation

	GD	2GD	SGD	2SGD
Time per iteration :	n	n	1	1
Iters to accuracy ρ :	$\log \frac{1}{\rho}$	$\log \log \frac{1}{\rho}$	$\frac{1}{\rho}$	$\frac{1}{\rho}$
Time to accuracy ρ :	$n \log \frac{1}{\rho}$	$n \log \log \frac{1}{\rho}$	$\frac{1}{\rho}$	$\frac{1}{\rho}$
Time to error ε :	$\frac{1}{\varepsilon^{1/\alpha}} \log^2 \frac{1}{\varepsilon}$	$\frac{1}{\varepsilon^{1/\alpha}} \log \frac{1}{\varepsilon} \log \log \frac{1}{\varepsilon}$	$\frac{1}{\varepsilon}$	$\frac{1}{\varepsilon}$

- 2GD optimizes much faster than GD.
- SGD optimization speed is catastrophic.
- SGD learns faster than both GD and 2GD.
- 2SGD only changes the constants.

V. Experiments with SGD

Benchmarking SGD

Many people associate SGD with trouble

- Historically associated with back-propagation.
- Multilayer networks are very hard problems (nonlinear, nonconvex)
- Notoriously hard to debug (always check the gradients!)
- What is difficult, SGD or MLP?



- Try PLAIN SGD on a simple learning problem.

Download from <http://leon.bottou.org/projects/sgd>.

These simple programs are very short.

Text Categorization with SVMs

- **Dataset**

- Reuters RCV1 document corpus.
- 781,265 training examples, 23,149 testing examples.
- 47,152 TF-IDF features.

- **Task**

- Recognizing documents of category CCAT.

- Minimize $\frac{1}{n} \sum_{i=1}^n \left(\frac{\lambda}{2} w^2 + \ell(w x_i + b, y_i) \right).$

- Update $w \leftarrow w - \eta_t \nabla(w_t, x_t, y_t) = w - \eta_t \left(\lambda w + \frac{\partial \ell(w x_t + b, y_t)}{\partial w} \right)$

Same setup as (Shalev-Schwartz et al., 2007) but plain SGD.

Text Categorization with SVMs

- **Results: Linear SVM**

$$\ell(\hat{y}, y) = \max\{0, 1 - y\hat{y}\} \quad \lambda = 0.0001$$


	Training Time	Primal cost	Test Error
SVMLight	23,642 secs	0.2275	6.02%
SVMPerf	66 secs	0.2278	6.03%
SGD	1.4 secs	0.2275	6.02%

- **Results: Log-Loss Classifier**

$$\ell(\hat{y}, y) = \log(1 + \exp(-y\hat{y})) \quad \lambda = 0.00001$$

	Training Time	Primal cost	Test Error
TRON(LibLinear, $\varepsilon = 0.01$)	30 secs	0.18907	5.68%
TRON(LibLinear, $\varepsilon = 0.001$)	44 secs	0.18890	5.70%
SGD	2.3 secs	0.18893	5.66%

The Wall



More SVM Experiments

From: Patrick Haffner

Date: Wednesday 2007-09-05 14:28:50

...I have tried on some of our main datasets...

I can send you the example, it is so striking!

– Patrick

Dataset	Train size	Number of features	% non-0 features	LIBSVM (SDot)	LLAMA SVM	LLAMA MAXENT	SGDSVM
Reuters	781K	47K	0.1%	210,000	3930	153	7
Translation	1000K	274K	0.0033%	days	47,700	1,105	7
SuperTag	950K	46K	0.0066%	31,650	905	210	1
Voicetone	579K	88K	0.019%	39,100	197	51	1


More SVM Experiments

From: Olivier Chapelle

Date: Sunday 2007-10-28 22:26:44

...you should really run batch with various training set sizes ...

Average Test Loss



Log-loss problem

Batch Conjugate
Gradient on various
training set sizes

Stochastic Gradient
on the full set

Why is SGD near the envelope?

Text Chunking with CRFs

- **Dataset**

- CONLL 2000 Chunking Task:
Segment sentences in syntactically correlated chunks
(e.g., noun phrases, verb phrases.)
- 106,978 training segments in 8936 sentences.
- 23,852 testing segments in 2012 sentences.

- **Model**

- Conditional Random Field (all linear, log-loss.)
- Features are n -grams of words and part-of-speech tags.
- 1,679,700 parameters.

Same setup as (Vishwanathan et al., 2006) but plain SGD.

Text Chunking with CRFs

- Results

	Training Time	Primal cost	Test F1 score
L-BFGS	4335 secs	9042	93.74%
SGD	568 secs	9098	93.75%

- Notes

- Computing the gradients with the chain rule runs faster than computing them with the usual forward-backward algorithm.

Choosing the Gain Schedule

Decreasing gains: $w_{t+1} \leftarrow w_t - \gamma_0(1 + \gamma_0\lambda t)^{-1} \frac{\partial Q}{\partial w}(w_t, x_t, y_t)$

Rationale



- Gain $\gamma_t \equiv \eta t^{-1}$ leads to worst-case rate $t^{-\min\{1, 2\eta\lambda_{\min}\}}$.
We want to avoid the slow convergence case $2\eta\lambda_{\min} < 1$.
- Choosing $\eta = 1/\lambda$ works because $\lambda \geq \lambda_{\min}$.
- We are then left to choose an initial gain γ_0 .

Example: the SVM benchmark

- Choose initial gain γ_0 to make sure that the expected initial updates are comparable with the expected size of the weights.
When $\|x_t\| = 1$, choosing $\gamma_0 = 0.1$ works nicely.

Example: the CRF benchmark

- Choose γ_0 with the secret recipe.

The Secret Recipe

The sample size n does not change the SGD maths!

Constant gain $w_{t+1} \leftarrow w_t - \gamma \frac{\partial Q}{\partial w}(w_t, x_t, y_t)$



At any moment during training, we can:

- Pick a **random subset** of examples with moderate size.
- Try various gains γ on the subsample.
- Pick the gain γ that most reduces the cost.
- Use it for the next 100000 iterations on the **full dataset**.

Examples

- The CRF benchmark code does this to choose γ_0 before training.
- We could also perform such cheap measurements every so often.
The selected gains would then decrease automatically.
- Do not forget to **check the gradients**.

Getting the Engineering Right

The very simple SGD update offers lots of engineering opportunities.



Example: Sparse Linear SVM

The update $w \leftarrow w - \eta(\lambda w + \nabla \ell(wx_i, y_i))$ can be performed in two steps:

- i) $w \leftarrow w - \eta \nabla \ell(wx_i, y_i)$ (sparse, cheap)
- ii) $w \leftarrow w(1 - \eta\lambda)$ (not sparse, costly)

• Solution 1

Represent vector w as the product of a scalar s and a vector v .
Perform (i) by updating v and (ii) by updating s .

• Solution 2

Perform only step (i) for each training example.
Perform step (ii) with lower frequency and higher gain.

VI. General Convergence Results

SGD Algorithms for everything...

Adaline (Widrow and Hoff, 1960)

$$Q_{\text{adaline}} = \frac{1}{2} (y - w^\top \Phi(x))^2$$
$$\Phi(x) \in \mathbb{R}^d, \quad y = \pm 1$$

$$w \leftarrow w + \gamma_t (y_t - w^\top \Phi(x_t)) \Phi(x_t)$$

Perceptron (Rosenblatt, 1957)

$$Q_{\text{perceptron}} = \max\{0, -y w^\top \Phi(x)\}$$
$$\Phi(x) \in \mathbb{R}^d, \quad y = \pm 1$$

$$w \leftarrow w + \gamma_t \begin{cases} y_t \Phi(x_t) & \text{if } y_t w^\top \Phi(x_t) \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

Multilayer perceptrons (Rumelhart et al., 1986) ...

SVM (Cortes and Vapnik, 1995) ...

Lasso (Tibshirani, 1996)

$$Q_{\text{lasso}} = \lambda |w|_1 + \frac{1}{2} (y - w^\top \Phi(x))^2$$
$$w = (u_1 - v_1, \dots, u_d - v_d)$$
$$\Phi(x) \in \mathbb{R}^d, \quad y \in \mathbb{R}, \quad \lambda > 0$$

$$u_i \leftarrow [u_i - \gamma_t (\lambda - (y_t - w_t^\top \Phi(x_t)) \Phi_i(x_t))]_+$$
$$v_i \leftarrow [v_i - \gamma_t (\lambda + (y_t - w_t^\top \Phi(x_t)) \Phi_i(x_t))]_+$$

with notation $[x]_+ = \max\{0, x\}$.

K-Means (MacQueen, 1967)

$$Q_{\text{kmeans}} = \min_k \frac{1}{2} (z - w_k)^2$$
$$z \in \mathbb{R}^d, \quad w_1 \dots w_k \in \mathbb{R}^d$$
$$n_1 \dots n_k \in \mathbb{N}, \quad \text{initially } 0$$

$$k^* = \arg \min_k (z_t - w_k)^2$$
$$n_{k^*} \leftarrow n_{k^*} + 1$$
$$w_{k^*} \leftarrow w_{k^*} + \frac{1}{n_{k^*}} (z_t - w_{k^*})$$

Convergence theory summary

- 1 Convex case
 - General convexity
 - Three proofs
 - Convergence speed
- 2 Nonconvex case
 - Global confinement
 - Convergence to extremal points
- 3 Second order stochastic gradient descent
 - Conditions on scaling matrices

Notations

Expected Risk

$$C(w) \triangleq \mathbb{E}_z L(z, w) \triangleq \int L(z, w) dP(z)$$

Stochastic Gradient Update


$$w_{t+1} = w_t - \gamma_t J(\mathbf{z}_t, w_t)$$

$$\mathbb{E}_{\mathbf{z}} J(\mathbf{z}, w) = \nabla_w C(w)$$

Hessian

$$H(w) = \frac{\partial^2}{\partial w^2} C(w)$$

General Convexity



General convexity assumption.

$$\forall \varepsilon > 0, \quad \inf_{(w-w^*)^2 > \varepsilon} (w - w^*) \nabla_w C(w) > 0$$

- Gradient point towards the right direction.
- Gradient does not vanish (no plateaus).

Learning rates

Requirement:

$$\gamma_t J(\mathbf{z}_t, w_t) \longrightarrow 0$$

Two possibilities:

- Decreasing learning rates: $\gamma_t \rightarrow 0$
- Decreasing gradients: $J(\mathbf{z}_t, w_t) \rightarrow 0$.

Note:

$$\begin{aligned}\mathbb{E}_z \left(J(\mathbf{z}_t, w_t)^2 \right) &\approx \mathbb{E}_z \left(J(\mathbf{z}_t, w^*)^2 \right) + (w_t - w^*)' H(w^*) (w_t - w^*) \\ \mathbb{E}_z \left(J(\mathbf{z}_t, w_t)^2 \right) &\leq A + B(w_t - w^*)^2\end{aligned}$$

In general $A \neq 0$. Therefore **learning rates must decrease**.

Three convergence proofs

I will present three proofs for the following cases.

- Continuous gradient,
- Batch gradient,
- Stochastic gradient

All share the same three-step structure,
but use increasingly sophisticated tools.

Reference: Metivier (1981).

Continuous Gradient: step A.

Differential equation defines $w(t)$:

$$\frac{dw}{dt} = -\nabla_w C(w)$$

Step A: define Lyapunov function:

$$h(t) \triangleq (w(t) - w^*)^2$$

Continuous Gradient: step B.

Step B: Lyapunov function converges

$$\frac{dh}{dt} = 2(w - w^*) \frac{dw}{dt} = -2(w - w^*) \nabla_w C(w) \leq 0$$

Function $h(t)$ is positive, decreasing \implies converges.

Continuous Gradient: step C.

Step C: Continuous gradient converges

Since $h(t)$ converges,

$$\frac{dh}{dt} = -2(w - w^*)\nabla_w C(w) \rightarrow 0$$

General convexity $\implies w(t) \rightarrow w^*$

Batch Gradient: step A.


Update rule defines w_t :

$$w_{t+1} = w_t - \gamma_t \nabla_w C(w)$$

Step A: define Lyapunov sequence:

$$h_t \triangleq (w_t - w^*)^2$$

Batch Gradient: lemma for step B.



$$S_t^+ = \sum_{i=1}^{t-1} [u_{t+1} - u_t]_+ \quad \text{with} \quad [x]_+ = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise.} \end{cases}$$

Bounded positive variations convergence theorem:

$$\left. \begin{array}{l} \forall t, u_t > 0 \\ S_t^+ \text{ converges} \end{array} \right\} \implies u_t \text{ converges.}$$

Batch Gradient: step B.

Step B: Convergence of Lyapunov sequence:

 Additional assumptions.

$$\left. \begin{array}{l} \sum \gamma_t^2 \text{ converges} \\ (\nabla_w C(w))^2 < A + B(w - w^*)^2 \end{array} \right\} \implies h_t \text{ converges.}$$

Proof (assuming $B = 0$):

$$\begin{aligned} [h_{t+1} - h_t]_+ &= [-2\gamma_t (w_t - w^*) \nabla_w C(w_t) + \gamma_t^2 (\nabla_w C(w_t))^2]_+ \\ &\leq \gamma_t^2 (\nabla_w C(w_t))^2 \leq \gamma_t^2 A \end{aligned}$$

and apply the bounded positive variations convergence theorem.

Proof sketch (assuming $B \neq 0$):

Define $\mu_t = \prod_{i=1}^t (1 - \gamma_i^2 B)$.
Show that $h'_t = h_t / \mu_t$ converges.
Then show that h_t converges as well.

Batch Gradient: step C.

Step C: Batch gradient converges.

 Additional assumption:

$$\sum \gamma_t = \infty \implies w_t \rightarrow w^*$$

Proof:

$$h_{t+1} - h_t = -2\gamma_t (w_t - w^*) \nabla_w C(w_t) + \gamma_t^2 (\nabla_w C(w_t))^2$$

Both blue terms have convergent sums.

Therefore $\sum \gamma_i (w_i - w^*) \nabla_w C(w_i)$ converges.

The assumption then means that $(w_t - w^*) \nabla_w C(w) \rightarrow 0$.

The general convexity hypothesis then implies the convergence.

Stochastic Gradient: step A.

Stochastic update rule defines w_t :

$$w_{t+1} = w_t - \gamma_t J(\mathbf{z}_t, w_t)$$

Step A: define Lyapunov process:

$$h_t \triangleq (w_t - w^*)^2$$

Stochastic Gradient: Quasi-martingales.

Let u_t be a stochastic process.

Let \mathcal{P}_t represent what is known on time t , (u_0, \dots, u_t)

Define indicator of expected positive differences.

$$\delta_t = \begin{cases} 1 & \text{if } \mathbb{E}(u_{t+1} - u_t | \mathcal{P}_t) > 0 \\ 0 & \text{otherwise.} \end{cases}$$

and the sum

$$S_t^+ = \sum_{i=1}^t \mathbb{E}(\delta_i (u_{i+1} - u_i))$$

Positive quasi-martingale convergence theorem:

$$\left. \begin{array}{l} \forall t, u_t > 0 \\ S_t^+ \text{ converges} \end{array} \right\} \implies u_t \text{ converges almost surely.}$$

Stochastic Gradient: step B.

Step B: Convergence of Lyapunov process:

 Additional assumptions.

$$\left. \begin{array}{l} \sum \gamma_t^2 \text{ converges} \\ \mathbb{E} \left(J(\mathbf{z}, w)^2 \right) < A + B(w - w^*)^2 \end{array} \right\} \implies h_t \text{ converges a.s.}$$

Proof (assuming $B = 0$):

$$\begin{aligned} \mathbb{E}(h_{t+1} - h_t \mid \mathcal{P}_t) &= -2 \gamma_t (w_t - w^*) \mathbb{E}(J(\mathbf{z}_t, w_t) \mid \mathcal{P}_t) + \gamma_t^2 \mathbb{E}(J(\mathbf{z}_t, w_t)^2 \mid \mathcal{P}_t) \\ &= -2 \gamma_t (w_t - w^*) \nabla_w C(w_t) + \gamma_t^2 \mathbb{E}_{\mathbf{z}}(J(\mathbf{z}_t, w_t)^2) \\ &\leq \gamma_t^2 A \end{aligned}$$

$$\mathbb{E}(\delta_t(h_{t+1} - h_t)) = \mathbb{E}(\delta_t \mathbb{E}(h_{t+1} - h_t \mid \mathcal{P}_t)) \leq \mathbb{E}(\mathbb{E}(h_{t+1} - h_t \mid \mathcal{P}_t)) \leq \gamma_t^2 A$$

and apply the quasi-martingale convergence theorem.

The case $B \neq 0$ can be treated as suggested for batch gradient.

Stochastic Gradient: step C.

Step B: Stochastic gradient converges almost surely.

 Additional assumption:

$$\sum \gamma_t = \infty \implies w_t \xrightarrow{a.s.} w^*$$

Proof:

$$\mathbb{E}(h_{t+1} - h_t) = -2\gamma_t \mathbb{E}((w_t - w^*) \nabla_w C(w_t)) + \gamma_t^2 \mathbb{E}_{\mathbf{z}}(J(\mathbf{z}_t, w_t)^2)$$

Both blue terms have convergent sums.

Therefore $\sum \gamma_i \mathbb{E}((w_i - w^*) \nabla_w C(w_i))$ converges.

The assumption then means that $\liminf (w_t - w^*) \nabla_w C(w) = 0$ (a.s.)

The general convexity hypothesis then implies the convergence.

Nonconvex case

Why worry about nonconvex objective functions?

- Multilayer networks.
- Clustering algorithms.
- Learning features.
- Semi-supervised learning.
- Mixture models.
- Hidden Markov Models.
- Selecting features (some).
- Transfer learning.

Nonconvexity issues

- Several local minima.
- Possibly several global minima.
- Critical subspaces.
- Saddle points.

Local confinement

- Partition space into attraction basins.
- Make sure w_t is confined to a specific basin.
- Define suitable Lyapunov functions on the basin

Reference: Krasovskii (1963)

Batch gradient: Confinement works well.

Stochastic gradient: Confinement works poorly.

Stochastic noise can always get w_t out of any attraction basin with small but non zero probability.

Standard assumptions

$$C(w) > 0$$

$C(w)$ has continuous second derivatives.

$$\sum_i \gamma_i = \infty$$

$$\sum_i \gamma_i \text{ converges.}$$

$$\mathbb{E}_{\mathbf{z}} \left(J(\mathbf{z}, w)^2 \right) \leq A + Bw^2$$

Global confinement

Additional Assumptions

$$\left. \begin{array}{l} \exists D > 0, \quad \inf_{w^2 > D} w \nabla_w C(w) > 0 \\ \exists E > D, \quad \sup_{w^2 < E} J(\mathbf{z}, w)^2 < K \end{array} \right\} \implies \mathbf{P} \left\{ \exists t_0, \forall t > t_0, w_t^2 < 2E \right\} = 1$$

Proof sketch:

Define Lyapunov process $h_t = \max(E, w_t^2)$.

Performs steps [a](#), [b](#), and [c](#). Conclude.

Benefits:

We are now working in a compact set $\{w_t^2 < E\}$.

Continuity implies that all continuous quantities are bounded a.s.

Convergence to Extremal Points

$$\dots \implies \begin{cases} C(w_t) \text{ converges to some value } C_\infty. \\ \nabla_w C(w_t) \xrightarrow{a.s.} 0 \end{cases}$$


Proof sketch:

- a- Define Lyapunov process $h_t = C(w_t)$.
- b- Lyapunov process converges.
- c- Define secondary Lyapunov process $g_t = (\nabla_w C(w_t))^2$.
Perform steps a, b, and c again.

Note:

We must either use global confinement
or assume suitable bounds on $\|\nabla \nabla C(w)\|$.

Convergence to Extremal Points



- Local maxima and saddle points are unstable.
But one can construct cases without noise.
- Infinite plateaus are ruled out by global confinement.
Otherwise they can spoil the day.

Second Order Stochastic Gradient Descent

Second Order Stochastic Gradient Update

$$w_{t+1} = w_t - \Gamma_t J(\mathbf{z}_t, w_t)$$

$$\mathbb{E}_{\mathbf{z}} J(\mathbf{z}, w) = \nabla_w C(w)$$

Gain matrix

The gain is now a positive definite matrix Γ_t .

The idea is to have $\Gamma_t \approx \frac{1}{t} H(w^*)^{-1}$.

But convergence occurs under much weaker assumptions.

Assumptions

Gain Assumptions

Let $0 < \lambda_t^{\min} \leq \lambda_t^{\max}$ bracket the eigenvalues of Γ_t .

$$\sum_t \lambda_t^{\min} = +\infty$$

$$\sum_t (\lambda_t^{\max})^2 \text{ converges.}$$

More Assumptions

$C(w)$ is positive

$C(w)$ has bounded second derivatives

$$\mathbb{E}_{\mathbf{z}} \left(J(\mathbf{z}, w)^2 \right) \leq A + Bw^2$$

Convergence

$$\dots \implies \begin{cases} C(w_t) \text{ converges to some value } C_\infty. \\ \nabla_w C(w_t) \xrightarrow{a.s.} 0 \end{cases}$$

Proof sketch:

- a- Define Lyapunov process $h_t = C(w_t)$.
- b- Lyapunov process converges.
- c- Define secondary Lyapunov process $g_t = (\nabla_w C(w_t))^2$.
Perform steps a, b, and c again.

Remarks:

- We did not say anything about the determination of Γ_t .
Any reasonable approximation of the inverse Hessian works!

Conclusion


We have a very well oiled machinery to establish the convergence of stochastic gradient algorithms under mild conditions.

General convergence results apply to:

- adalines
- multilayer networks
- perceptrons
- kmeans
- lvq2
- ...

VII. Learning with a Single Epoch

Batch and online paths




Effect of one Additional Example (i)

Compare

$$w_n^* = \arg \min_w E_n(w)$$

$$w_{n+1}^* = \arg \min_w E_{n+1}(w) = \arg \min_w \left[E_n(w) + \frac{1}{n} \ell(f(x_{n+1}, w), y_{n+1}) \right]$$



Effect of one Additional Example (ii)

First Order Calculation

$$w_{n+1}^* = w_n^* - \frac{1}{n} H_{n+1}^{-1} \frac{\partial \ell(f(x_n, w_n), y_n)}{\partial w} + \mathcal{O}\left(\frac{1}{n^2}\right)$$

where H_{n+1} is the empirical Hessian on $n+1$ examples.

Compare with Second Order Stochastic Gradient Descent

$$w_{t+1} = w_t - \frac{1}{t} H^{-1} \frac{\partial \ell(f(x_n, w_n), y_n)}{\partial w}$$

Could they converge with the same speed?

– We need to generalize our simple one-dimensional analysis.

Speed of Scaled Stochastic Gradient (i)

Scaled Stochastic Gradient

$$w_{t+1} = w_t - \frac{1}{t} B_t \frac{\partial \ell(f(x_n, w_t), y_n)}{\partial w} + \mathcal{O}\left(\frac{1}{t^2}\right)$$

with $B_t \rightarrow B \succ 0$, $BH \succ I/2$,
and H is the Hessian in w^* .

1- Establish convergence a.s.

- Using the general convergence results.

2- Define error term

- Let $U_t \triangleq H (w_t - w^*) (w_t - w^*)^\top$.
- Observe $E(w_t) - E(w^*) = \text{tr}(U_t) + o(\text{tr}(U_t))$

Speed of Scaled Stochastic Gradient (ii)

3- Derive error recursion

$$\mathbb{E}[U_{t+1}] = \left[I - \frac{2BH}{t} + o\left(\frac{1}{t}\right) \right] \mathbb{E}[U_t] + \frac{HBGB}{t^2} + o\left(\frac{1}{t^2}\right)$$

where the Fisher matrix

$$G \triangleq \int \left[\left(\frac{\partial \ell(f(x, w^*), y)}{\partial w} \right) \left(\frac{\partial \ell(f(x, w^*), y)}{\partial w} \right)^\top \right] dP(x, y).$$

4- Establish lemma on sequences satisfying

$$u_{t+1} = \left(1 + \frac{\alpha}{t} + o\left(\frac{1}{t}\right) \right) u_t + \frac{\beta}{t^2} + o\left(\frac{1}{t^2}\right)$$

- When $\alpha > 1$ show $u_t = \frac{\beta}{\alpha-1} \frac{1}{t} + o\left(\frac{1}{t}\right)$ (nasty proof!).
- When $\alpha < 1$ show $u_t \sim t^{-\alpha}$ (up to log factors).

Speed of Scaled Stochastic Gradient (iii)

5- Bracket $\mathbb{E} [\text{trace} (()) U_{t+1})]$

- Using two sequences of the aforementioned type.
- The idea is to bracket the eigenvalues of BH .

6- Apply the lemma and conclude:

$$\begin{aligned} & \frac{\text{tr}(HBGB)}{2\lambda_{BH}^{\max} - 1} \frac{1}{t} + o\left(\frac{1}{t}\right) \\ & \leq \mathbb{E}[E(w_t) - E(w^*)] \leq \\ & \frac{\text{tr}(HBGB)}{2\lambda_{BH}^{\min} - 1} \frac{1}{t} + o\left(\frac{1}{t}\right) \end{aligned}$$

- Note that we know the constants.
- Interesting special cases: $B = I/\lambda_H^{\min}$ and $B = H^{-1}$.


Asymptotic Efficiency of 2SGD.

“Empirical optima”

“Second-order SGD”

$$n \mathbb{E}[E(f_{w_n^*}) - E(f_{\mathcal{F}})] = \lim_{t \rightarrow \infty} t \mathbb{E}[E(f_{w_t}) - E(f_{\mathcal{F}})]$$

$$\lim_{n \rightarrow \infty} n \mathbb{E}[\|w_{\infty}^* - w_n^*\|^2] = \lim_{t \rightarrow \infty} t \mathbb{E}[\|w_{\infty} - w_t\|^2]$$





Optimal Learning in One Pass



**A Single Pass of Second Order Stochastic Gradient
generalizes as well as the Empirical Optimum.**

Experiments on synthetic data



Unfortunate Issue

Second order SGD is costly

Repeat: (a) Pick random example x_t, y_t

$$(b) \ w \leftarrow w - \gamma_t \mathbf{H}^{-1} \frac{\partial \ell(f(x_t, w), y_t)}{\partial w}$$

- Estimate and store $d \times d$ matrix \mathbf{H}^{-1} .
- Multiply the gradient for each example by this matrix \mathbf{H}^{-1} .

Solutions

Limited storage approximations of H^{-1}



- Diagonal Gauss-Newton (Becker and Lecun, 1989)
- Low rank approximation [oLBFGS], (Schraudolph et al., 2007)
- Diagonal approximation [SGDQN], (Bordes et al., 2009)

Averaged stochastic gradient

- Perform SGD with slowly decreasing gains, e.g. $\gamma_t \sim t^{-0.75}$.
- Compute averages $\bar{w}_{t+1} = \frac{t}{t+1}\bar{w}_t + \frac{1}{t+1}w_{t+1}$
- Same asymptotic speed as second order SGD (Polyak and Juditsky, 92)
- Can take a while to “reach” the asymptotic regime.

Experiment: ALPHA dataset



- From the 2008 Pascal Large Scale Learning Challenge.
- Loss: $\ell(\hat{y}, y) = (\max\{0, 1 - y \hat{y}\})^2$.



Optimal expected risk after a single epoch!

Experiment: Conditional Random Field

- CRF for the CONLL 2000 Chunking task.
- 1.7M parameters. 107,000 training segments.



SGDQN more attractive than ASGD.

Training times: 500s (SGD), 150s (ASGD), 75s (SGDQN).

Standard LBFGS optimizer needs 72 minutes.

Getting the Engineering Right



Gains

- SGD $\gamma_t = \gamma_0(1 + \gamma_0\lambda t)^{-1}$.
- SGDQN $\gamma_t = \gamma_0(1 + \gamma_0\lambda t)^{-1}$.
- ASGD: $\gamma_t = \gamma_0(1 + \gamma_0\lambda t)^{-0.75}$

Sparsity preserving implementation

- SGD: Write $\mathbf{w}_t = s_t \mathbf{W}_t$ with $s_t = \prod_{i=0}^{t-1} (1 - \gamma_i \lambda)$.
- ASGD: Write $\mathbf{w}_t = s_t \mathbf{W}_t$ with $s_t = \prod_{i=0}^{t-1} (1 - \gamma_i \lambda)$
and $\bar{\mathbf{w}}_t = (\mathbf{U}_t + r_t \mathbf{W}_t)/t$ with $r_t = \sum_{i=1}^t s_i$.
- SGDQN: Decoupled loss and regularization updates.

Learning in a single epoch

One epoch learning in practice

- For convex objective functions,
given a sufficiently large dataset,
one pass learning can be achieved in practice.

Other considerations

- Shuffling the data can be hard.
- How large a dataset we need to reach this regime?

Reduced claim

- A couple epochs should suffice.
(e.g., five...)

VIII. SGD for Neyman-Pearson classification

with Gilles Gasso, Aristidis Pappaiaonnou, and Marina Spivak.

ACM TIST 2(3), 2011.

Asymmetric cost problem

Binary classification.

- Positive class $y = +1$, negative class $y = -1$.

Examples of positive classes.

- fraudulent credit card transaction
- relevant document for a given query
- heart failure detection

Different kinds of errors have different costs.

- False positive, false detection, false alarm.
- False negative, non detection.

Imbalanced datasets.

Asymmetric cost problem

Misclassification probabilities

- Probability of non detection (false negatives)

$$P_{\text{nd}}(f) \triangleq \mathbb{P} \{ f(x) \leq 0 \mid y = +1 \}$$

- Probability of false alarm (false positives)

$$P_{\text{fa}}(f) \triangleq \mathbb{P} \{ f(x) \geq 0 \mid y = -1 \}$$

Asymmetric classification (AC)

$$\min_f C_+ P_{\text{nd}}(f) + C_- P_{\text{fa}}(f)$$

Neyman-Pearson Classification

Costs are difficult to assess

- Example: the pacemaker problem.

Neyman-Pearson classification (NP)

- It is often more convenient to solve

$$\min_f P_{\text{nd}}(f) \quad \text{subject to} \quad P_{\text{fa}}(f) < \rho$$

AC and NP classification share the same solutions

- Let $r^*(x) = \mathbb{P}\{y = +1|x\}$

$$f_{AC}^*(x) = r^*(x) - C_- / (C_+ + C_-)$$


$$f_{NP}^*(x) = r^*(x) - \min \{r \text{ such that } P_{\text{fa}}(r^* - r) < \rho\}$$

Adjusting only the Decision Threshold

Decision threshold

- Discriminant function $f_b(x) = f(x) - b$.
- The function f is fixed; we are only adjusting the threshold b .
- True positives: $F_+(b) = \mathbb{P}\{f(x) - b \geq 0 | Y = +1\} = 1 - P_{nd}(f_b)$
- False positives: $F_-(b) = \mathbb{P}\{f(x) - b \geq 0 | Y = -1\} = P_{fa}(f_b)$

ROC curve



Learning the classifier

Empirical counterparts of P_{nd} and P_{fa}


$$\tilde{P}_{\text{nd}}(f) = \frac{1}{n_+} \sum_{y_i > 0} \mathbb{I}\{f(x_i) \leq 0\} \quad \tilde{P}_{\text{fa}}(f) = \frac{1}{n_-} \sum_{y_i < 0} \mathbb{I}\{f(x_i) \geq 0\}$$

Empirical Neyman-Pearson problem

$$\min_w \tilde{P}_{\text{nd}}(f_w) \quad \text{subject to} \quad \tilde{P}_{\text{fa}}(f_w) < \rho$$

The step functions make a very hard optimization problem.

Learning the classifier



$$\hat{P}_{\text{nd}}(f) = \frac{1}{n_+} \sum_{y_i > 0} \sigma(y_i f(x_i)) \quad \hat{P}_{\text{fa}}(f) = \frac{1}{n_-} \sum_{y_i < 0} \sigma(y_i f(x_i))$$

Empirical Neyman-Pearson problem revisited

$$\min_w \hat{P}_{\text{nd}}(f_w) \quad \text{subject to} \quad \hat{P}_{\text{fa}}(f_w) < \rho$$

Regularized variant

$$\min_w \frac{\lambda}{2} \|w\|^2 + \hat{P}_{\text{nd}}(f_w) \quad \text{subject to} \quad \hat{P}_{\text{fa}}(f_w) < \rho$$

Nonconvex objective function with nonconvex constraints

Nonconvex optimization with constraints

Lagrangian

$$\mathcal{L}(f, \mu) = \frac{\lambda}{2} \Omega(f) + \hat{P}_{\text{nd}}(f) + \mu(\hat{P}_{\text{fa}}(f) - \rho) \quad \mu \geq 0$$

Two useful theorems

- The local saddle points of \mathcal{L} are feasible local minima of the constrained problem.
- Assuming **differentiability** and **KKT qualification** the feasible local minima of the constrained problem are critical points of the Lagrangian \mathcal{L} . (Ciarlet, 1989)

Let us find local saddle points of \mathcal{L} . . .

Uzawa algorithm

Uzawa algorithm

- Set initial values for f , $\mu > 0$.
- Choose **very small** gain ν .
- Repeat

$$\begin{aligned} f &\leftarrow \arg \min_f \mathcal{L}(f, \mu) \\ \mu &\leftarrow \mu \left(1 + \nu \frac{\partial \mathcal{L}(f, \mu)}{\partial \nu} \right) \end{aligned}$$

Convergence of Uzawa algorithm

- Nonobvious because $f_\mu^* = \arg \min_f \mathcal{L}(f, \mu)$ can be noncontinuous.
- However we can show that $\hat{P}_{fa}(f_\mu^*)$ is a nonincreasing function of μ .
- Therefore the sign of $\frac{\partial \mathcal{L}(f, \mu)}{\partial \nu}$ indicates the right direction.

Stochastic NP algorithm

Stochastic NP algorithm

- Set initial values for f , $\mu > 0$.
- Choose decreasing gains γ_t .
- Choose **very small** gain ν .
- Repeat
 - Pick random training example x_t, y_t
 - Set $a_t = \begin{cases} n/n_+ & \text{if } y_i = +1 \\ \mu n/n_- & \text{if } y_i = -1 \end{cases}$
 - Update $w \leftarrow (1 - \gamma_t \lambda)w - \gamma_t a_t \sigma'(y_t f_w(x_t)) \frac{\partial f}{\partial w}(x_t)$
 - If $y_i = -1$ update $\mu \leftarrow \mu (1 + \nu (\sigma(y_t f_w(x_t)) - \rho))$

Convergence, etc...

- This is not a gradient descent algorithm (saddle point \neq minimum.)
- Such stochastic approximations have been studied extensively
e.g. (Tsypkin, 1971; Andrieu et al. 2007)

Rebalancing the data

Imbalanced Dataset

- Suppose positive examples from a small proportion of the training set. (Asymmetric costs and imbalanced data often come together.)

Rebalancing

- Instead we draw positive and negative examples with equal probabilities.
- \hat{P}_{nd} , \hat{P}_{fa} are not affected: they depend on $\mathbb{P}\{f(x)|y\}$, not $\mathbb{P}\{y\}$.
- Lagrange coefficient μ auto-adjusts to satisfy the constraint.

$$\mathcal{L}(f, \mu) = \frac{\lambda}{2} \Omega(f) + \hat{P}_{nd}(f) + \mu(\hat{P}_{fa}(f) - \rho) \quad \mu \geq 0$$


- Potential improvement: adjust the initial value of μ .

Equivalently: redefine $a_t = \begin{cases} 1 & \text{if } y_i = +1 \\ \mu & \text{if } y_i = -1 \end{cases}$

Results: Pageblocks

10 features, 4913 negatives, 560 positives.

Linear, 50% training set, 25% validation set, 25% test set.



ONP-SVM: stochastic Neyman Pearson classification.

BNP-SVM: batch Neyman Pearson.

AC-SVM: asymmetric cost SVMs (Davenport et al., 2010)


GEN: generative method (Huang et al., 2006)

SVMPerf: (Joachims, 2005)

Results: Covertypes

54 features, 211840 negatives, 20510 positives.

Linear, 50% training set, 25% validation set, 25% test set.



ONP-SVM: stochastic Neyman Pearson classification.

AC-SVM: asymmetric cost SVMs (Davenport et al., 2010)


GEN: generative method (Huang et al., 2006)

AC-SVM was modified to use the new liblinear solver (Hsieh et al., 2008) whose speed is comparable with SGD.

Results: Covertypes

54 features, 211840 negatives, 20510 positives.

Linear, 50% training set, 25% validation set, 25% test set.



ONP-SVM: stochastic Neyman Pearson classification.

SVMPerf: (Joachims, 2005)

Extension: Q-value optimization

(Spivak et al., 2009, 2010)

Problem

- Mass spectrometer analyses a preparation and produces a lot of spectras.
- Spectras are matched against dictionaries of known peptides.
- Which matches are worth verifying in costly wet lab experiments?

Approach

- Augment dictionary with a large number of **decoy peptides**.
- Matches against decoys are known negatives.
- Construct a classifier that returns as many good matches as possible subject to a constraint on the proportion (q-value) of decoys returned among the posited good matches.

Extension: Q-value optimization

Q-value optimization

$$\min_f P_{\text{nd}}(f) \quad \text{subject to} \quad P_{\text{fa}}(f) < q(1 - P_{\text{nd}}(f))$$

Same approach as NP classification

- Write the Lagrangian
- Create a stochastic approximation to the Uzawa algorithm.

Results

- 70K true matches, 70K decoy matches.
- Same features and same model as the state-of-the-art QRanker system:
A small multilayer network with 5 hidden units.

q	QRanker	This
0.0025	4449	5005
0.01	5462	5707
0.1	7473	7491

Number of positives returned for various q-values.