

Intro: Bigdata/SQL

Andrei Arion, LesFurets.com, tp-bigdata@lesfurets.com

Plan

1. Intro
2. From SQL to NoSQL
3. Scalling MySQL : @LesFurets.com
4. TP: PostgreSQL

Background

- XML databases
- software developper/consultant
- data engineer @LF



tp-bigdata@lesfurets.com

Planning

16/11 : Intro + PostgreSQL

22/11 : MongoDB Intro + Data Modelling

23/11 : MongoDB Application

30/11 : Cassandra Intro + Replication

06/12 : Cassandra modeling (timeseries)

07/12 : Apache Spark Intro + Shell

13/12 : Apache Spark 2 - Netbook

14/12 : Apache Spark 3 - MLlib

04/01 : Projet 1

10/01 : Projet 2

11/01 : Projet 3

18/01 : Projet 4

25/01, 01/02 : Soutenances

- **LesFurets.com:**
 - Andrei Arion, Geoffrey Berard , Charles Herriau
- **Télécom ParisTech:**
 - Ons Jelassi, Siwar Garouachi, Guillaume Pape

How

- Intro: 1h-1h30
- TP: 30 min + 2h30 pratique
- Travail en petit group (2p TP, 3p Projet)
- Évaluation en continuu (questions/etapes projet)
- Soutenance 30 minutes + 5 min questions

Plan

1. Intro

2. From SQL to NoSQL

1. Scalling a simple application

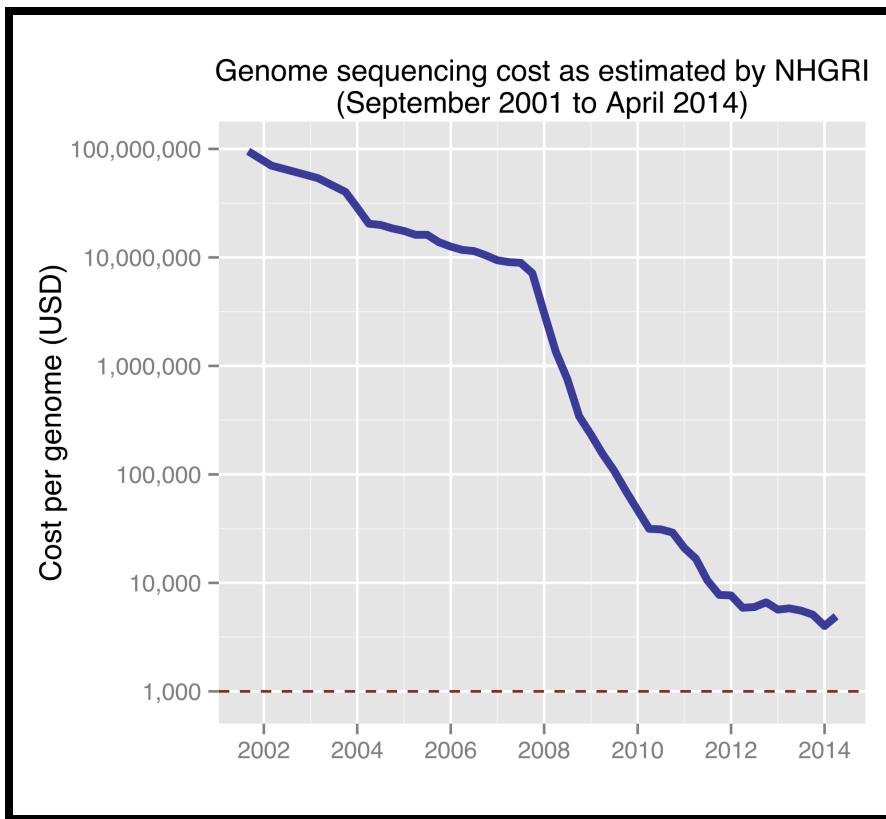
2. From relational to non-relational databases

3. Scalling MySQL : @LesFurets.com

4. PostgreSQL + TP

Why !

ADN sequencing



Processing the sequenced data has become the bottleneck!

– Andy Petrella

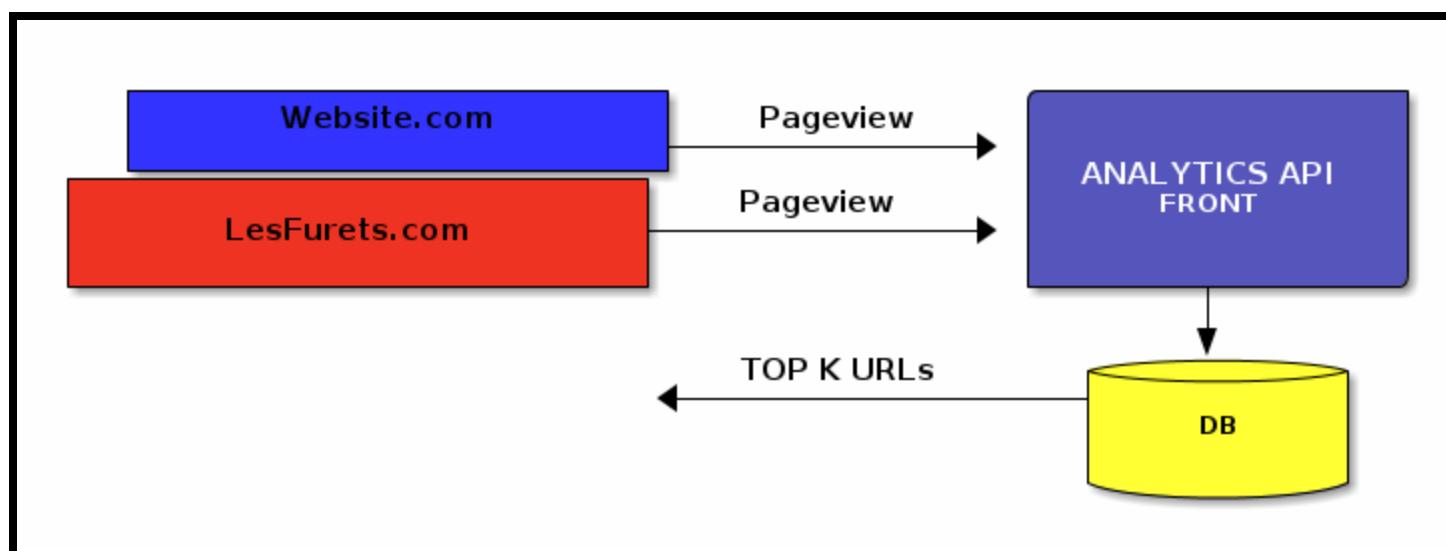
Lightning fast genomics with Spark Adam and Scala

Building a simple web analytics application

Bigdata by Nathan Marz

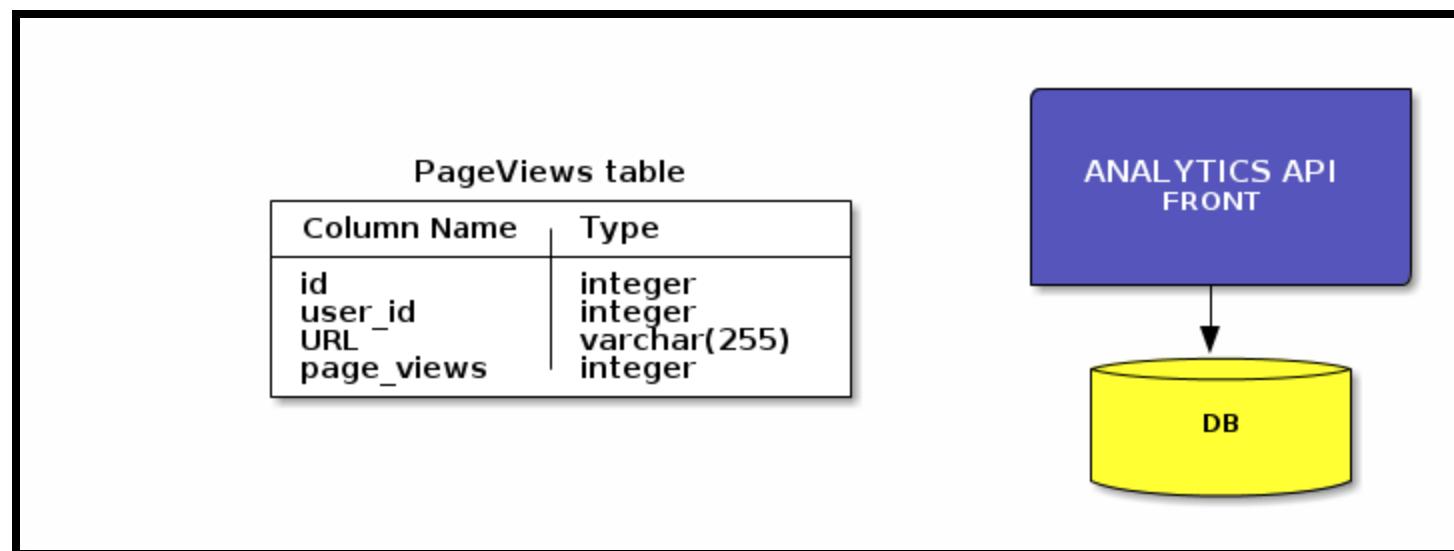
Simple web analytics application

- track the number of pageviews for each URL
- what are the top 100 URLs



Simplest architecture

- track the number of pageviews for each URL
- what are the top 100 URLs



Queries

- insert a pageview
- update pageviews
- top 100 URLs for a client

Insert pageviews

PageViews table	
Column Name	Type
<code>id</code>	<code>integer</code>
<code>user_id</code>	<code>integer</code>
<code>URL</code>	<code>varchar(255)</code>
<code>page_views</code>	<code>integer</code>

```
INSERT INTO PageViews VALUES
    (1,1,"http://www.lesfurets.com/index.html",1);
INSERT INTO PageViews VALUES
    (2,2,"http://Website.com/base.html",1);
INSERT INTO PageViews VALUES
    (3,1,"http://www.lesfurets.com/assurance-auto",1);
```

Update pageviews

PageViews table	
Column Name	Type
<code>id</code>	<code>integer</code>
<code>user_id</code>	<code>integer</code>
<code>URL</code>	<code>varchar(255)</code>
<code>page_views</code>	<code>integer</code>

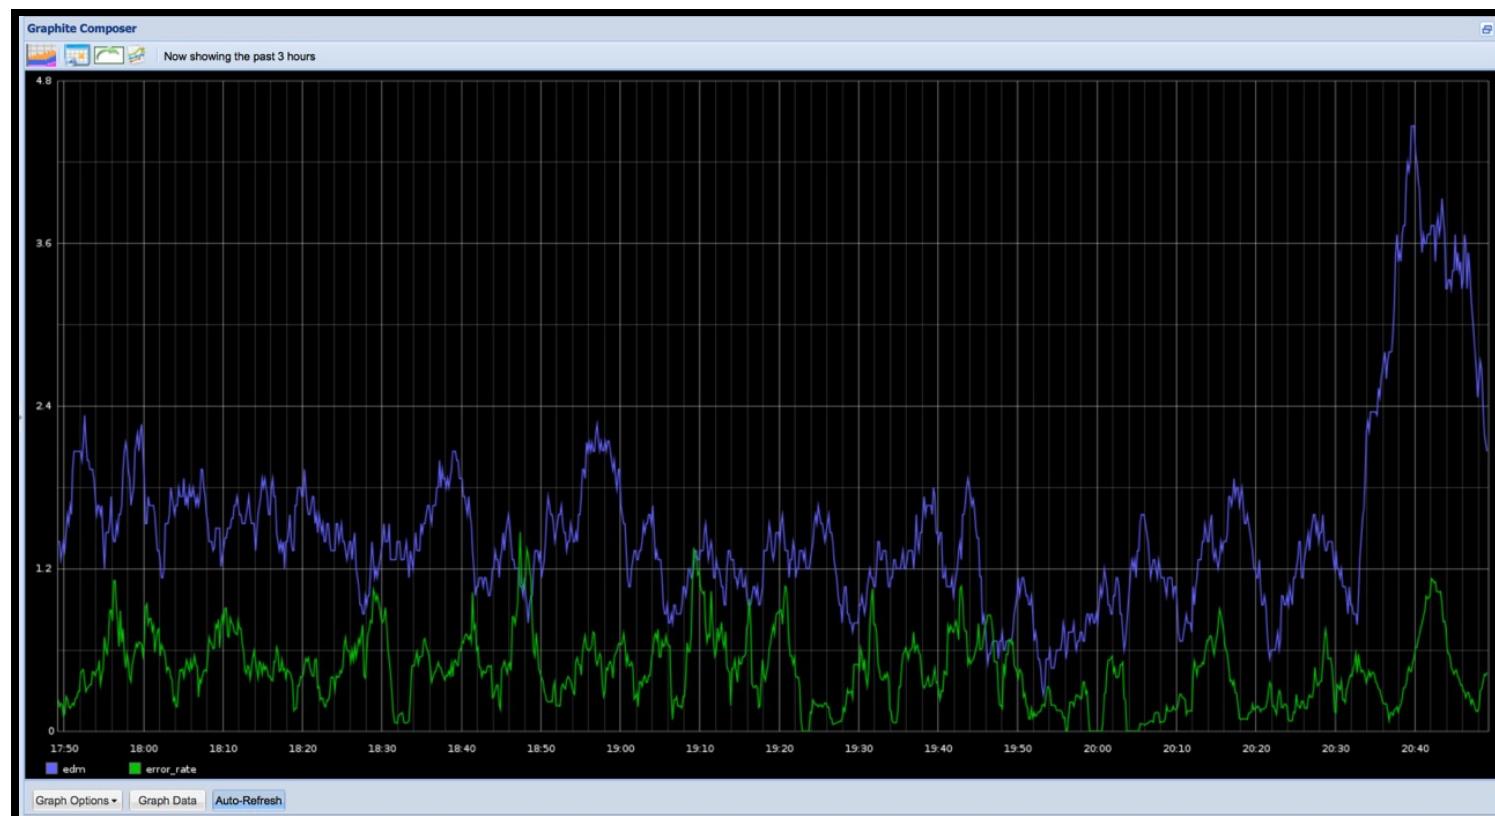
```
UPDATE PageViews SET page_views = page_views+1  
WHERE user_id="1" AND URL="http://www.lesfurets.com/index.html";
```

Top 100 URLs for a client

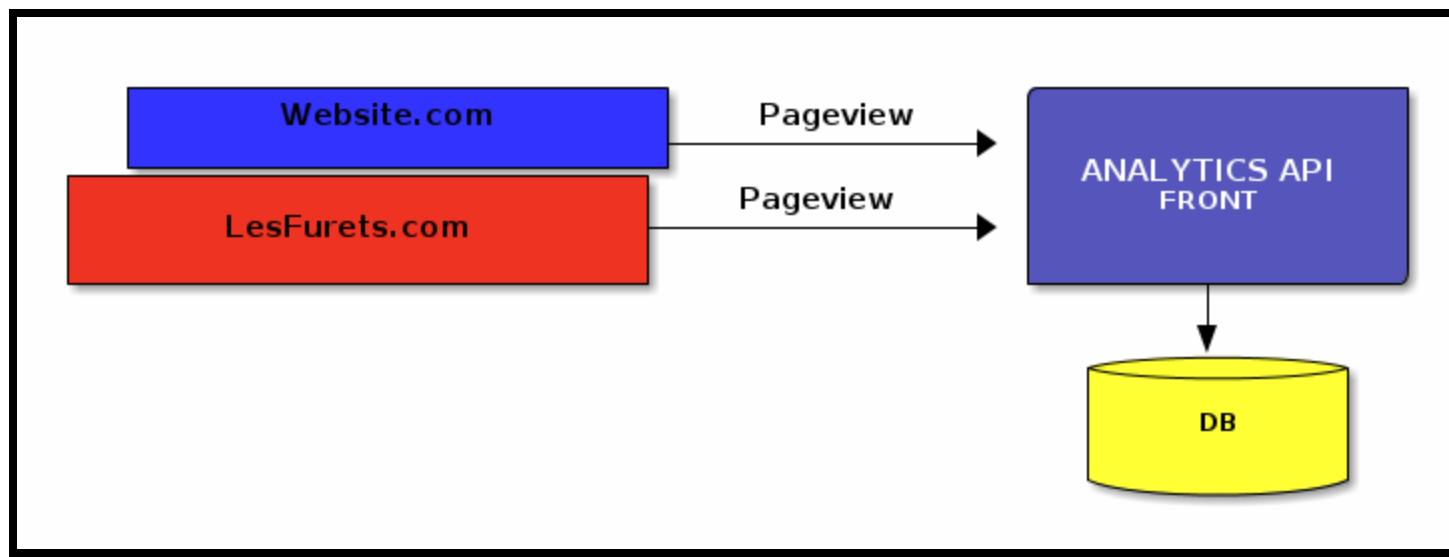
PageViews table	
Column Name	Type
<code>id</code>	<code>integer</code>
<code>user_id</code>	<code>integer</code>
<code>URL</code>	<code>varchar(255)</code>
<code>page_views</code>	<code>integer</code>

```
SELECT URL,page_views FROM PageViews  
WHERE user_id='1'  
ORDER BY page_views DESC LIMIT 100
```

Production load

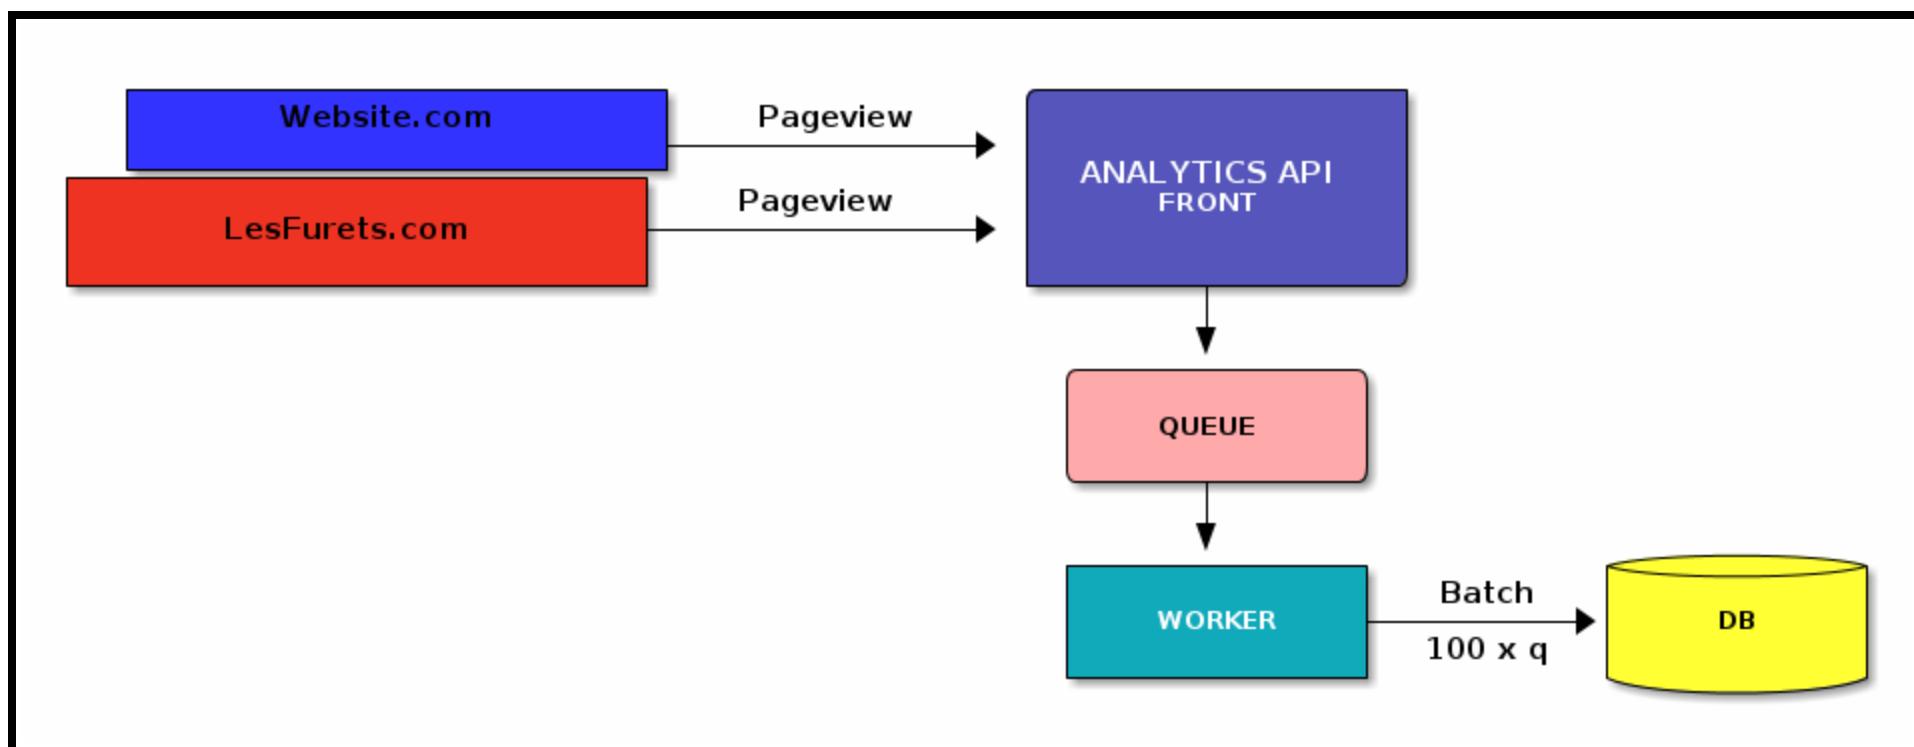


Timeouts



Timeout error on inserting/updating the database

Fix#1 Queuing



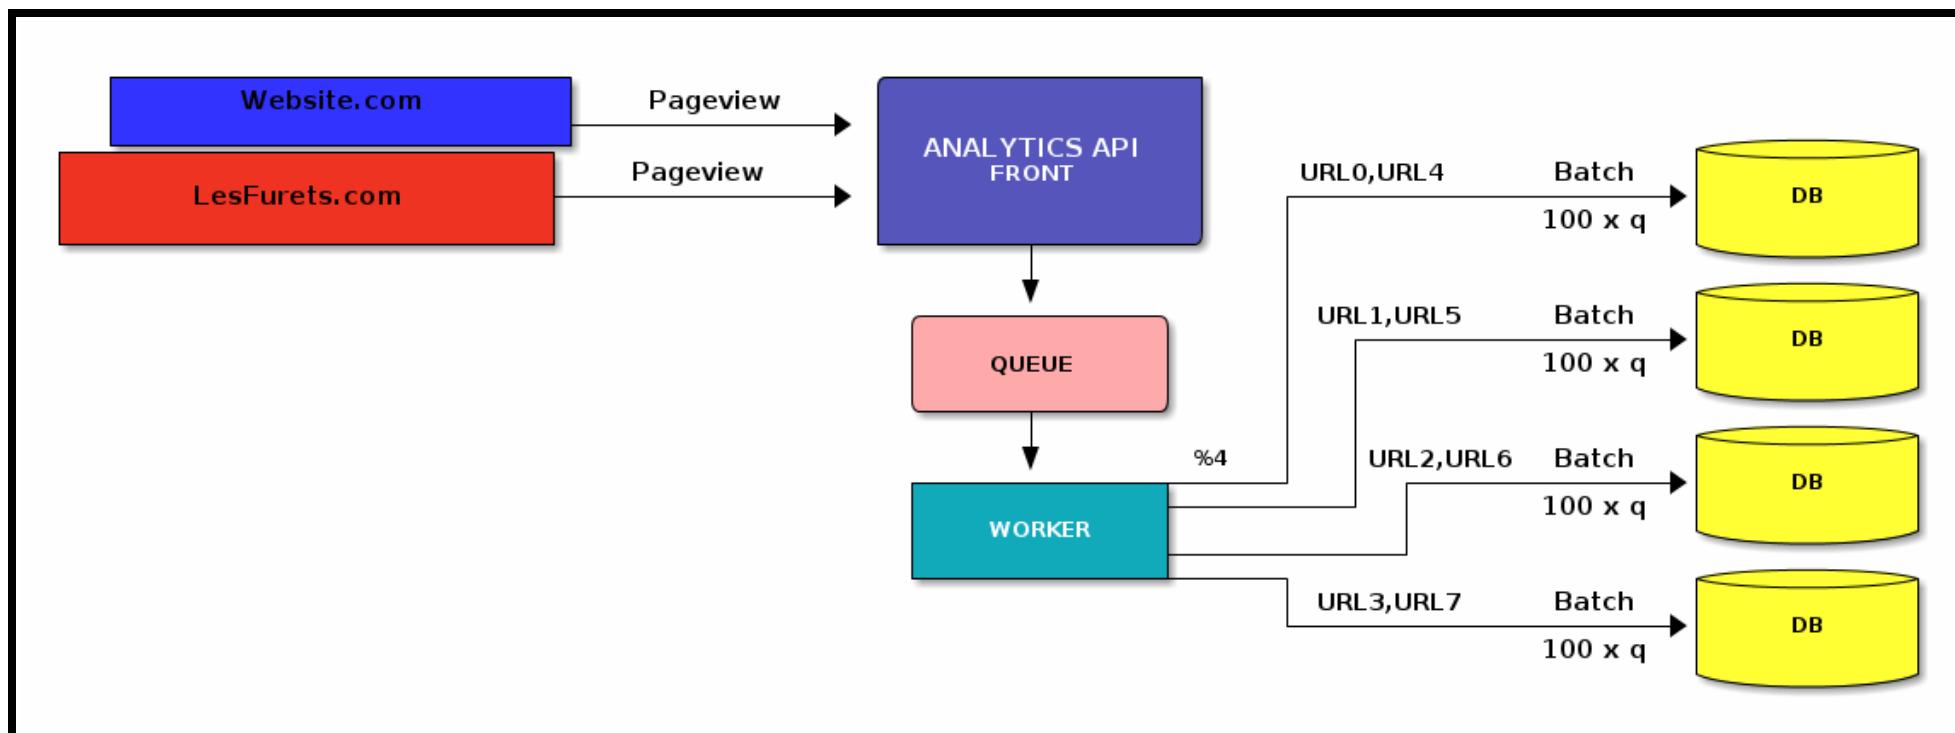
Fix#1 Queuing

- Modify the application ⇒ batch 100 queries
- Latency / queue size
- handle DB/queue failures ⇒ persistent queuing with event logging
- **cannot accomodate high load**

Fix#2 Sharding / Vertical partitioning

- Spread the load
 - use multiple database servers
 - spread the PageView table across the servers
 - mapping keys to shards using a hash function

Fix#2 Sharding



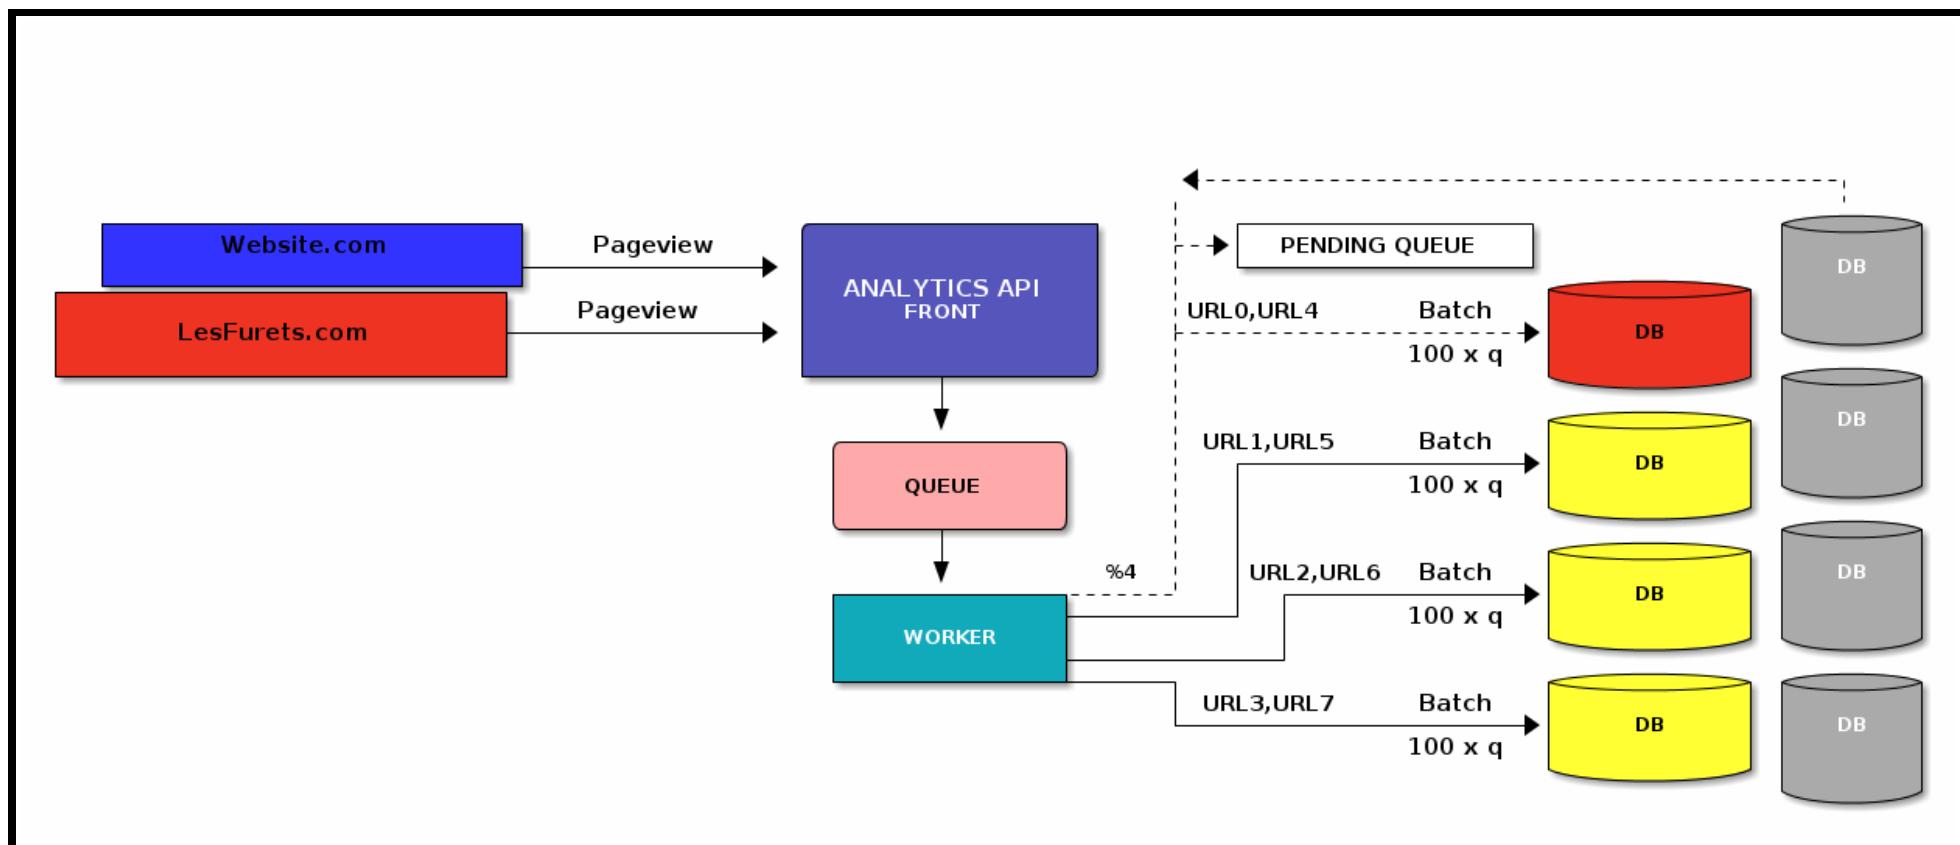
Fix#2 Sharding

- Impacts
 - distribute the keys to the new servers
 - write to the "right" DB instance
 - aggregate data from all the shards !
- **Sharding more and more**
 - new shards to follow the load
 - repeat the last steps

Fix#2 Sharding

- **Server failures** are more likely
 - **WRITES**: use a pending queue flushed less frequently
 - **READS**: a portion of the data is unavailable
 - ⇒ replication

Fix#3 Replication



Human failures

Distribution hash function = %3

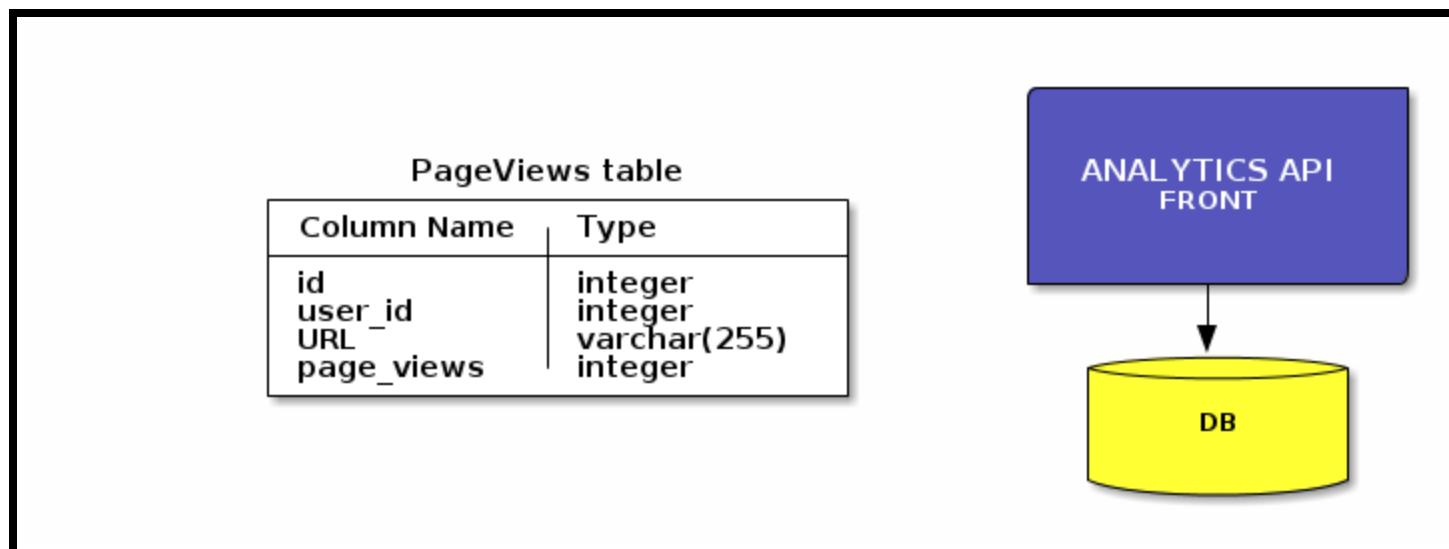
- **Data written to the wrong shards**
 - redistribute data to the right shard
 - while still accepting queries ?!

Human failures

```
UPDATE PageView SET page_views = page_views + 2  
WHERE user_id='42' AND URL='myurl';
```

Increments the number of pageviews + 2

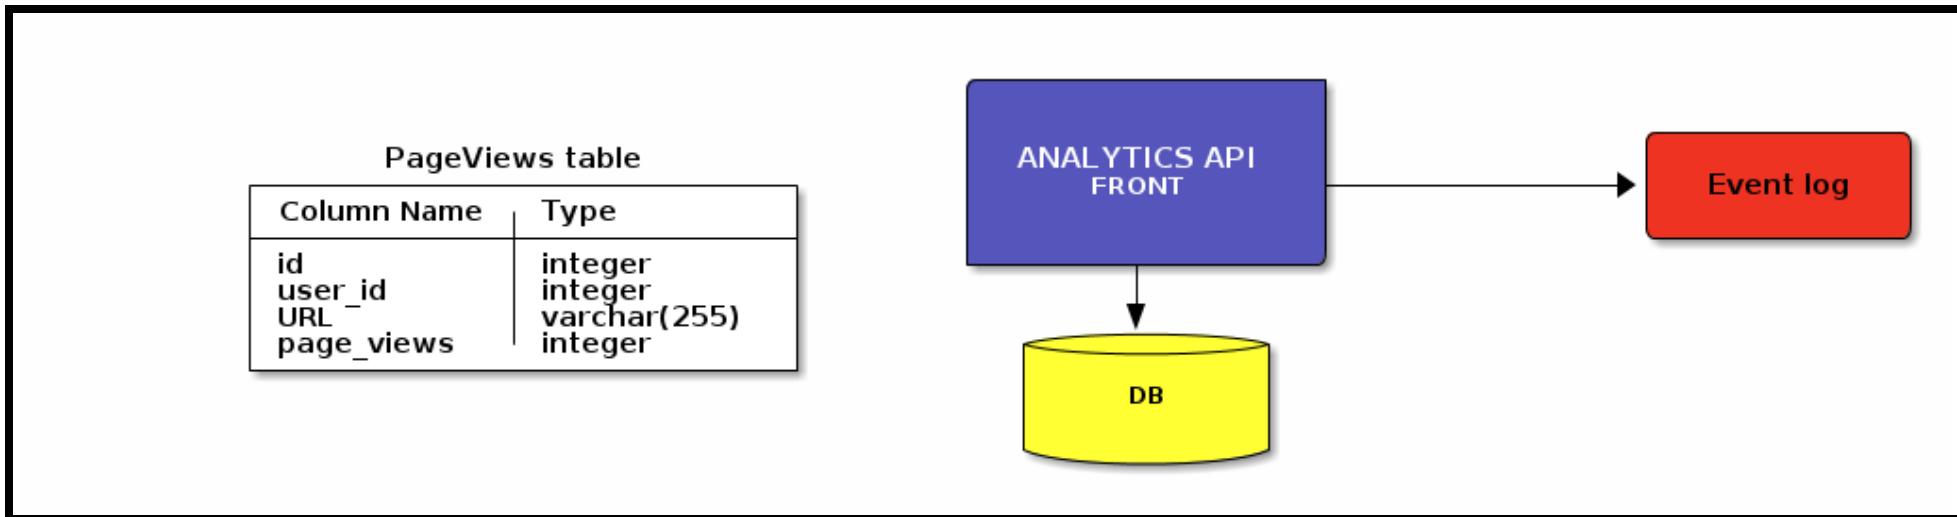
Human failures



```
UPDATE PageView SET page_views = page_views + 2  
WHERE user_id='42' AND URL='myurl';
```

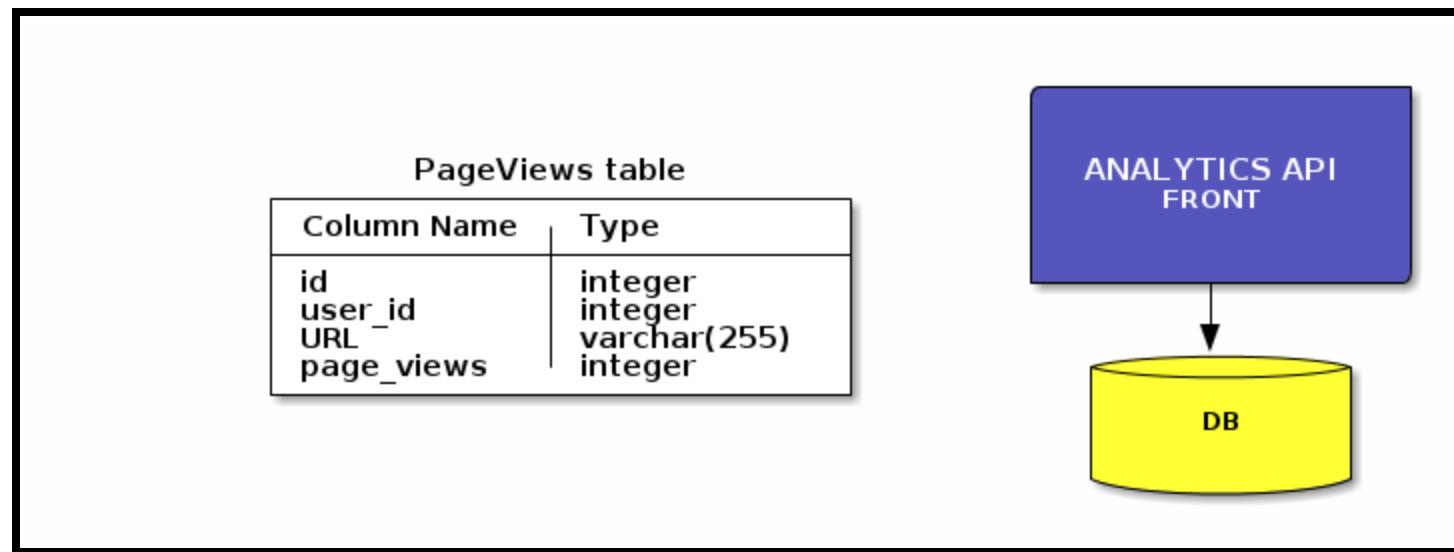
Human failures

- event logging



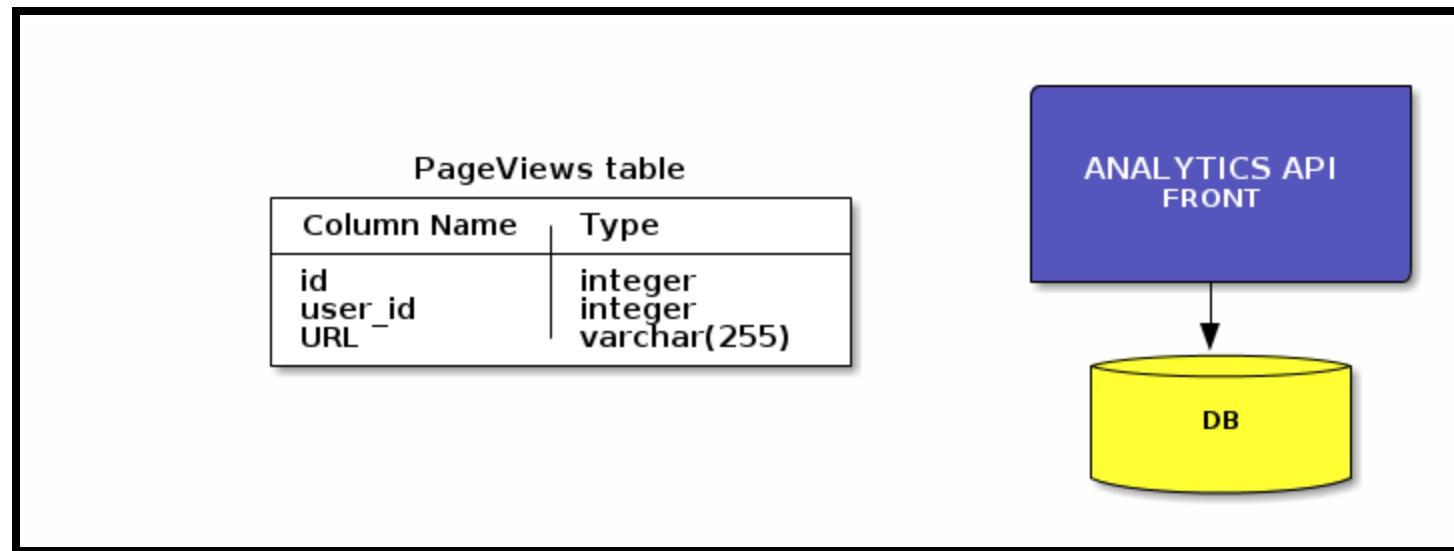
Human failures

- Incremental data model



Human failures

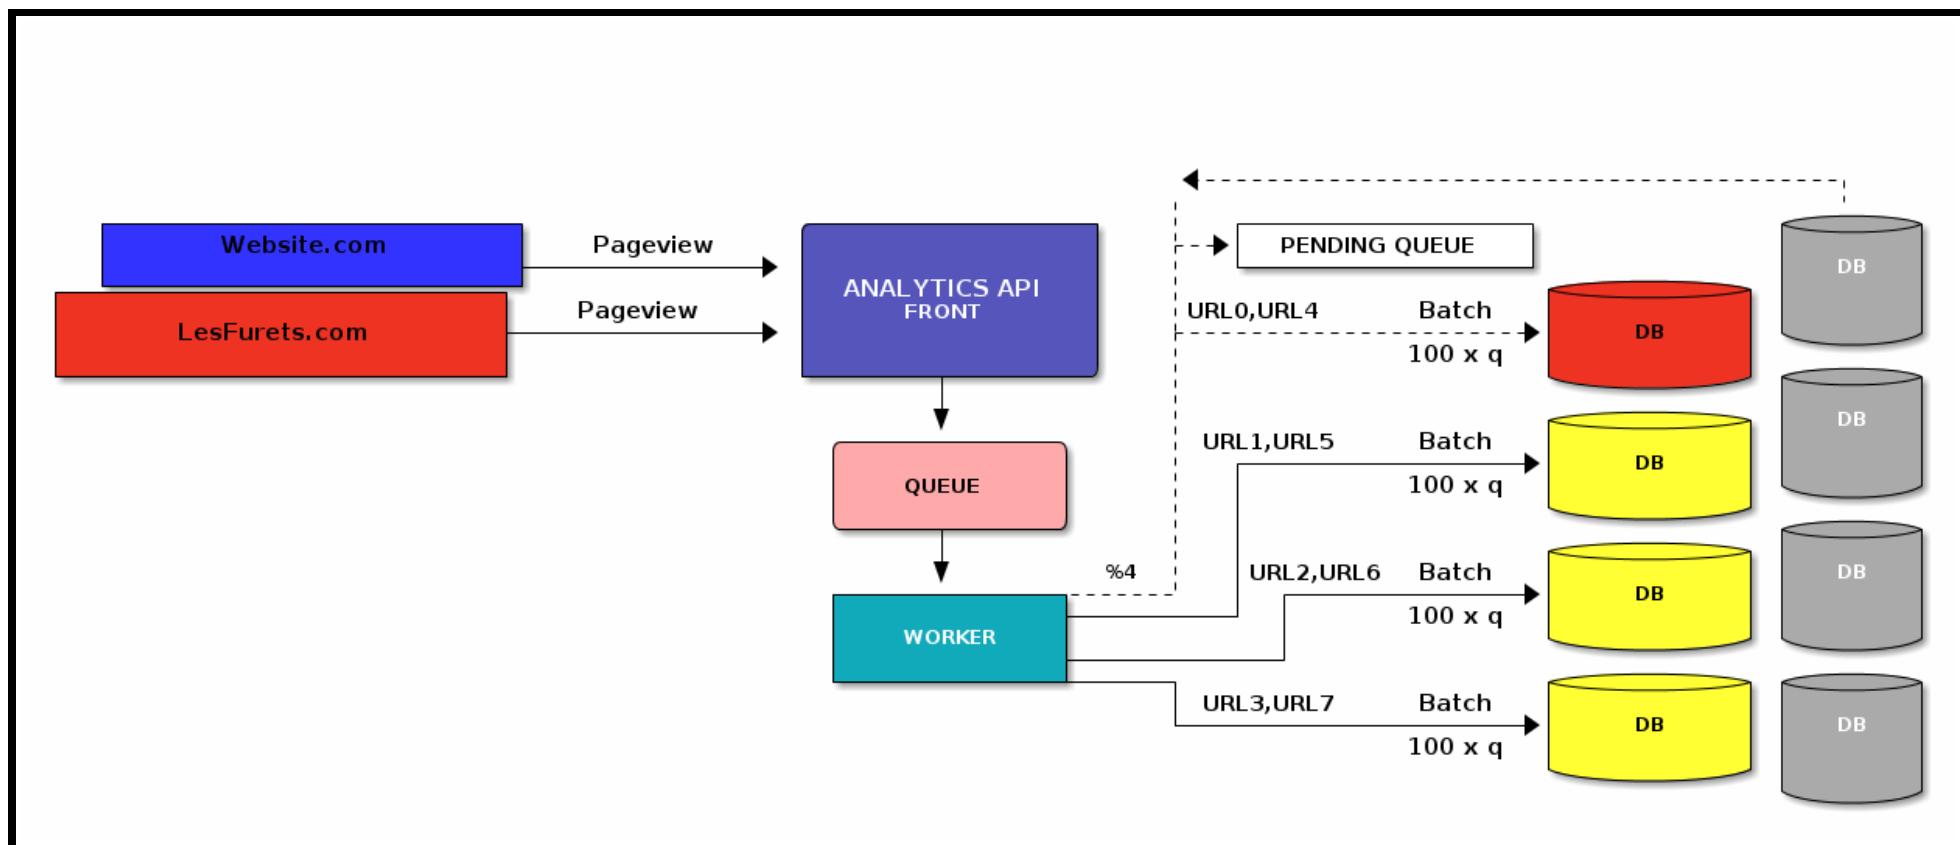
- Incremental data model \Rightarrow immutable data model



Human failures

- incremental data model \Rightarrow **data corruption**
- hard to correct

What went wrong?



- do I build new features for customers?
- or just dealing with reading/writing the data?

What went wrong?

- A single server cannot take the load \Rightarrow solution / complexity
 1. *distributed storage*
 2. querying distributed data
 3. built a data model that is not resilient

Wishlist 1 : Storage

- *Better* storage:*
 - easy to add/remove nodes (**scalling**)
 - transparent data distribution (**auto-sharding**)
 - handle failures (**auto-replication**)

⇒ **Distributed databases:** *Redis, Cassandra, MongoDB, ...*

Wishlist 2 : Queries

- *General purpose (distributed) computing:*
 - distributed queries + **parallel processing**

⇒ **Distributed data processing engines : MapReduce, Spark**

Wishlist 3 : Data model

- *We want a resiliant data model:*
 - human error is **unavoidable**
 - an **incremental data model** is not resiliant
- ⇒ **Immutability**

Plan

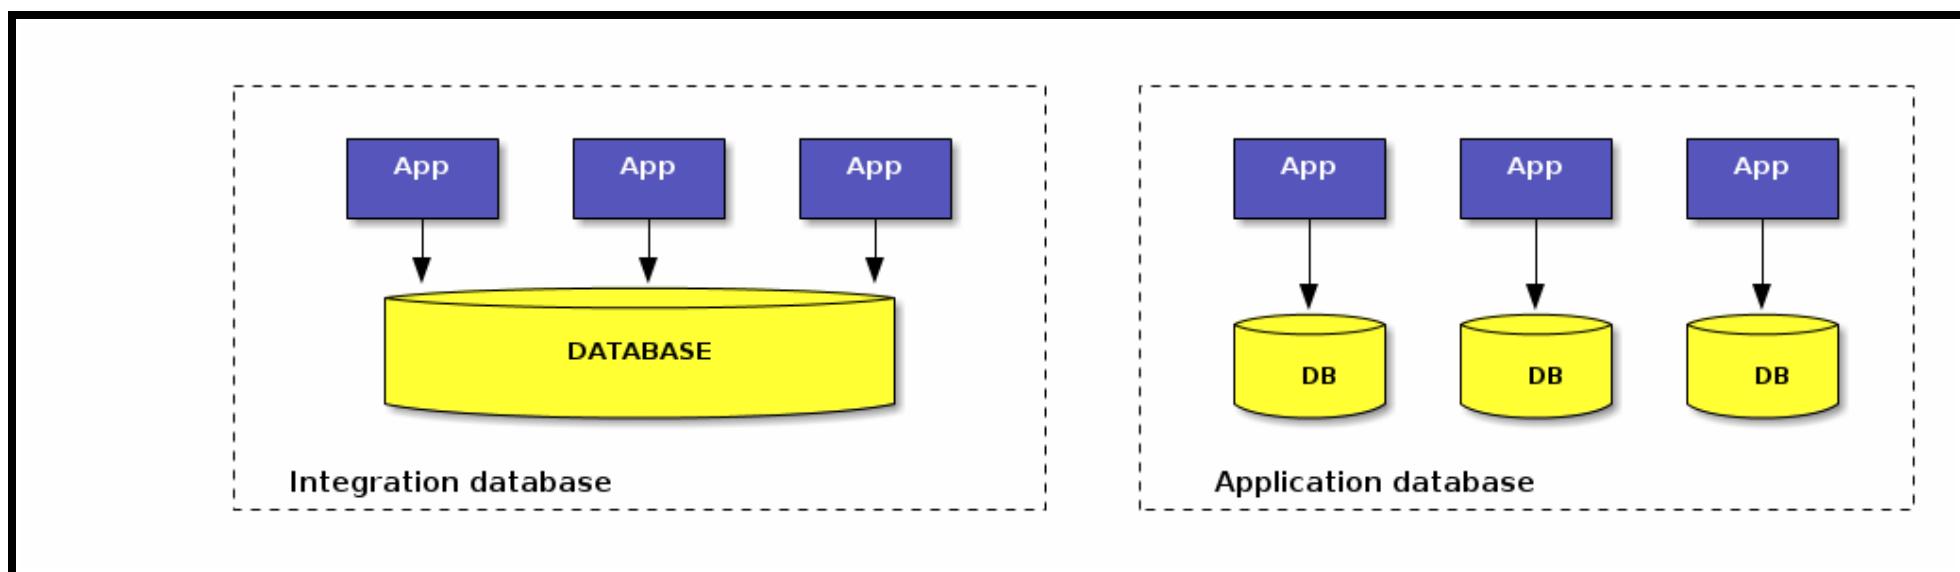
1. Intro
2. From SQL to NoSQL
 1. Scalling a simple application
 - 2. From relational to non-relational databases**
3. Scalling MySQL : @LesFurets.com
4. PostgreSQL + TP

From relational to non-relational databases

- RDBMS (Relational database management systems)
- RDBMS/Memory : model missmatch and scaling concerns
- RDBMS : modeling complex data

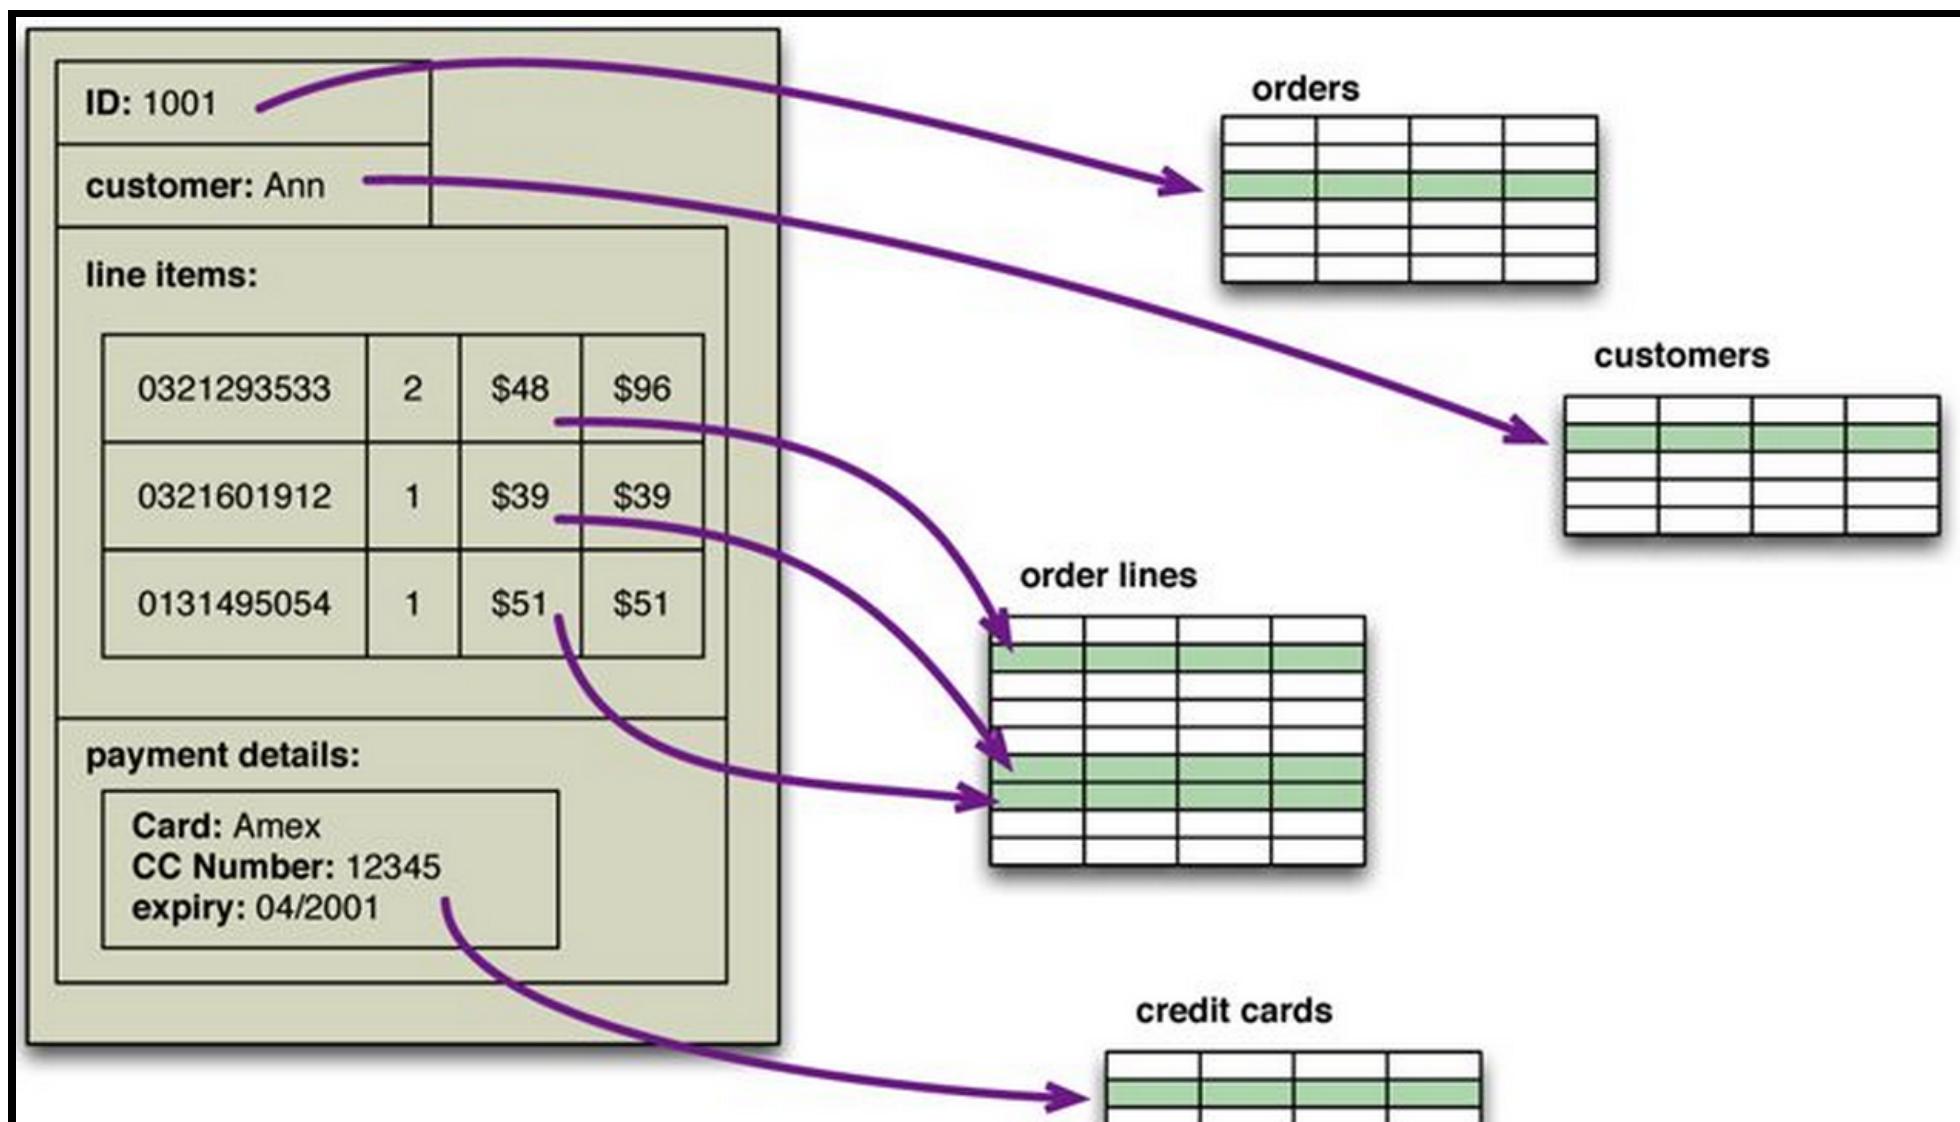
RDBMS: the good parts

- simple model with sound mathematical properties (ACID)
- 45+ years of research (Codd '70, System R '74, Oracle '79)
- industry patterns : integration database / shared database pattern



Model missmatch (1)

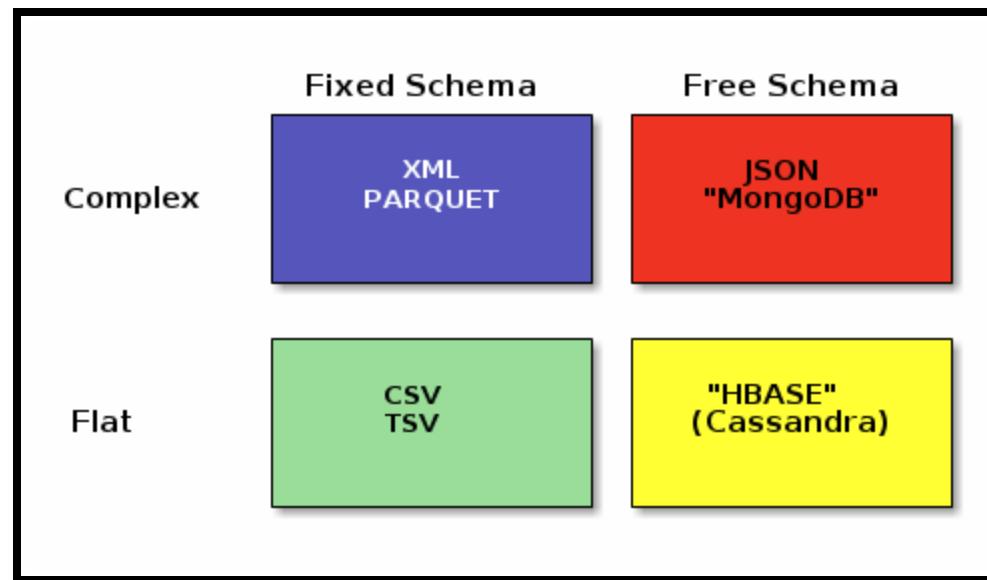
- Relational model: relations / tuples + normalisation (avoid duplication)
- Memory: rich data structures !



Model missmatch (2)

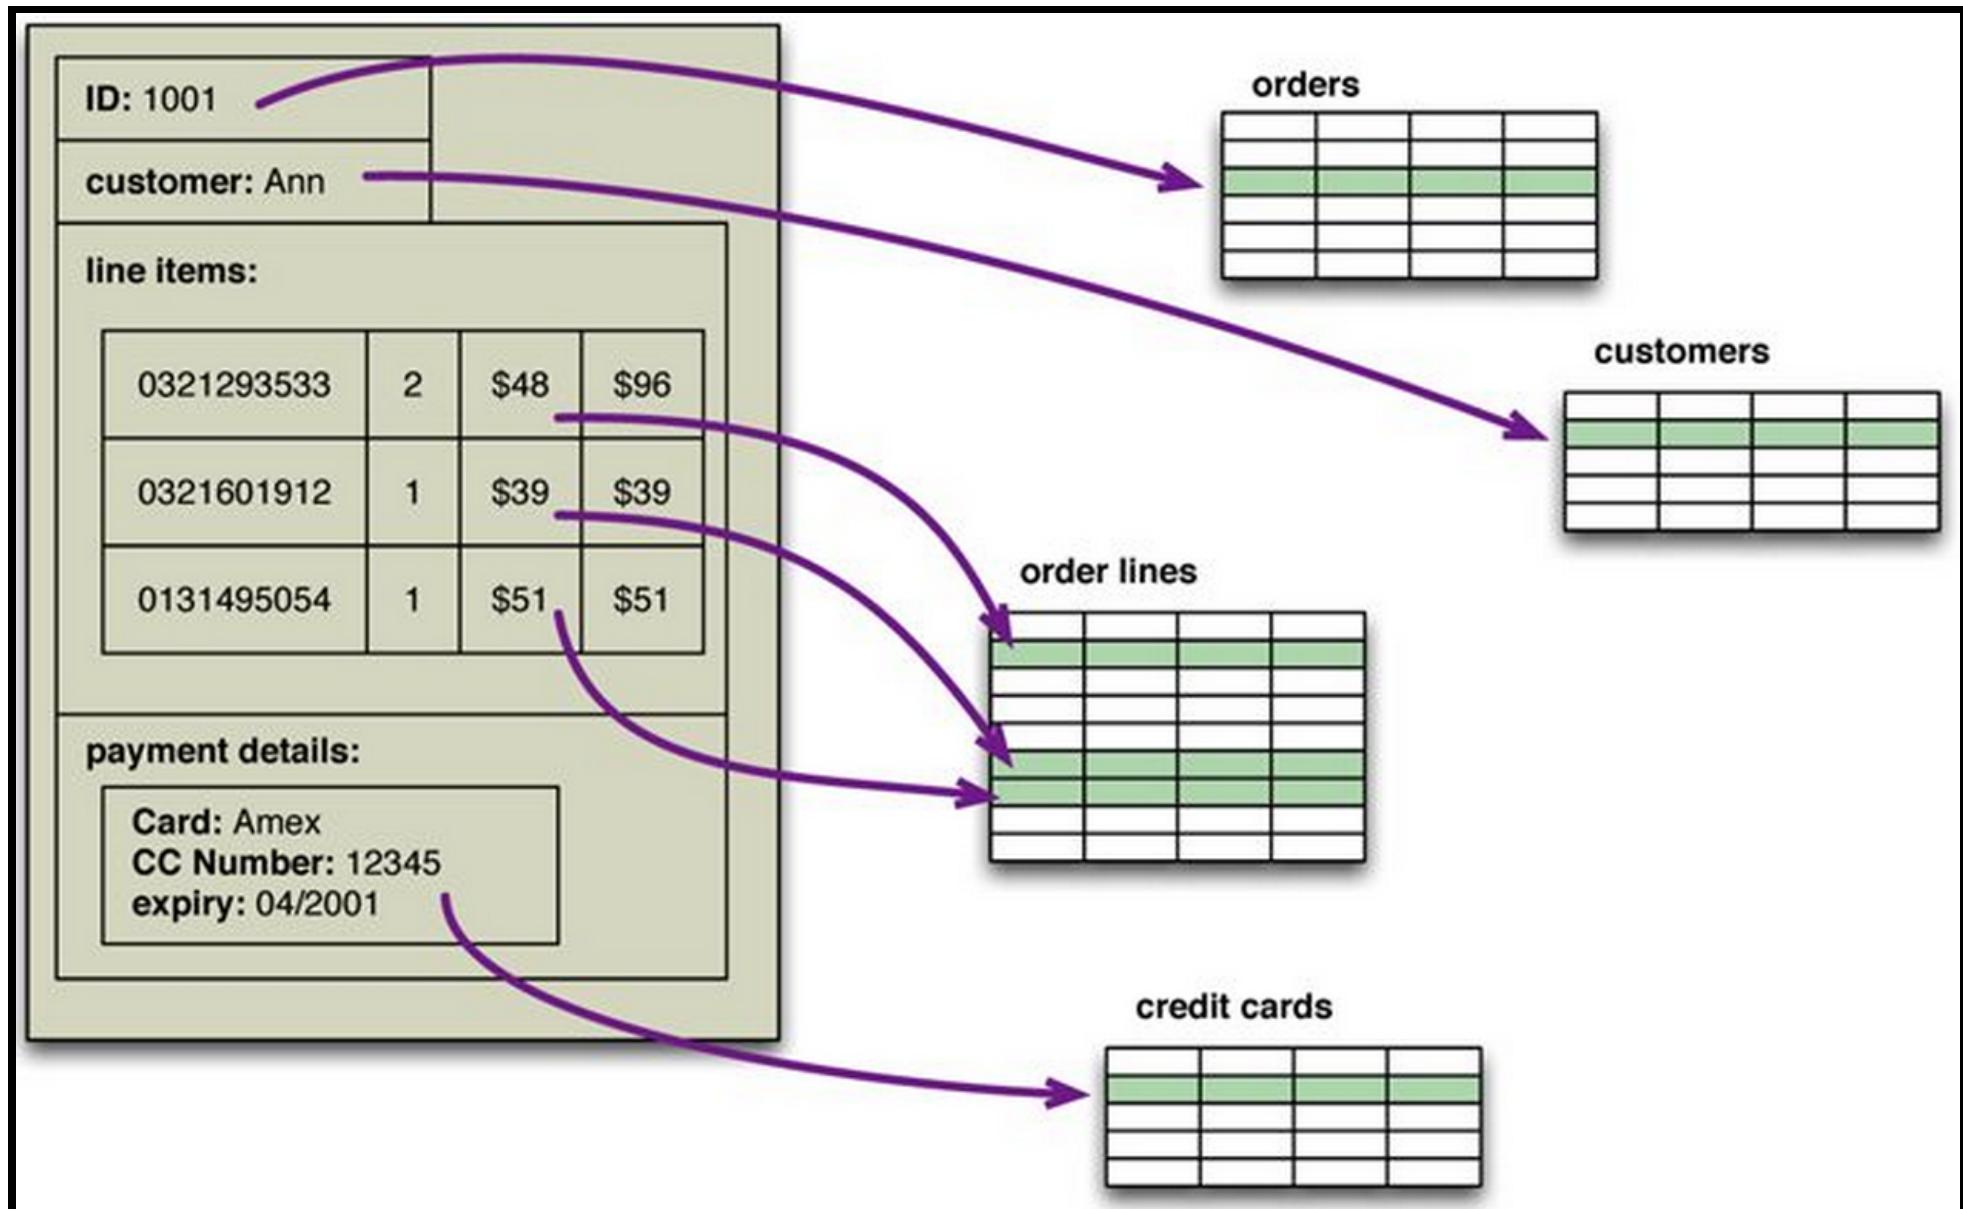
- Memory - object graphs
 - complex mapping from object to relations → ORMs
 - leads to performance issues
 - many rows/tables/**JOINS**
- Schema evolution:
 - adding an attribute → adding a whole column
 - expensive locks → application downtime

Modeling complex data (1)



Modeling complex data = model missmatch

- **Relational model:** relations / tuples + normalisation (avoid duplication)
- **Semi-structured data** hard to model (eg. XML, JSON):
 - blob vs name/value



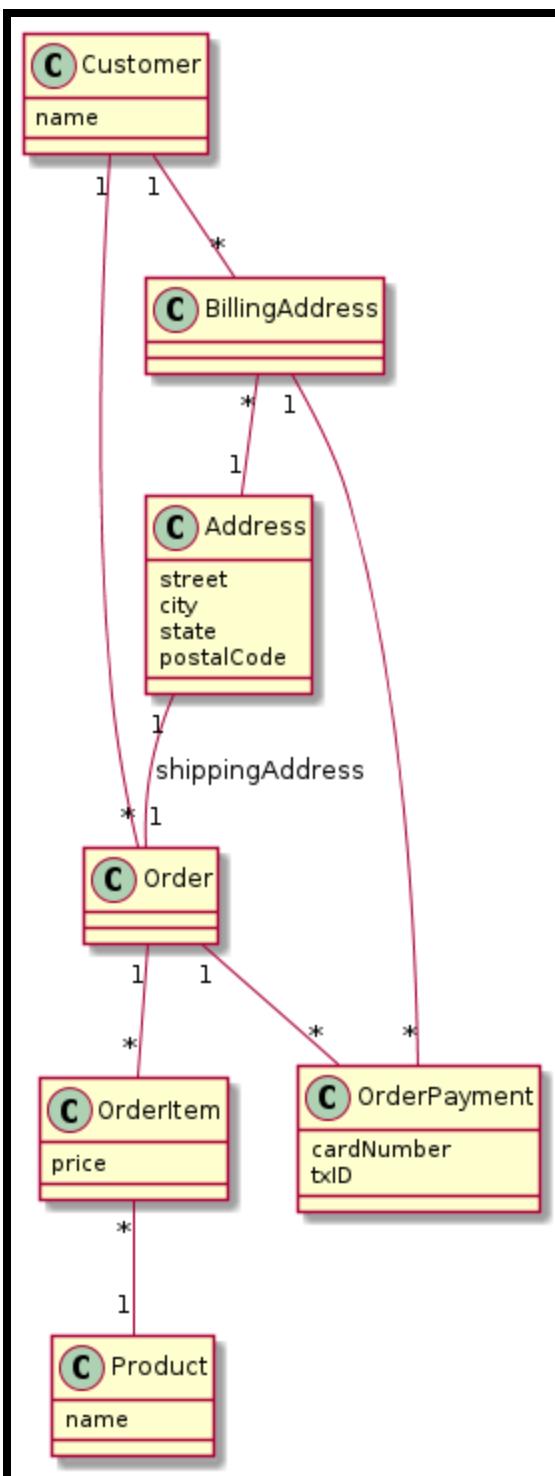
RDBMs: scaling

- **Scaling is expensive:**
 - scaling writes → locking + partitioning (sharding)
 - scaling reads → replication (latency, error recovery)
- **Not a good fit for RDBMs**
 - clustered databases ⇒ shared disk paradigm / cluster aware filesystem
 - application level sharding and replication

Not only SQL

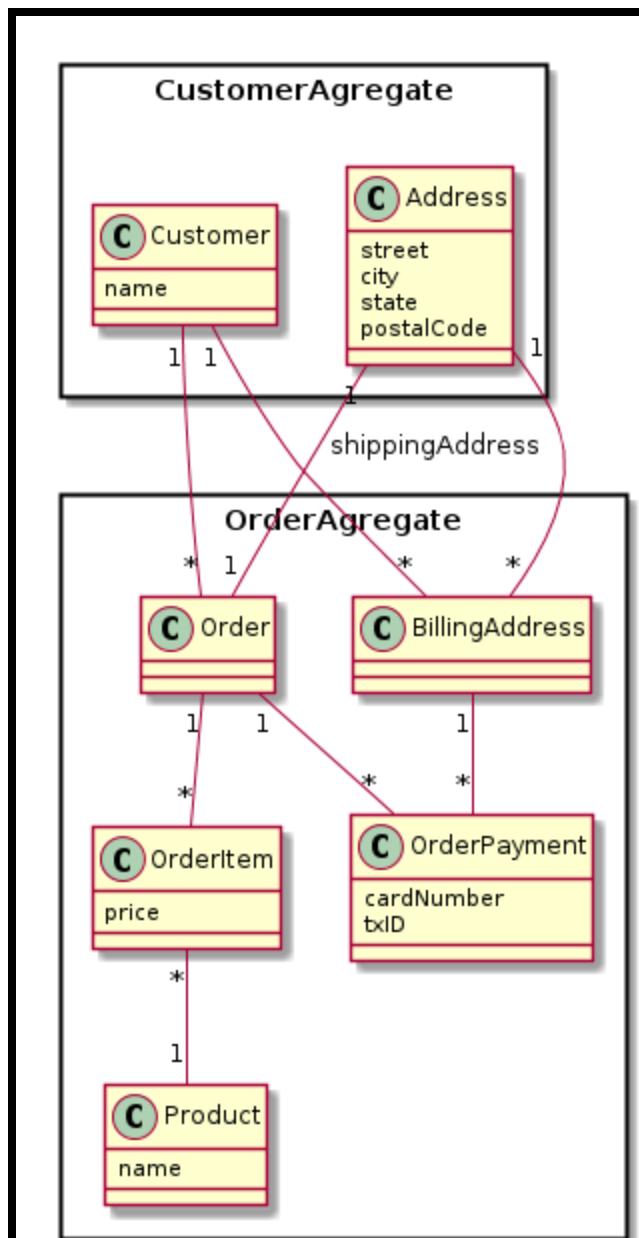
- 2009, Johan Oskarson, #NoSQL meetup for distributed, open-source, non-relational databases
 - not using relational model SQL
 - run on clusters
 - ACID \Rightarrow tunnable consistency
 - fixed schema \Rightarrow flexible schema
 - **polyglot persistance**

SQL models tuples and joins



NoSQL models aggregates

- collection of related objects that should be treated as a unit (consistency / data management)



Modelling SQL vs NoSQL

- *SQL:*
 - **model first**
 - 1 model to govern them all
- *NoSQL:*
 - **query first**
 - 1 data access pattern for each aggregate

NoSQL aggregate types

1. Key-value databases
2. Document-oriented databases
3. Column-oriented databases
4. Graph databases

Key-value databases

- Store and retrieve Blobs based on a primary Key
- Simplest API that matches REST : PUT/GET/DELETE

Model: HashTable ($K \rightarrow V$)

Use case: Session information, user profiles, ...

Document oriented databases

- Stores and retrieves **documents/fragments** (XML, JSON...):
 - self-describing, hierarchical tree data structures
 - maps, collections and scalar values

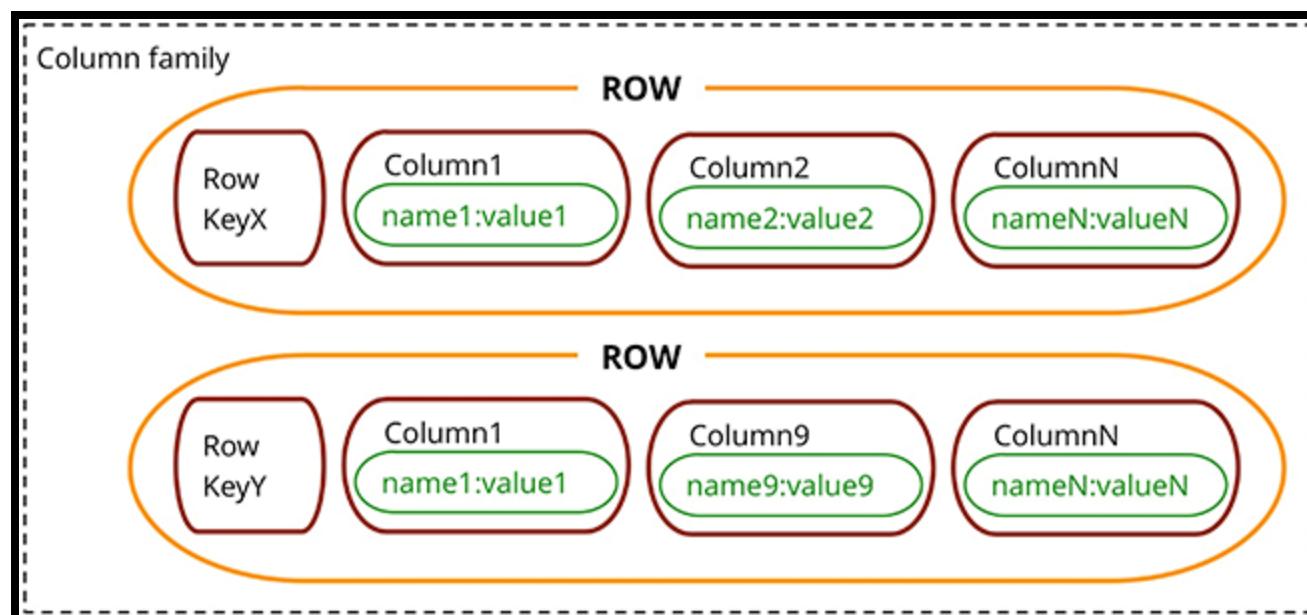
Model: key-value stores where the value is queryable

Use case: CMS, Product catalog, ...

Column oriented databases

- Stores data in **column families** as rows that have many columns associated to a row key
- Free format: two rows can have different columns

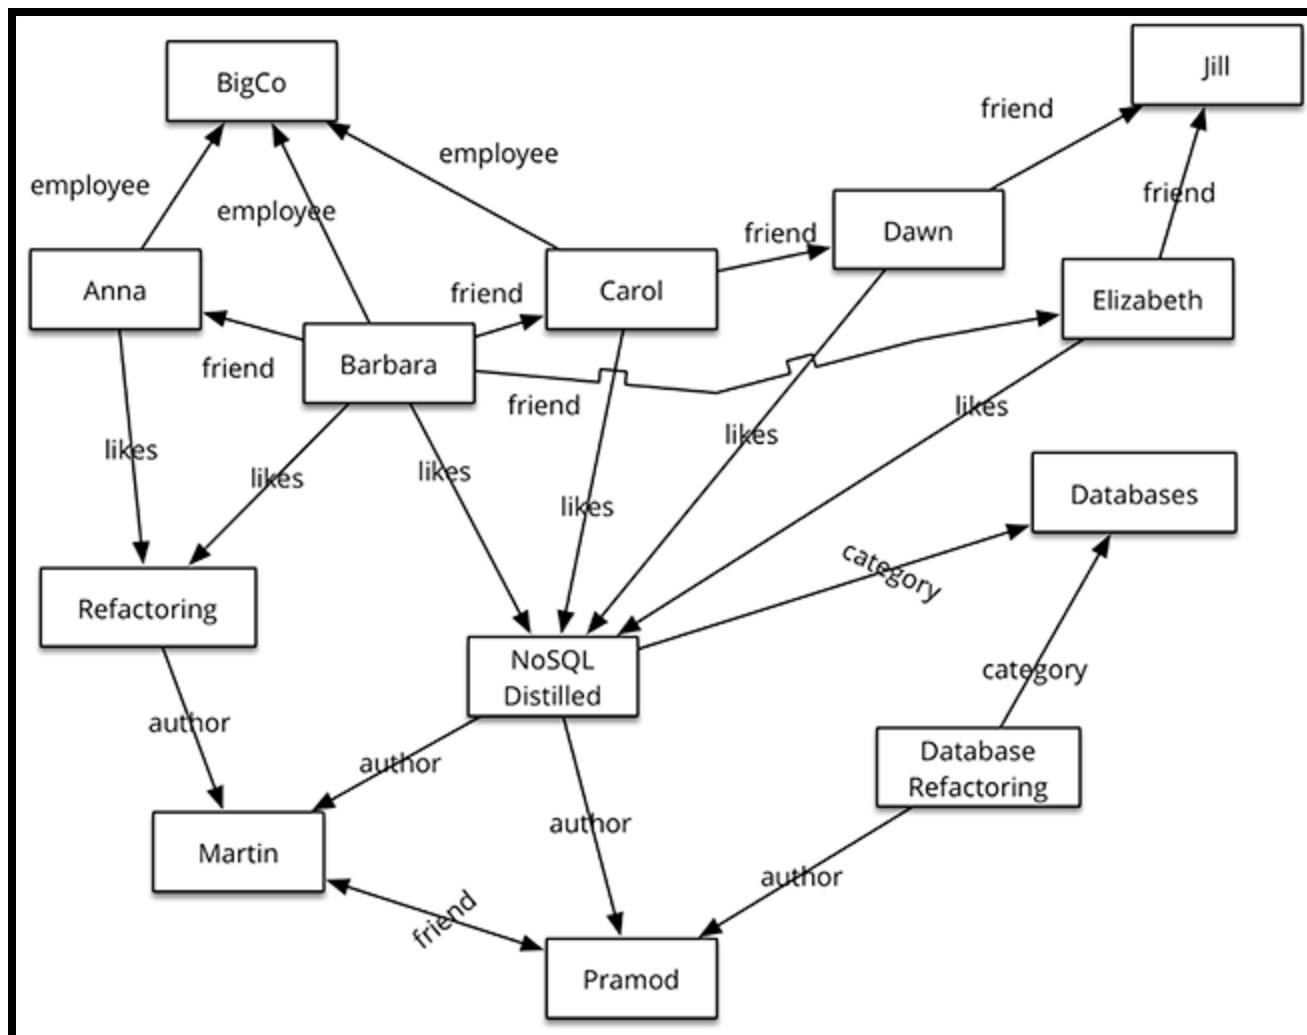
"Map of map model ($\text{rowId} \rightarrow (\text{columnName} \rightarrow \text{columnValue})$)"



Use cases: Time series, event logging, ...

Graph databases

- Stores **entities** and **relations** between entities **Query:** traversal of the graph



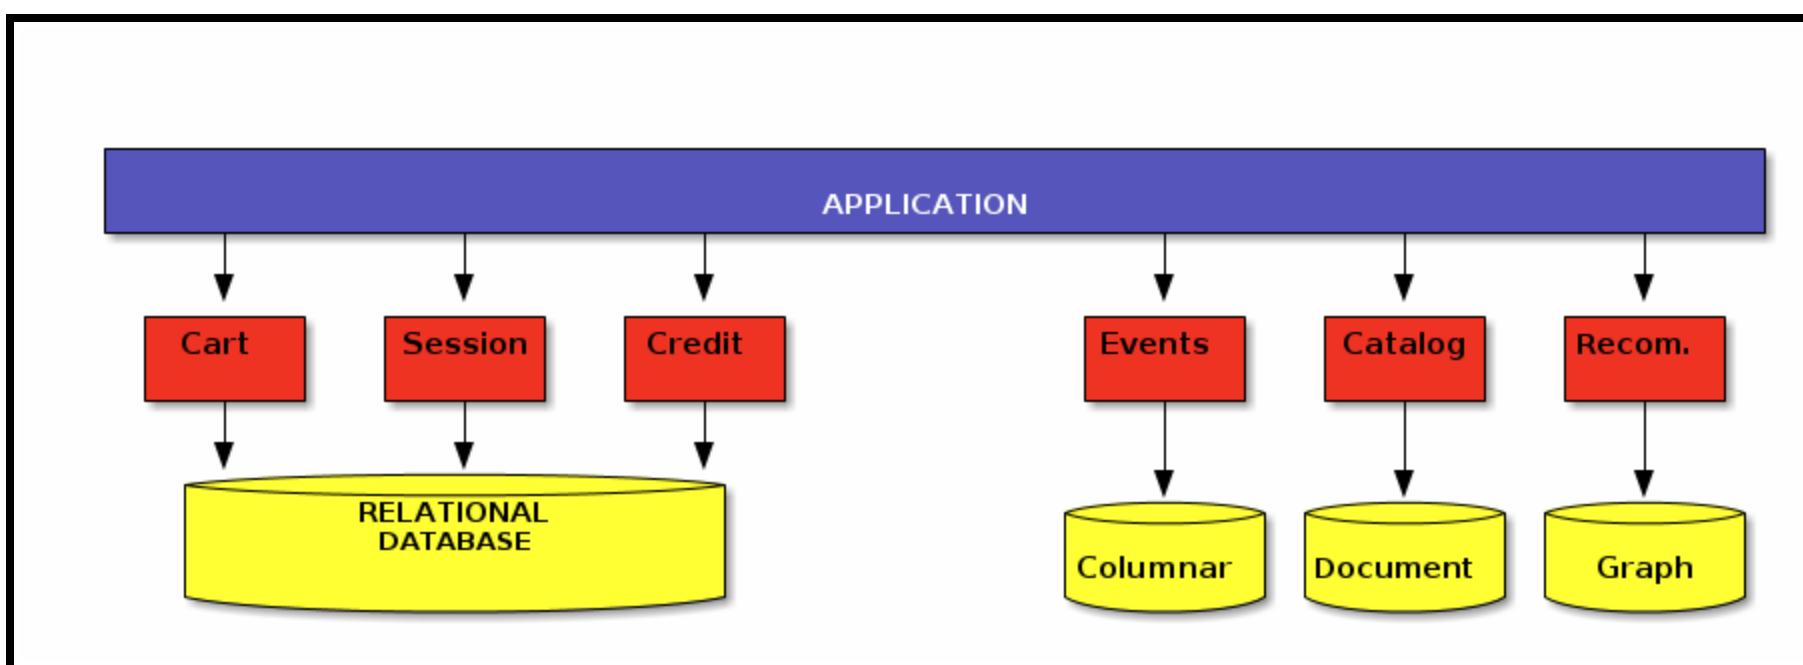
Use cases: Social networks, recommendation engines

NoSQL not a panacea

- Performance:
 - Hadoop: large scale batch computation but **high latency**
 - Cassandra: low latency, fine grain storage but **limited data model**
- Consistency: **ACID** \Rightarrow tunnable consistency (**CAP**)
- Model:
 - Incremental architectures \Rightarrow human failures

Polyglot persistence

- aggregates have different requirements
(availability/consistency/backup)
- Mix and match relational and non-relational storages



Plan

1. Intro
2. From SQL to NoSQL
3. Scalling MySQL : **@LesFurets.com**
4. PostgreSQL + TP

Scalling MySQL

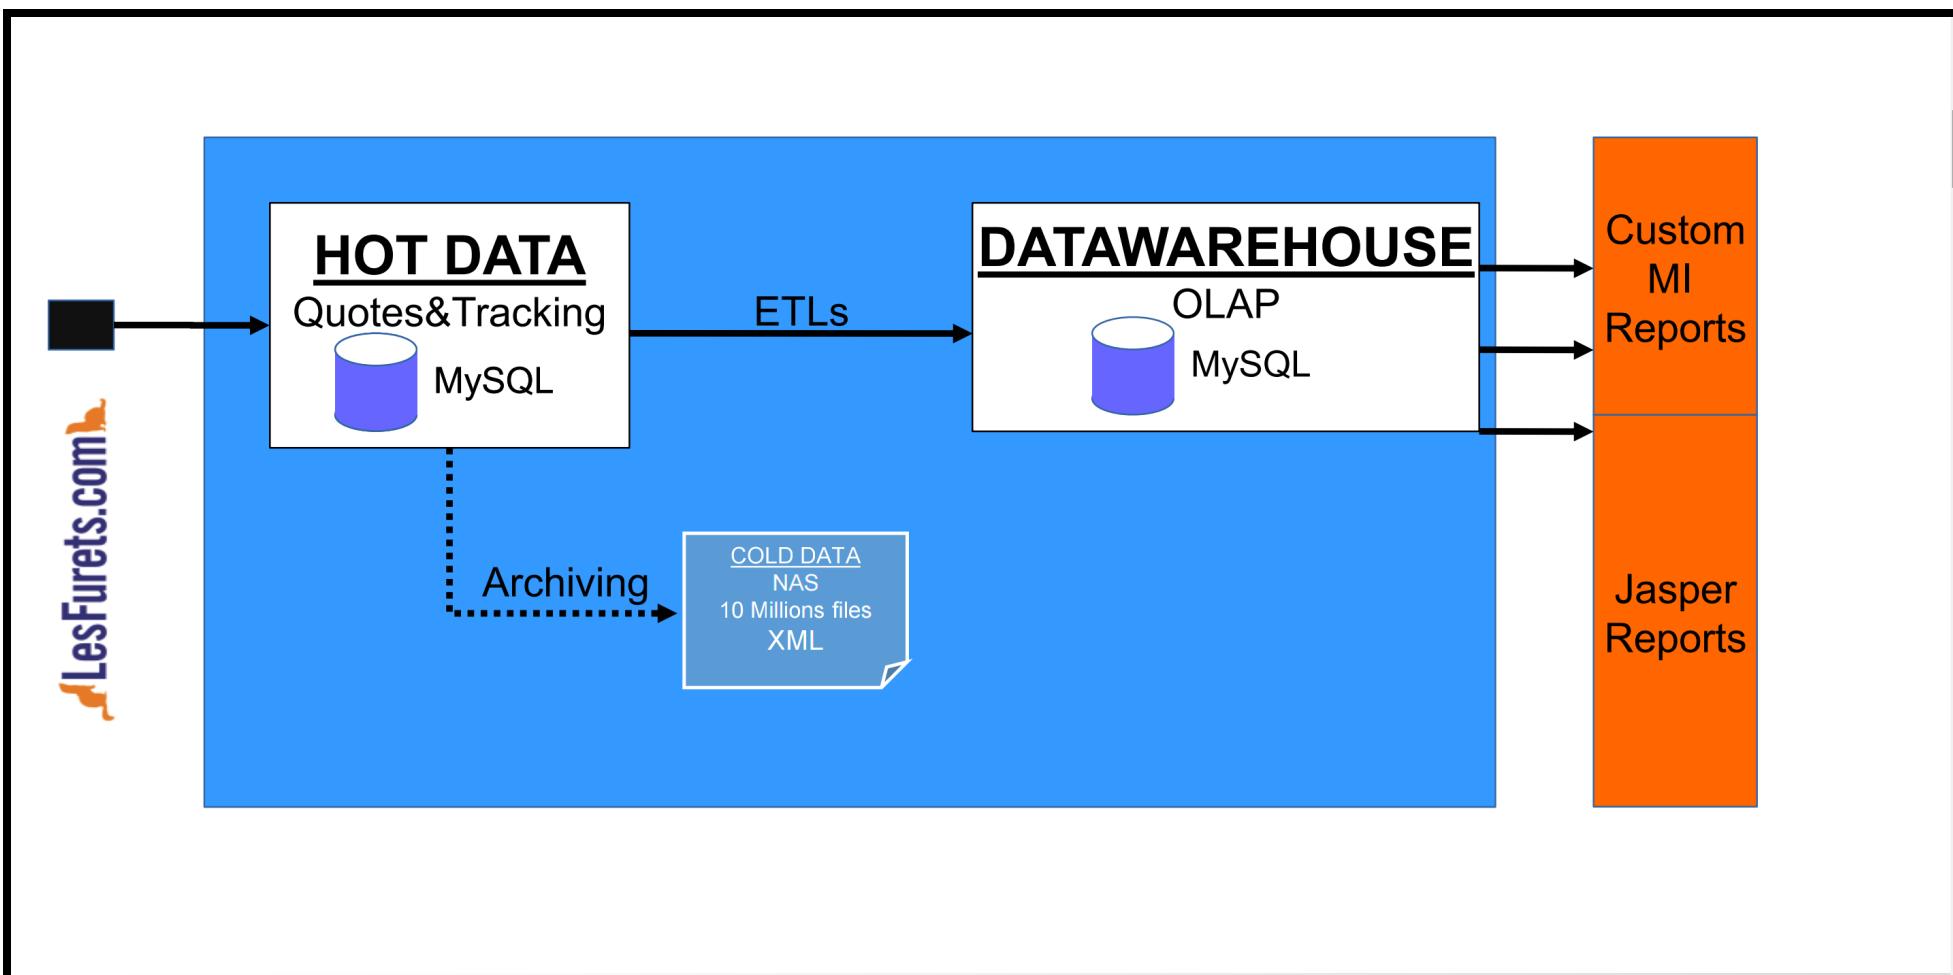
LesFurets.com

- Independent insurance aggregator
- Persistance:
 - replicated MySQL db ⇒ DRBD ⇒ Galera Cluster
 - Cassandra / Spark (migration ongoing)
- BI/reporting
 - Olap (JasperReports)
 - QlikView dashboarding
 - Spark/Zeppelin over MySQL/Cassandra: ad-hoc analyses

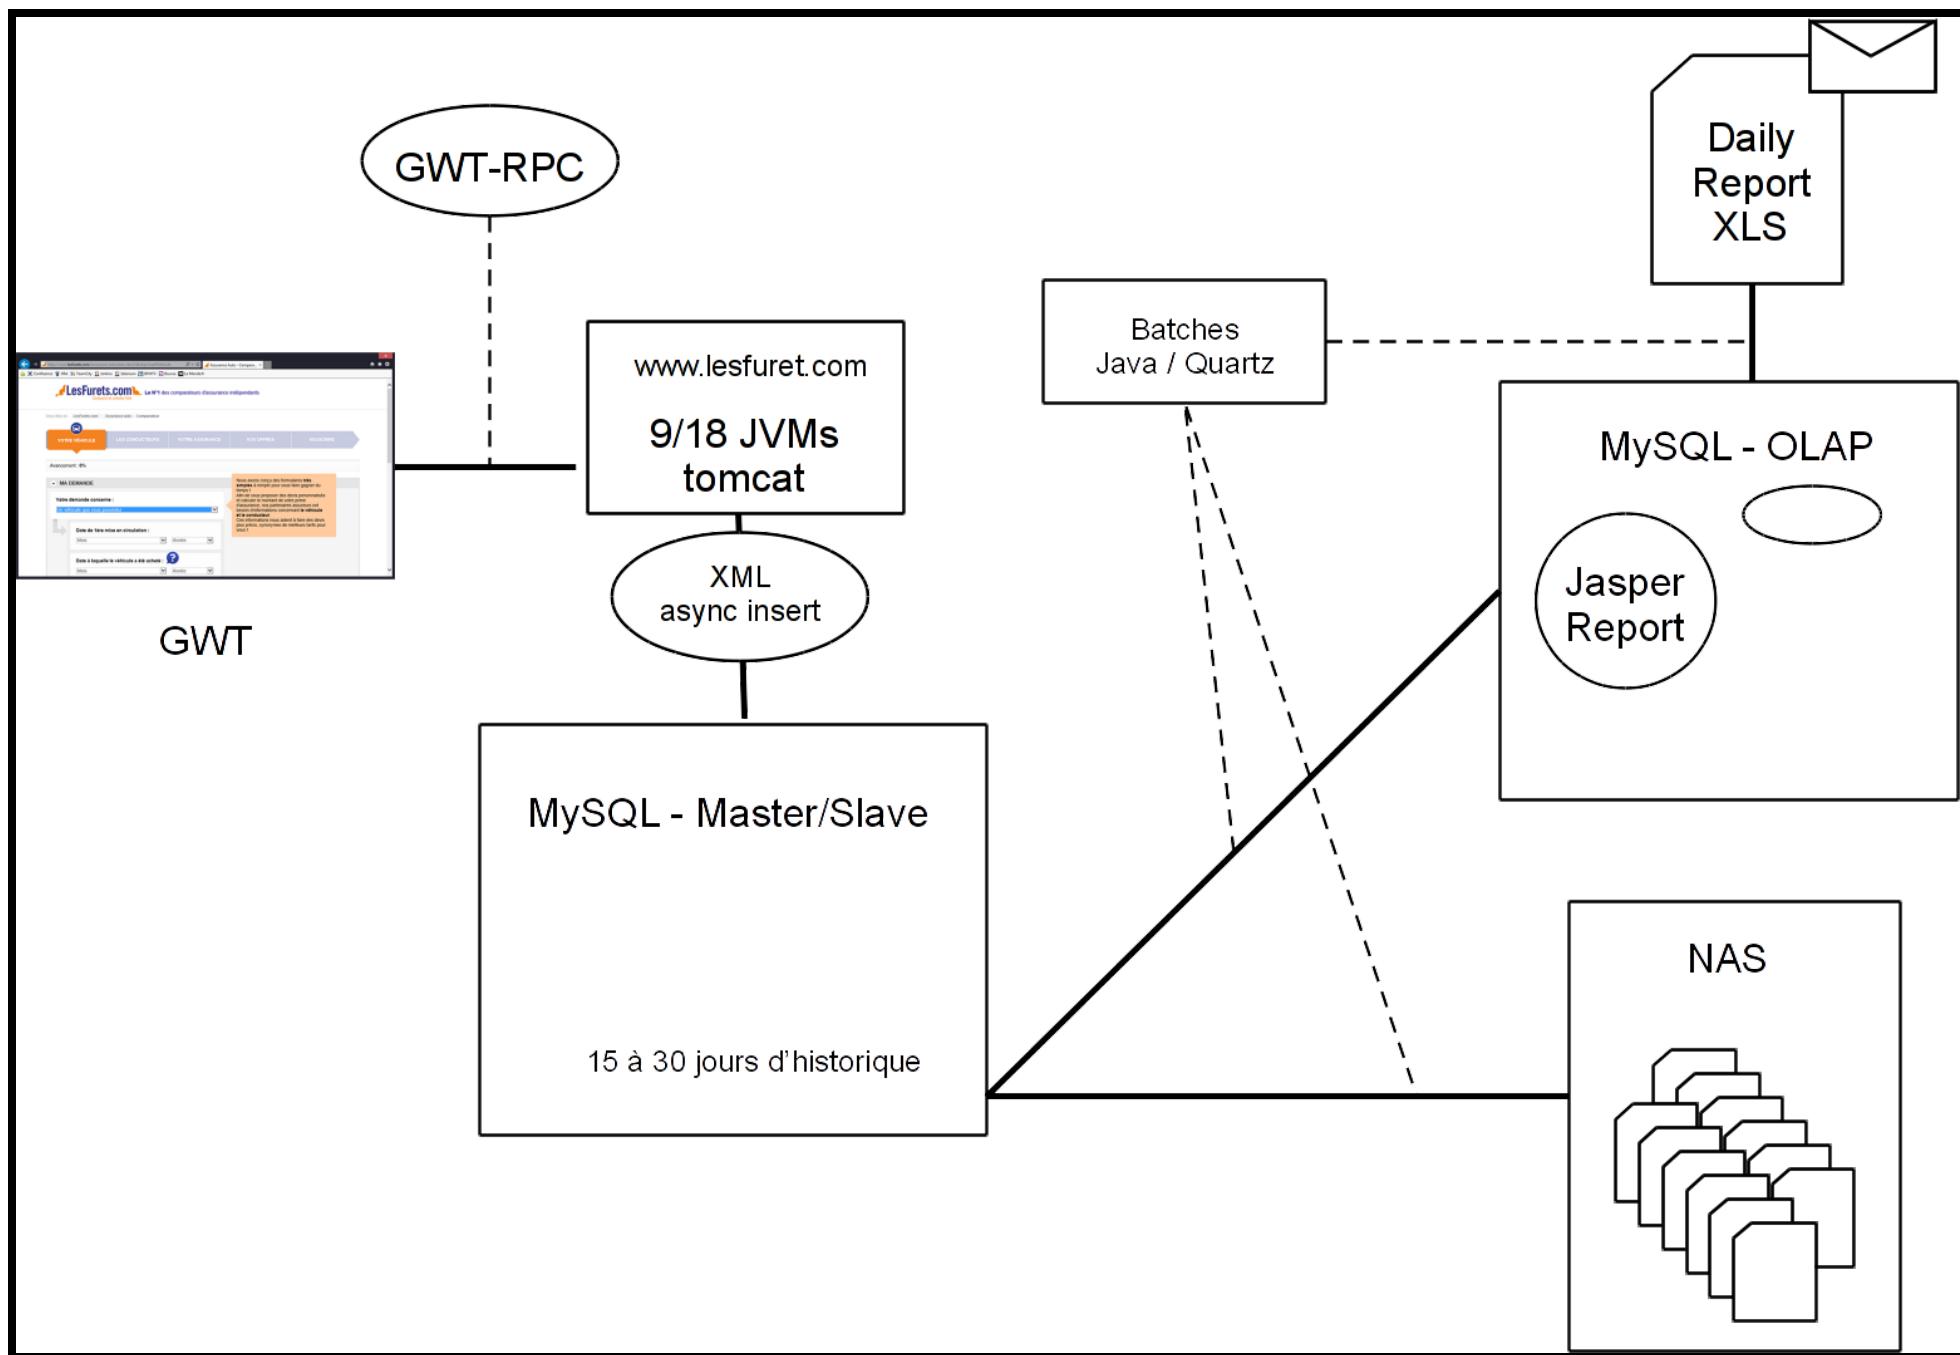
Data characteristics

- Specific workload
 - few updates
 - few synchronous queries
 - applicative caches
- Polyglot persistance
 - SQL
 - XML - blob (MySQL) and file (NFS)
 - ElasticSearch - centralized logging
 - Cassandra double run

Data workflow



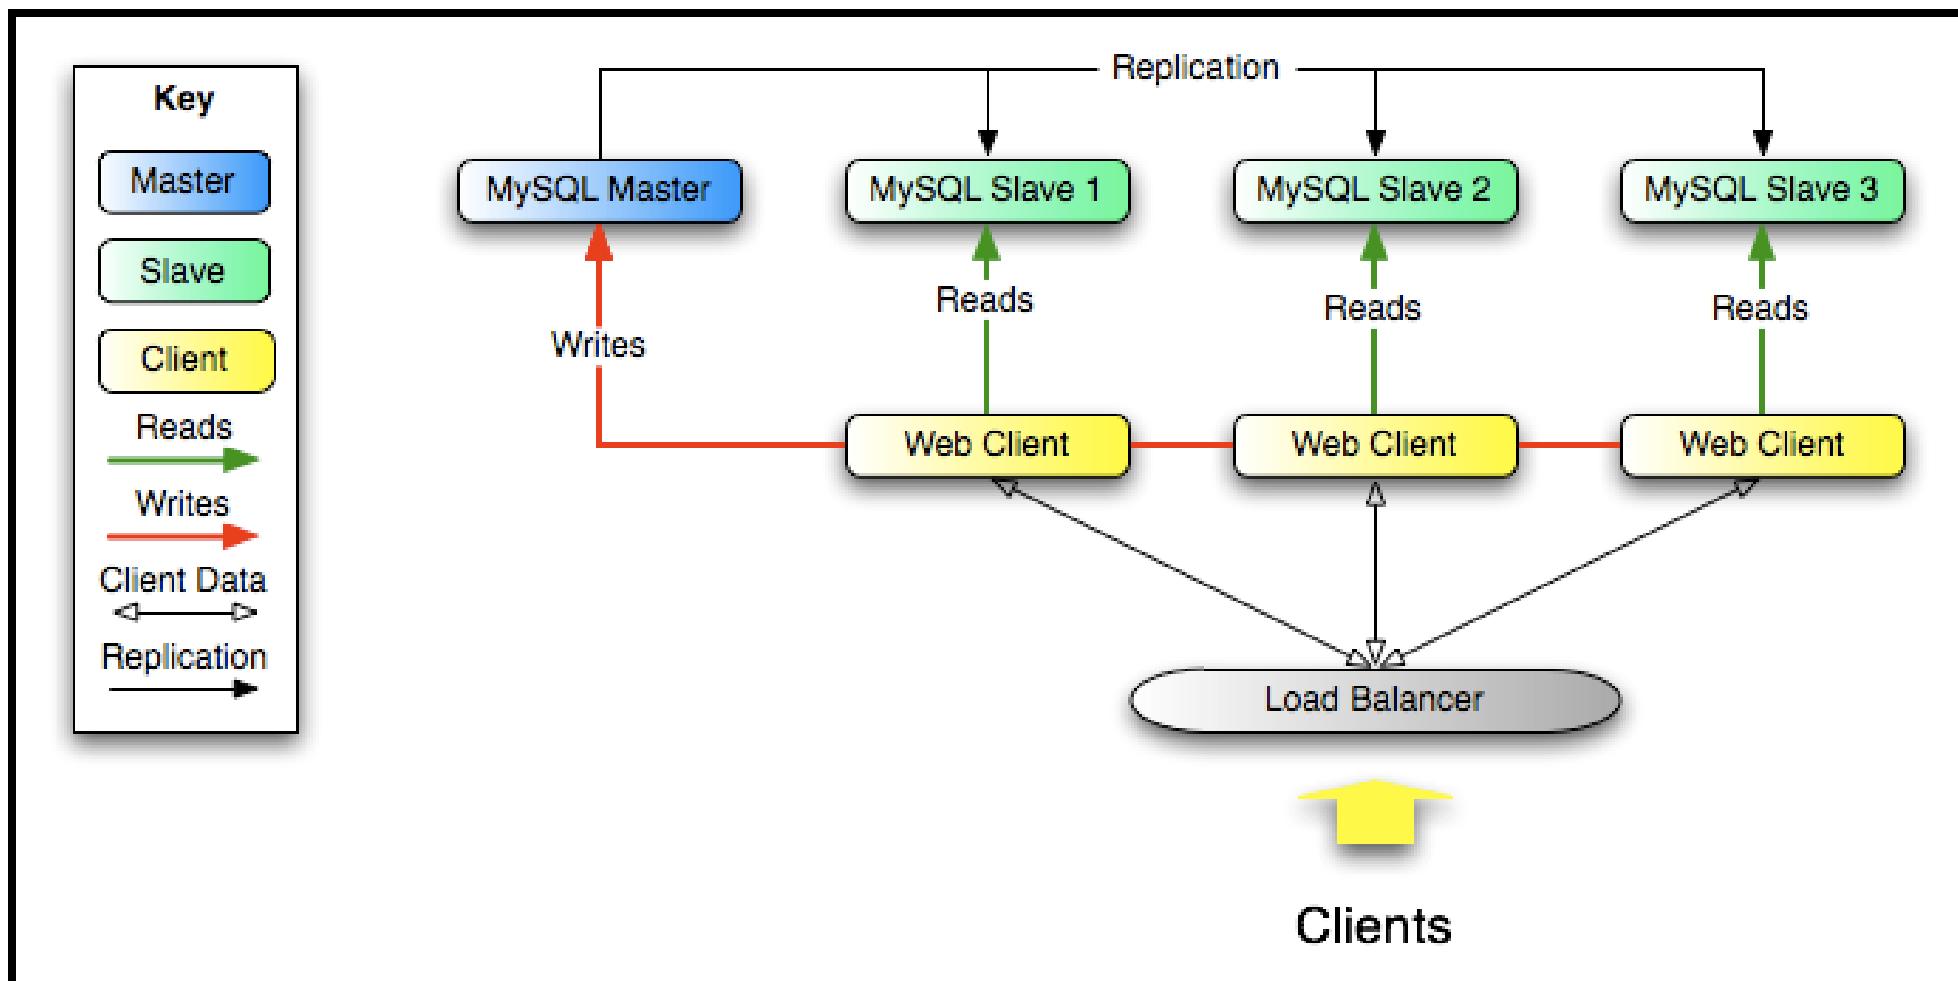
Architecture



Why replicate?

- scale out
- data security
- segregate workload (writes on master/ analytics on slaves)
- data distribution
- backups
- restore in a point in time (delayed replication)

MySQL scale-out



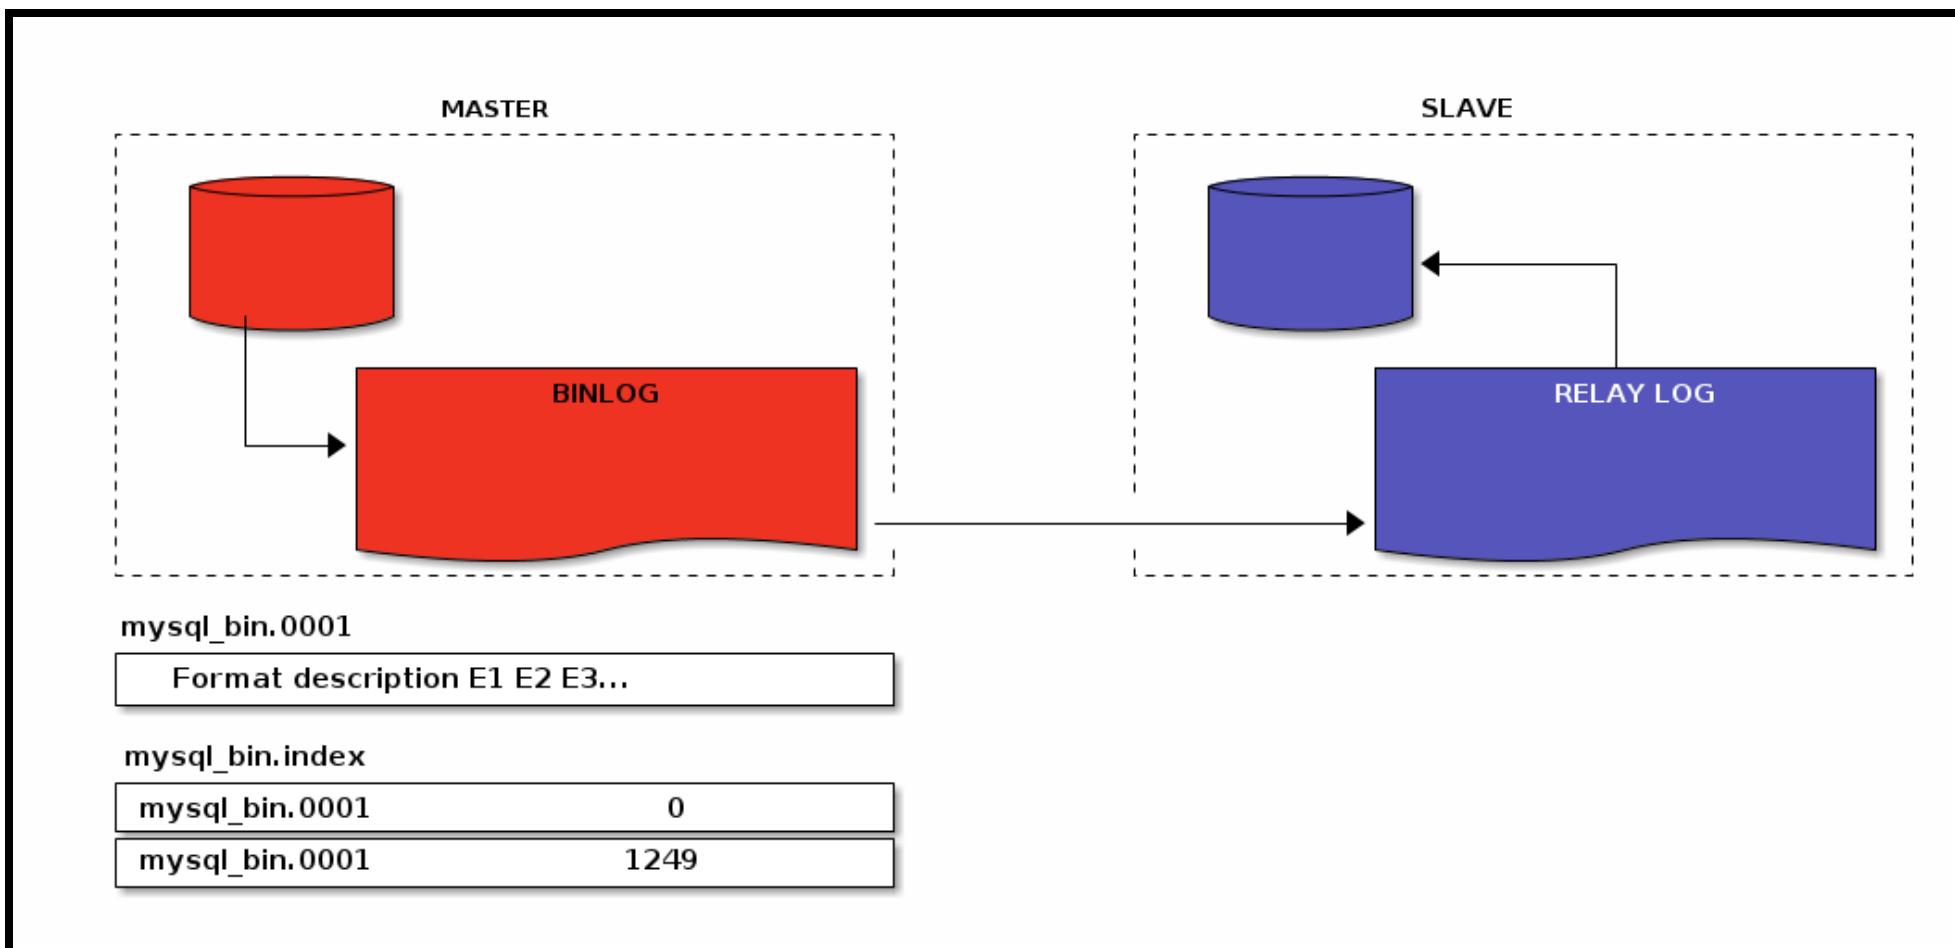
MySQL replication properties

- **1 Master ⇒ n Slaves**
- **Application transparent**
- **Full replication:** the full dataset is replicated on every node

MySQL replication

- **Master**
 - writes/updates only on master
 - log changes (events) on a bin-log
- **Slave(s)**
 - retrieve events from the master
 - reply the events

MySQL replication: binlog



MySQL replication

- one-way M→S replication
 - single threaded/multi-threaded (5.6+ 1 worker thread per database/slave)
 - **asynchronous** : wait until change recorded in local binlog
 - **semi-synchronous** : wait until one Slave ack received and stored event!)

What is replicated?

- SBR (statement based replication)
- RBR (row base replication)
- Mixed (choose by event size for every transaction)

SBR

- the slave replays the queries
- not all queries are safe for replication (!)
 - NOW, AUTOINCREMENT, TRIGGERS, TRANSACTIONS, SERVER STATE

RBR

- send all modified rows to the slave
- safer but costly
- 5.6+ optimizations (ignore blobs, increment, hashing..)

Semi-synchronous replication

- since 5.6+
- log SERVER_ID+TxID in a regular table
 - master blocks after the commit and waits until one semisynchronous slave acknowledges that it has received all events for the transaction or timeout
 - slave acknowledges receipt of a transaction's events only after the events have been written to its relay log

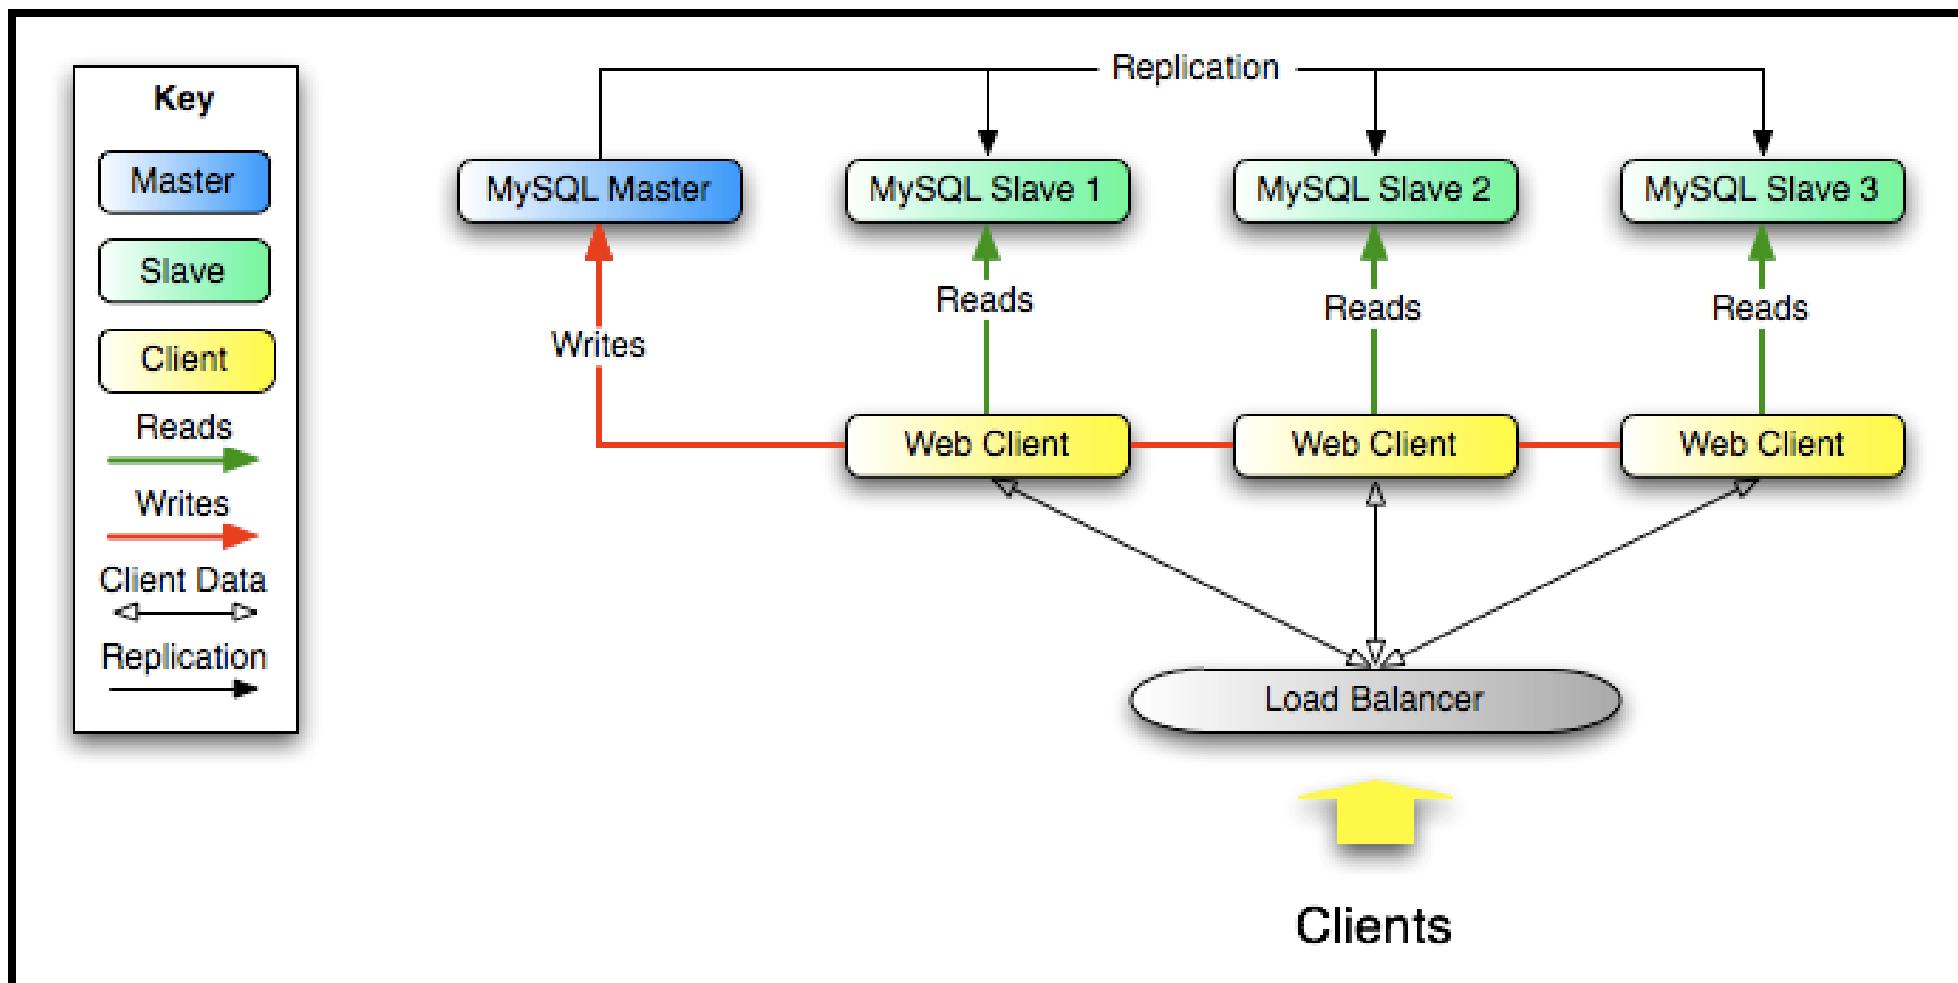
Semi-synchronous replication

- temporary switch to async-replication if no ACK from the slave
- **guarantees consistency between master and slave**
 - as long as all transactions committed on the master have also been applied on a slave

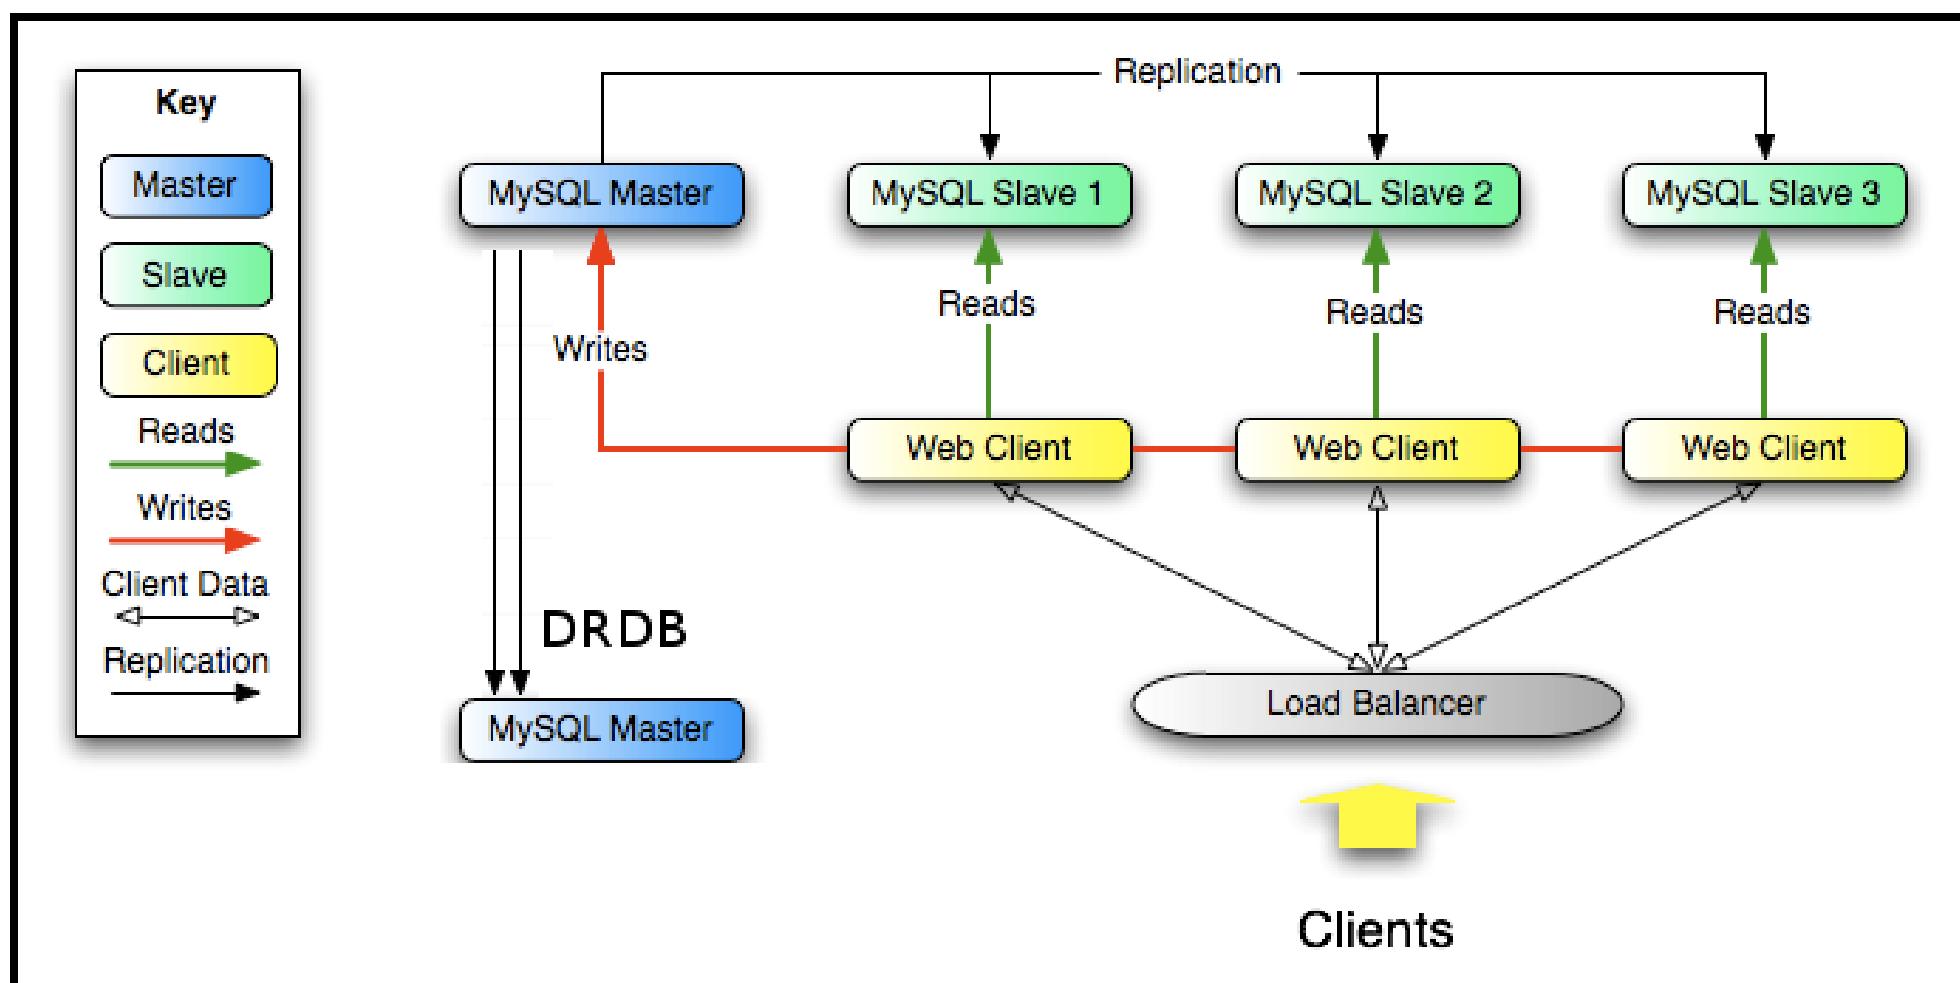
Crash recovery (M)

- manually promote a slave as a new master (<5.6)
- automatic slave promotion via `mysqlfailover` (5.6+)

Master/Slave

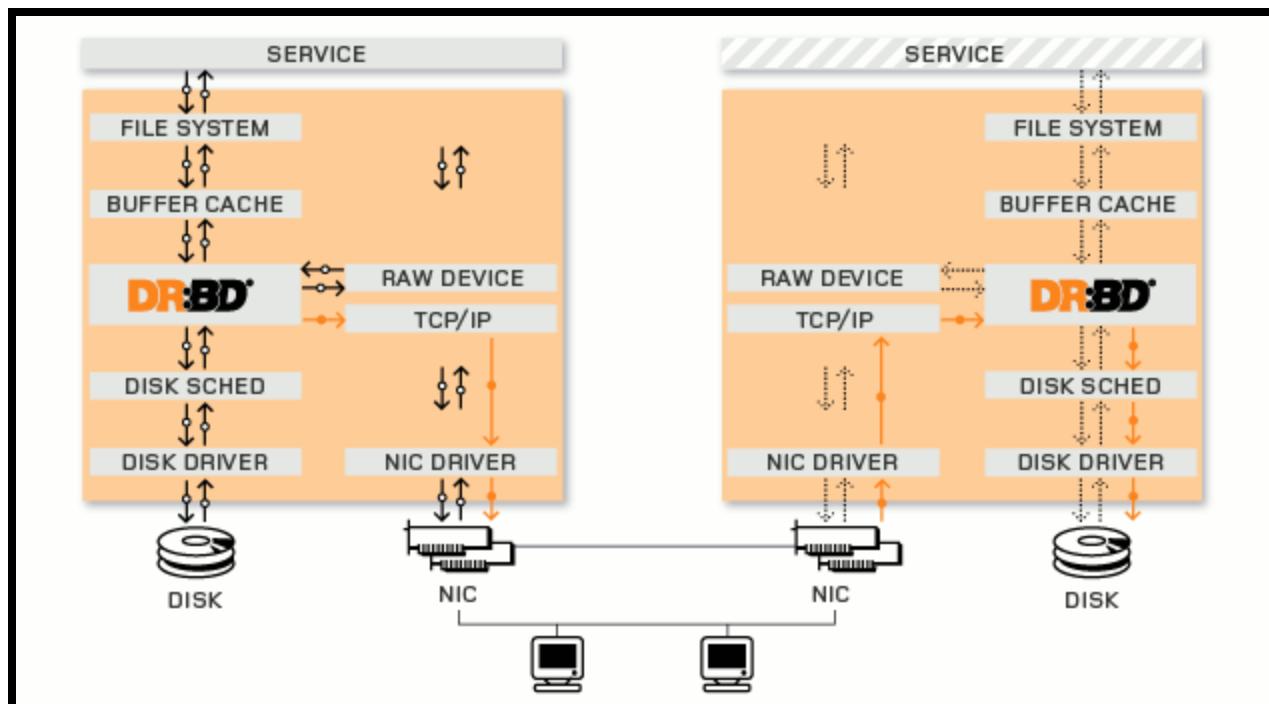


Multi - Master replication with DRBD



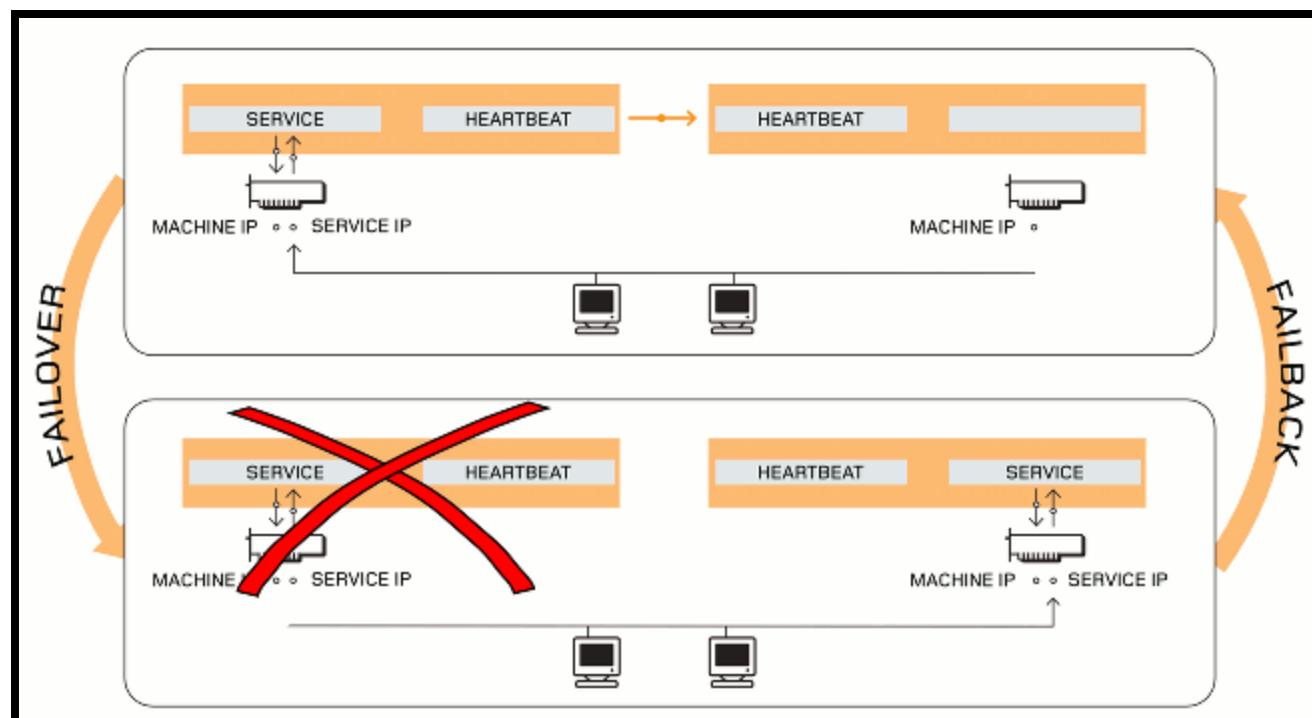
DRBD mirroring

- mirroring a linux partition over IP (sync/async)



DRBD HA

- heartbeat protocol monitors failures
- triggers service switch via IPFOs



more...

DRBD recovery

- node(s) outage
 - background sync (most up-to date node if both were down)
 - replication network outage
 - automatic recovery
 - storage subsystem
 - mostly transparent
- network partition
 - *split brain!* both nodes switched to the primary role while disconnected
 - *Manual intervention needed*

Scalling Writes

- Sharding : MySQL Fabric
- Multi-master replication
 - MySQL cluster
 - *Galera Cluster*

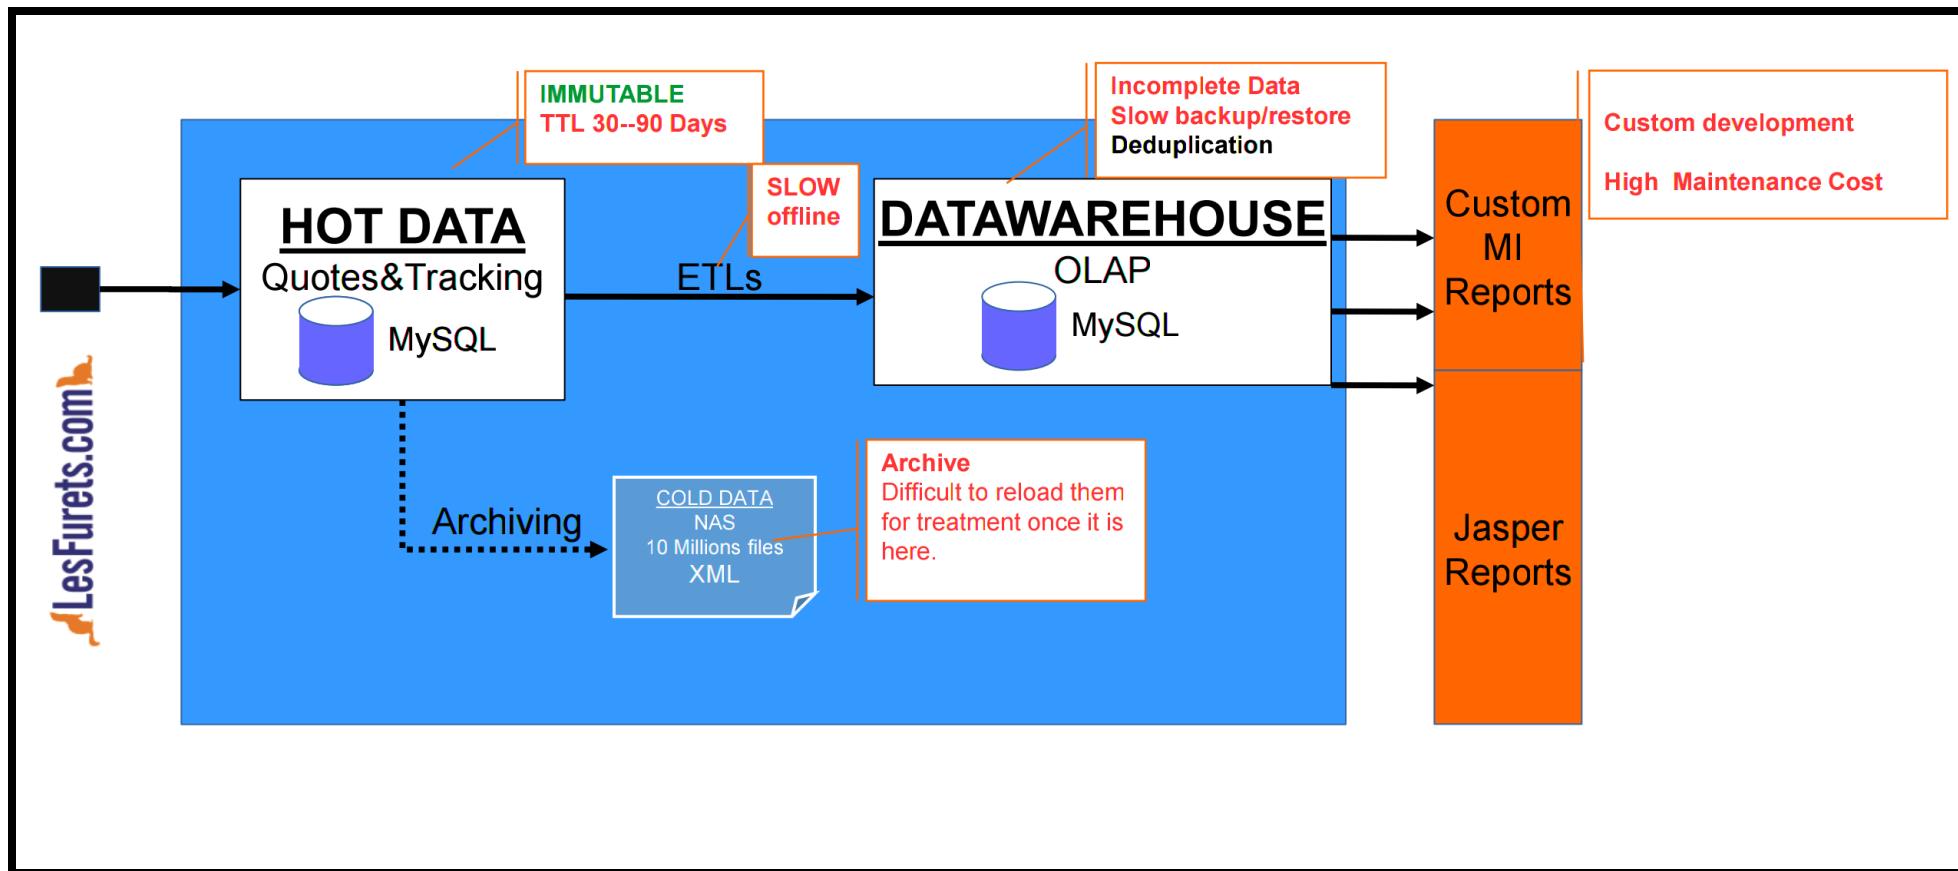
Galera Cluster

- multi-master: R/W at any node
- synchronous replication (tx commit)
 - traditionally implemented by distributed locking (2PH commit)
 - slow and complex
 - Galera ⇒ Certification Based Replication Method
 - good latency with increased consistency (*Don't be lazy be consistent paper*)
 - Easy migration from MySQL
 - combine with MySQL binlog replication
 - automatic node provisioning (XtraBackup)
 - transparent to applications*

LF Mariadb Galera Cluster

- 3 nodes cluster (3M)
- relaxed ACID (read after write)
- wsrep_sync_wait: ensures synchronous read view before executing an operation of the type specified by bitmask:
 1. READ (SELECT, SHOW & BEGIN/START TRANSACTION);
 2. UPDATE and DELETE;
 3. INSERT and REPLACE.
- Benefits
 - recover from failures ⇒ < 1s
 - typical replication lag < 10ms

Current architecture concerns



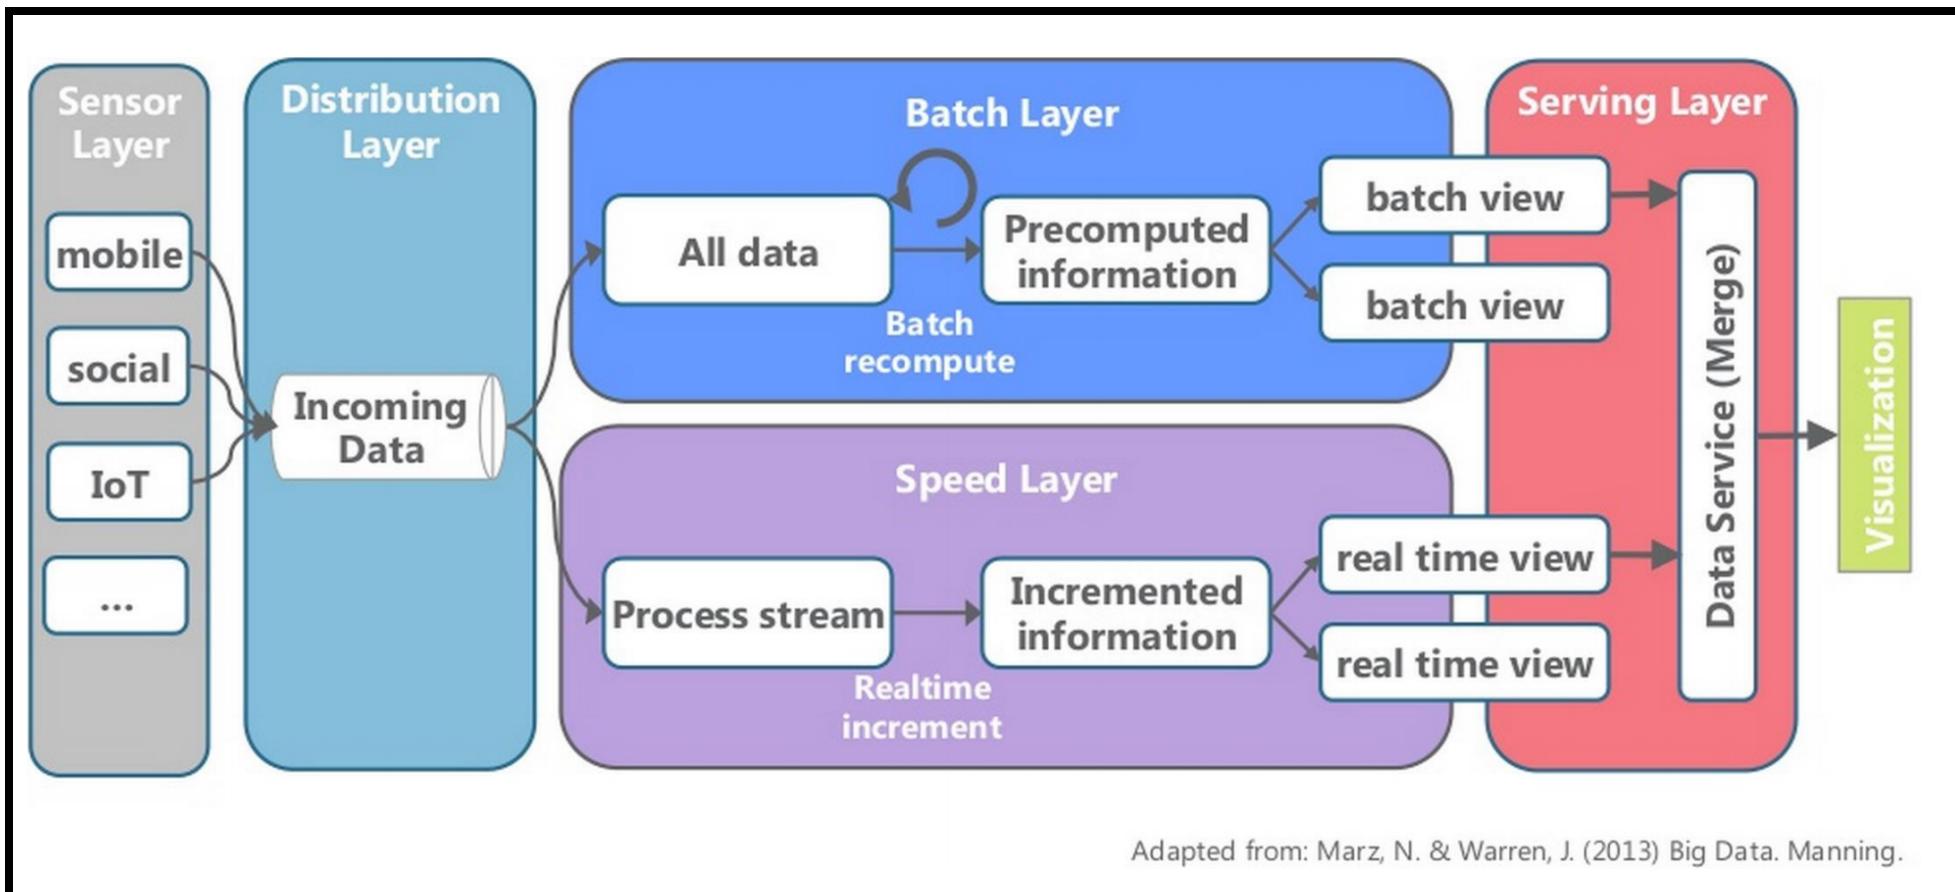
Feedback on scaling MySQL

- Scale UP limits ⇒ Olap size
- Scale OUT ⇒ replication and sharding
- Replication incidents
 - DRDB split brain
 - broken replication: network incidents
- **MySQL replication vs consistency:**
 - slowly diverging Master/Slave
 - not automatic check/resynchronisation → costly automatic
 - problematic when failover switch
- **Galera Cluster:** 1 year of transparent, incident free production

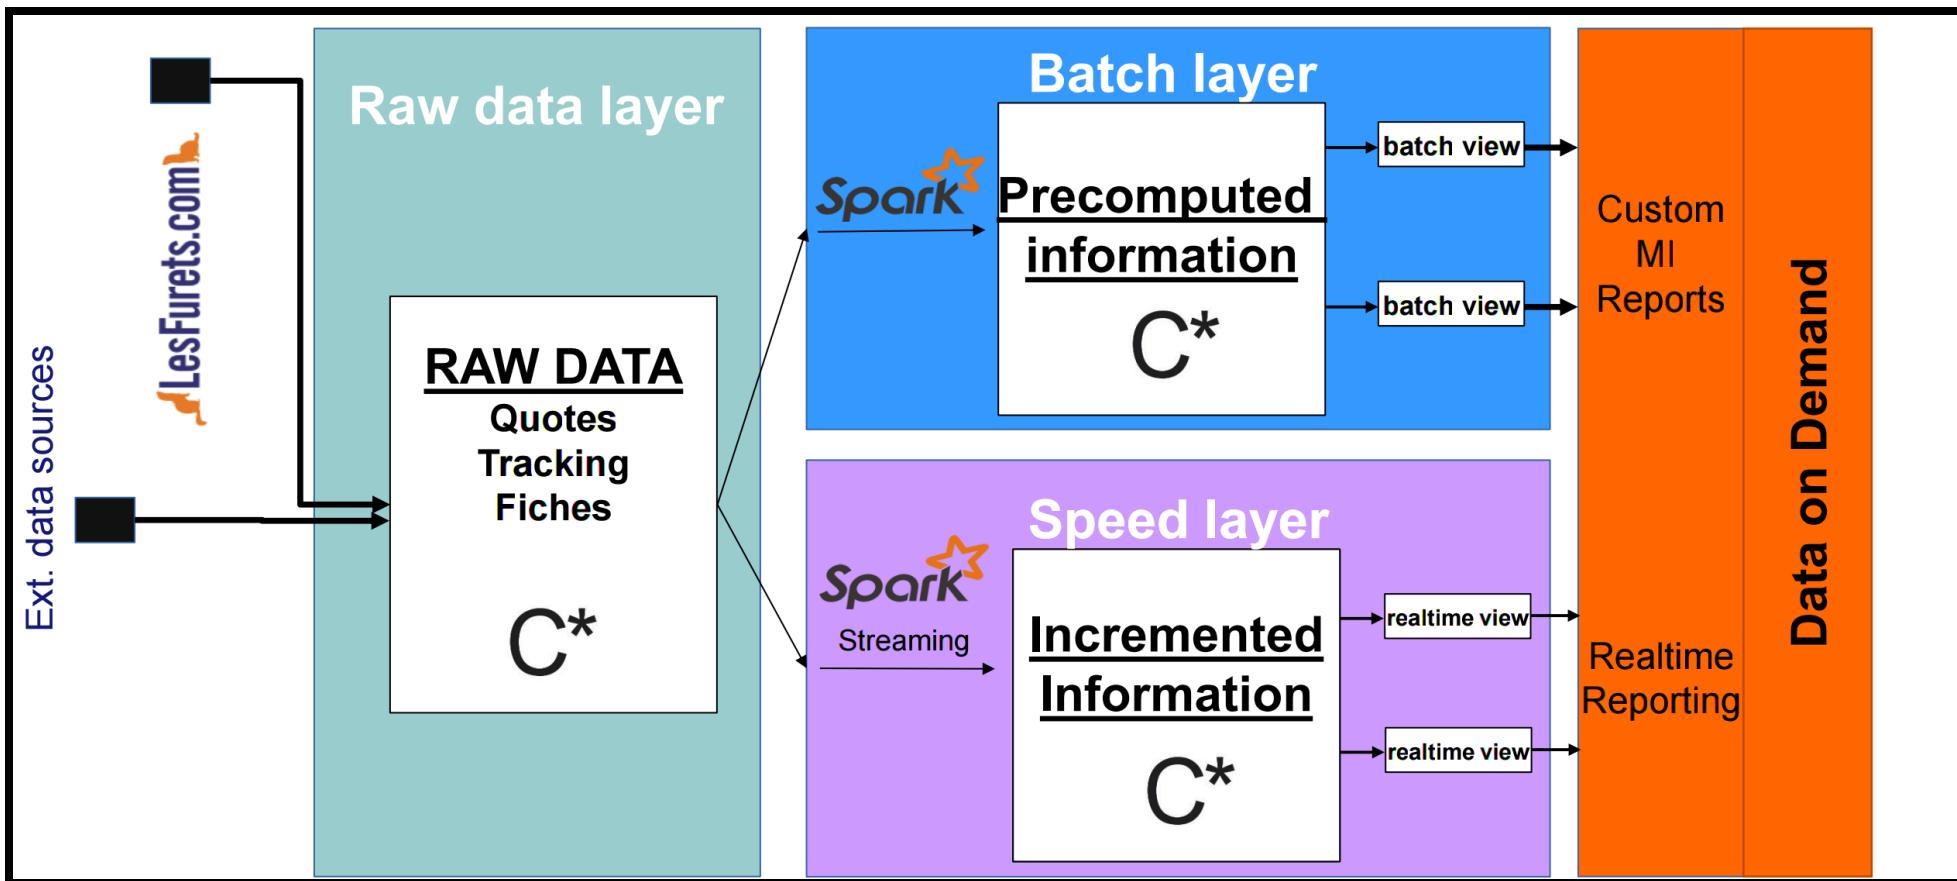
Lambda architecture (1)

- Combine best from No-SQL world to achieve
 - scalable systems
 - for arbitrary data problems
 - with human fault tolerance
 - and minimum complexity

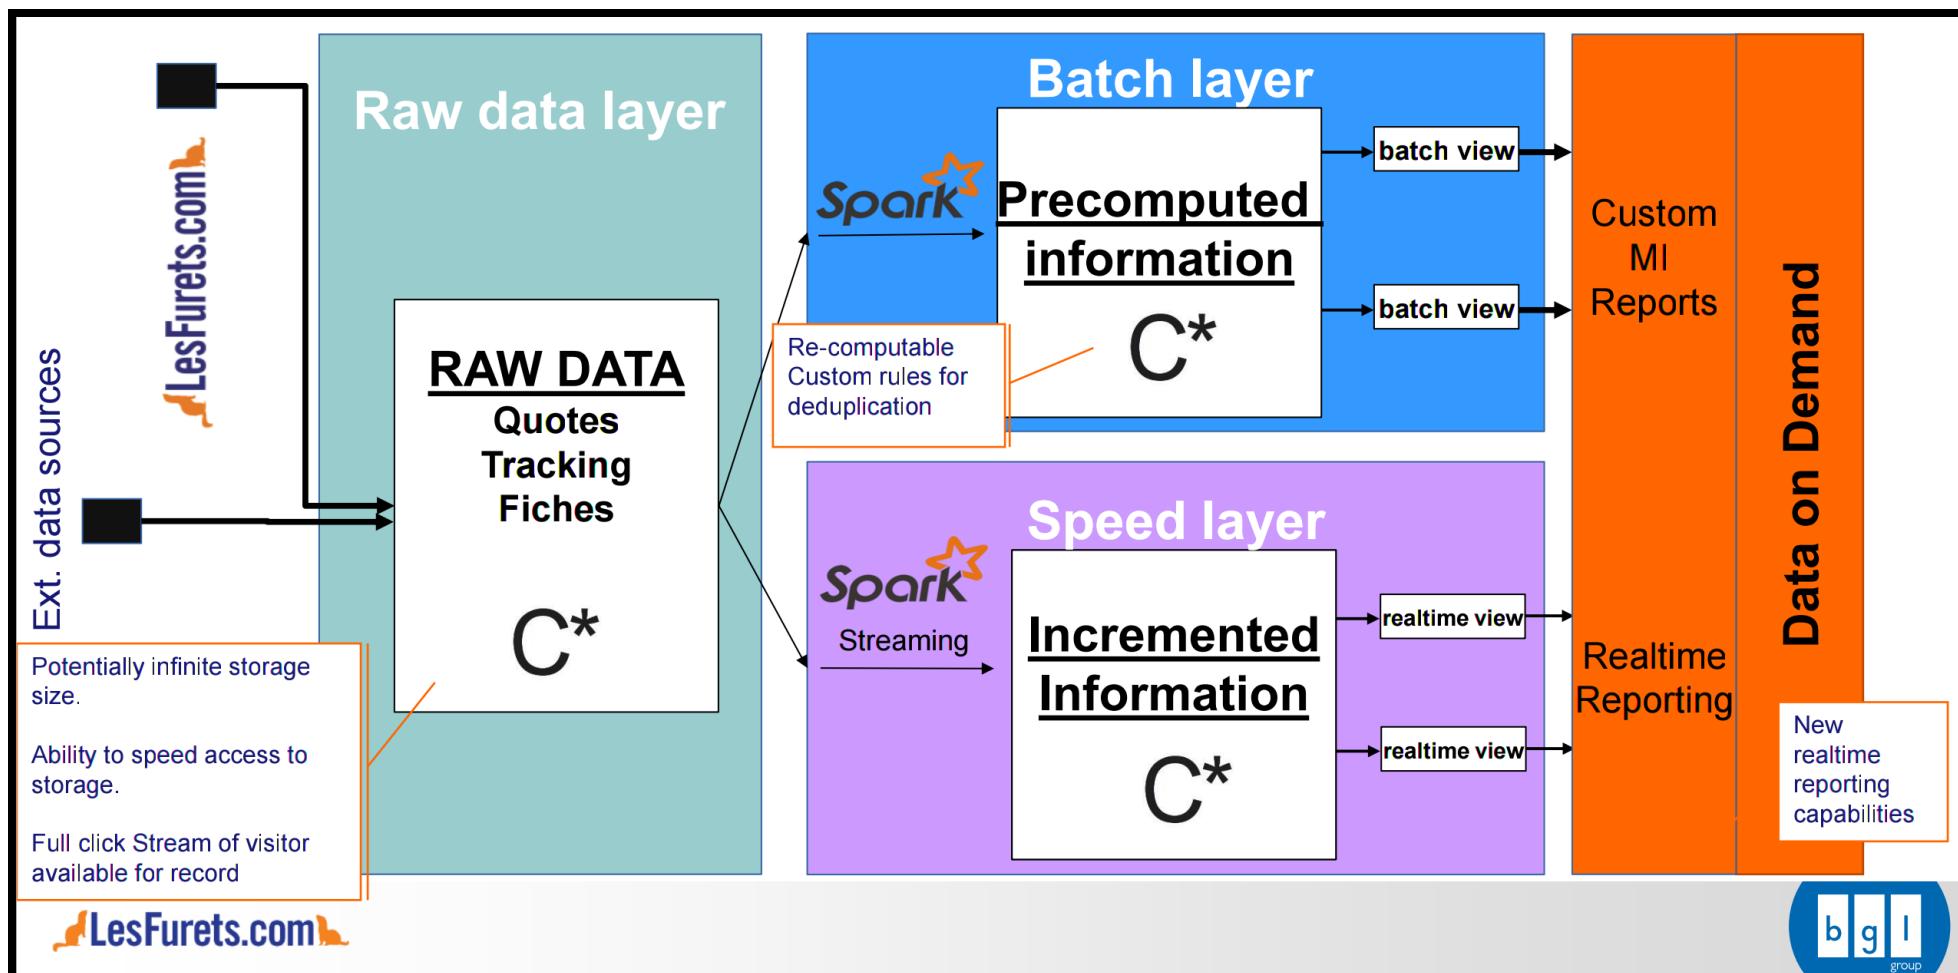
Lambda architecture (2)



On-going migration



On-going migration benefits



Plan

1. Intro
2. From SQL to NoSQL
3. Scalling MySQL : @LesFurets.com
4. PostgreSQL + TP

PostgreSQL

Postgres intro,

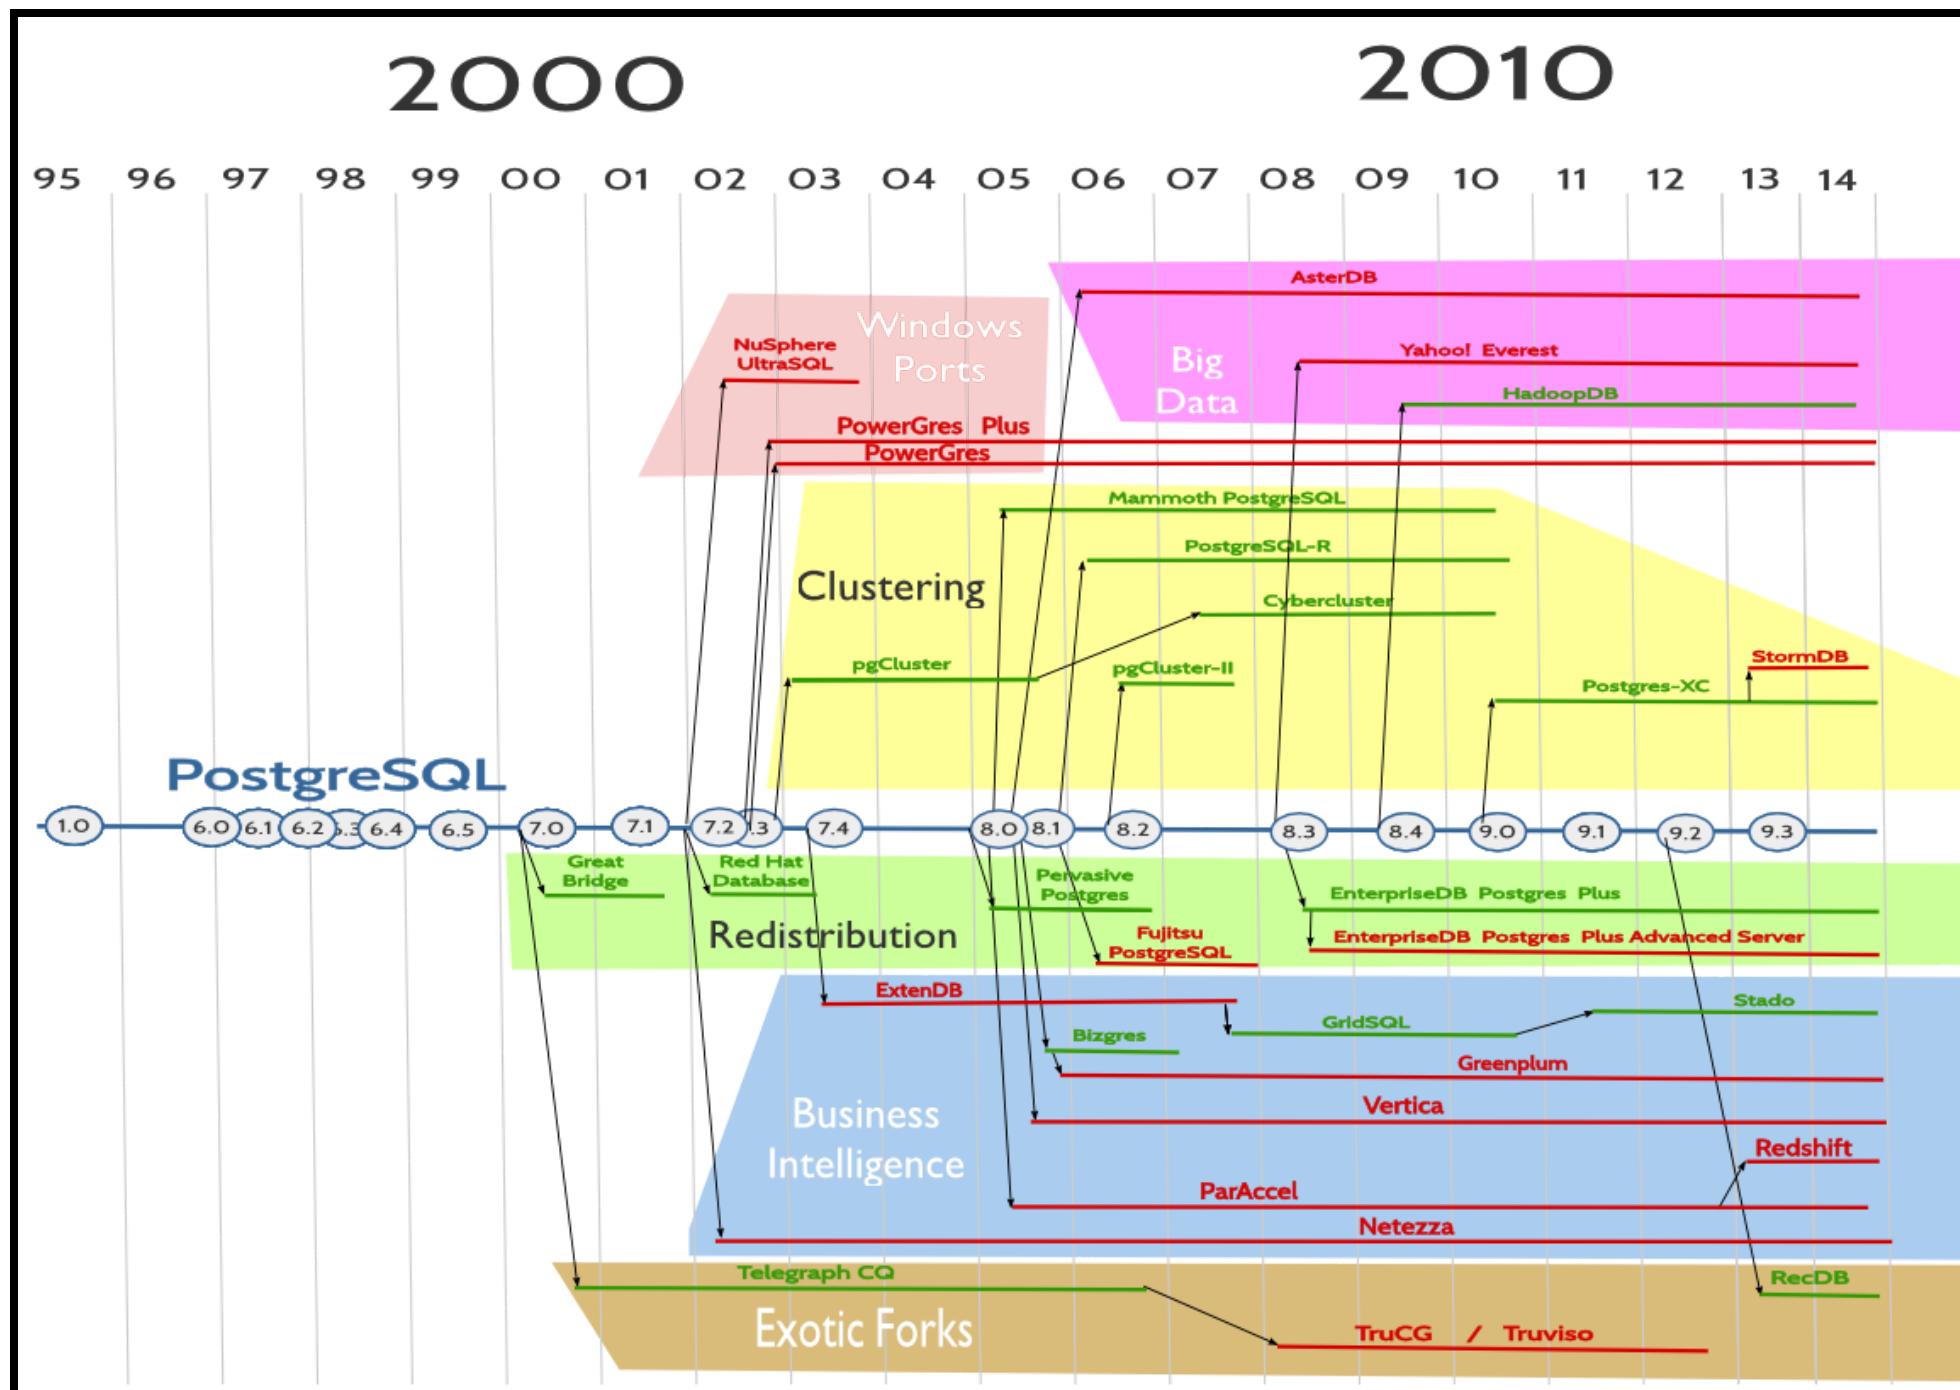
<http://momjian.us/main/presentations/overview.html>

PostgreSQL

- PostgreSQL
 - Berkeley : Ingres 1974
- Extensions
 - custom data types / operators
 - natural-language parsing
 - multidimensional indexing
 - geographic queries
- TP ⇒ PostgreSQL & extensions (*tablefunc*, *dict_xsyn*, *fuzzystrmatch*, *pg_trgm*, *cube*)

more...

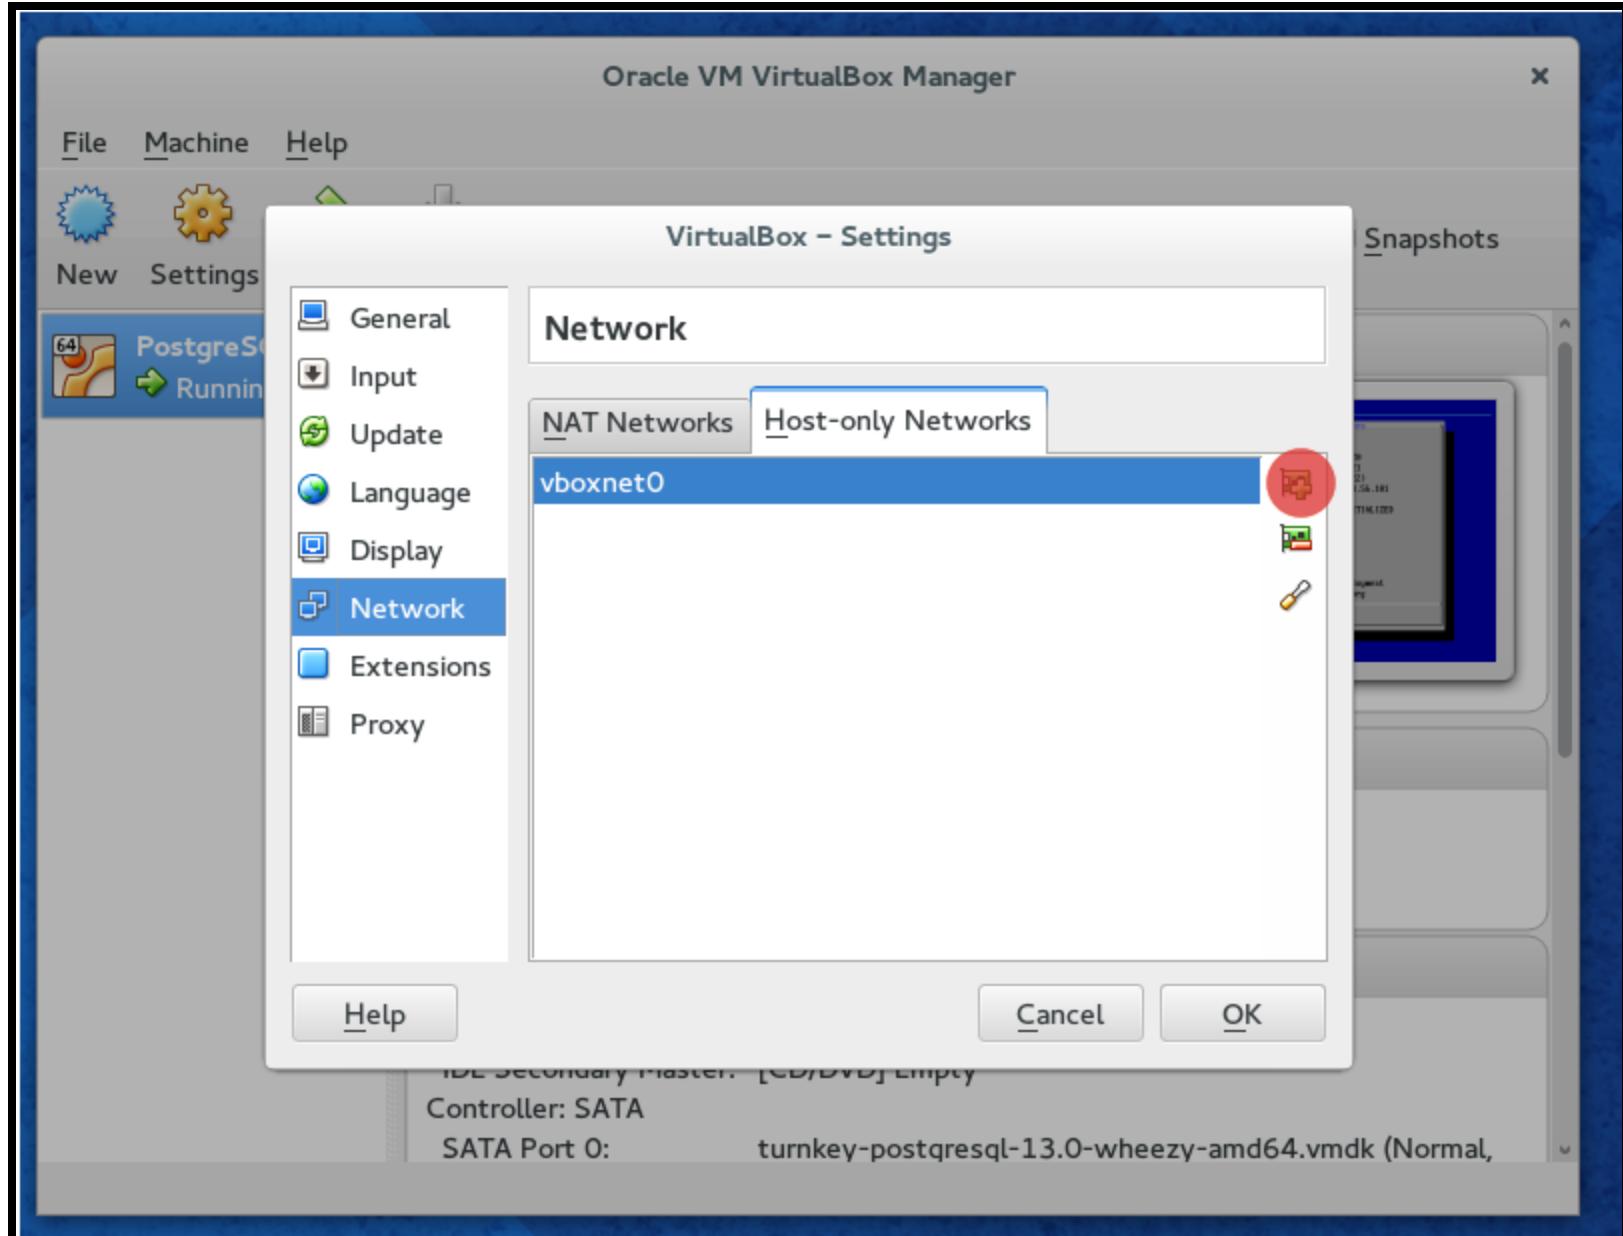
PostgreSQL offsprings



TP1: Installation de l'environnement

1. Démarrez l'application **Virtualbox**
2. Configurez un réseau host-only **vboxnet0** (File/Preferences/Network puis voir image) ou via les commandes suivantes:

```
VBoxManage hostonlyif create  
VBoxManage hostonlyif ipconfig vboxnet0 --ip 192.168.56.1  
VBoxManage dhcpserver add --ifname vboxnet0 --ip 192.168.56.1\  
--netmask 255.255.255.0 --lowerip 192.168.56.100\  
--upperip 192.168.56.200  
VBoxManage dhcpserver modify --ifname vboxnet0 --enable
```



TP1: Telecharger la VM pour le tp

1 Telecharger l'image Virtualbox: <http://bit.ly/telecom-postgresql>

TP1: Importer la VM

2 File/Import appliance ...



TP1: Start login to the VM

3 Démarrez la VM

4 Accédez a l'interface web <https://192.168.X.Y/> (voir la console de la VM)

```
Login  
User: postgres  
Pass: bigdata14
```

TP1: phpPgAdmin

The screenshot shows the phpPgAdmin interface for PostgreSQL 9.1.9. The left sidebar displays a tree view of the database structure under the 'PostgreSQL' server node, specifically the 'postgres' schema. The visible nodes include 'Schemas' (with 'public' selected), 'Tables' (containing 'actors', 'genres', 'movies', and 'movies_actors'), 'Views', 'Sequences', 'Functions', 'Full Text Search', and 'Domains'. Other nodes like 'Slony' and 'root' are also listed. The main panel shows the results of a query on the 'genres' table, which lists 18 genres with their corresponding positions. The total runtime of the query is 0.607 ms.

name	position
Action	1
Adventure	2
Animation	3
Comedy	4
Crime	5
Disaster	6
Documentary	7
Drama	8
Eastern	9
Fantasy	10
History	11
Horror	12
Musical	13
Romance	14
SciFi	15
Sport	16
Thriller	17
Western	18

18 row(s)

Total runtime: 0.607 ms

SQL executed.

[Edit SQL](#) | [Create report](#) | [Download](#)

TP1: Recherche et recommandation

Moteur de recherche et recommandation des films:

- recherche: textuelle, approximative, phonétique
- moteur de recommandation très basique

TP1: Schéma (déjà créé)

```
CREATE TABLE genres (
    name text UNIQUE,
    position integer
);
CREATE TABLE movies (
    movie_id SERIAL PRIMARY KEY,
    title text,
    genre cube
);
CREATE TABLE actors (
    actor_id SERIAL PRIMARY KEY,
    name text
);

CREATE TABLE movies_actors (
    movie_id integer REFERENCES movies NOT NULL,
    actor_id integer REFERENCES actors NOT NULL,
    UNIQUE (movie_id, actor_id)
```

```
CREATE INDEX movies_actors_movie_id ON movies_actors (movie_id);  
CREATE INDEX movies_actors_actor_id ON movies_actors (actor_id);  
CREATE INDEX movies_genres_cube ON movies USING gist (genre);
```

Create schema script Import data script

TP1: Recherche

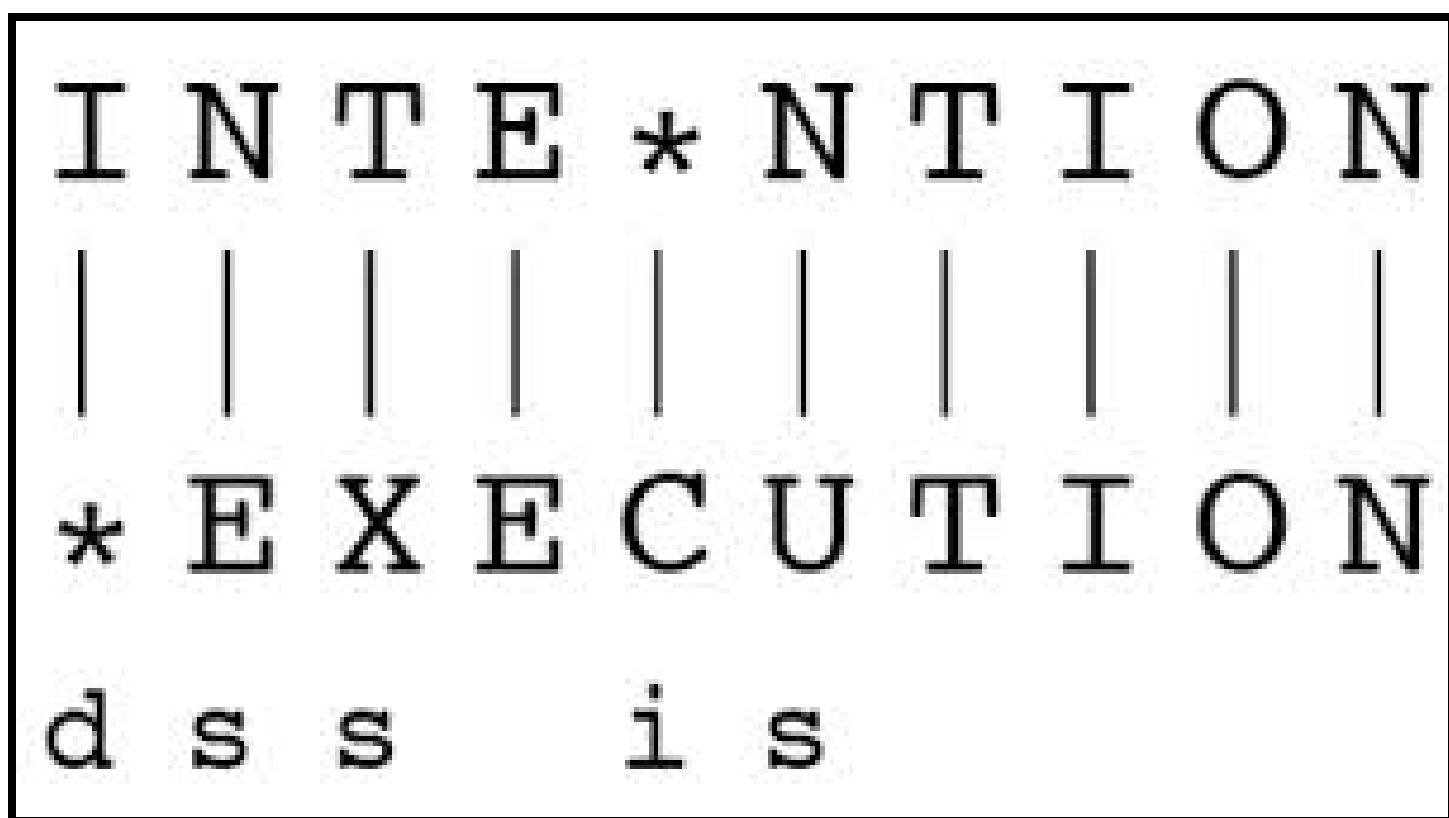
- *Recherche exacte / pattern matching*
- *Distance de Levenshtein* → typos simples
- *N-gram/similarité* → trouver les erreurs modérées
- *Full text match @@* → similarité grammaticale
- *Métagone* → similarité phonétique

TP1: Recherche textuelle/patterns

Utilisez les opérateurs **LIKE** ou **RegEX** pour les requêtes suivantes:

1. Tous les films qui ont le mot ***stardust*** dans leur nom.
2. Compter tous les films dont le titre ne commence pas par le mot ***the***
3. Tous les films qui ont le mot ***war*** dans le titre mais pas en dernière position

TP1: Distance Levenshtein



- Opérations: Substitute, Insert, Delete
- Distance *Levenshtein* : nb minimal d'opérations

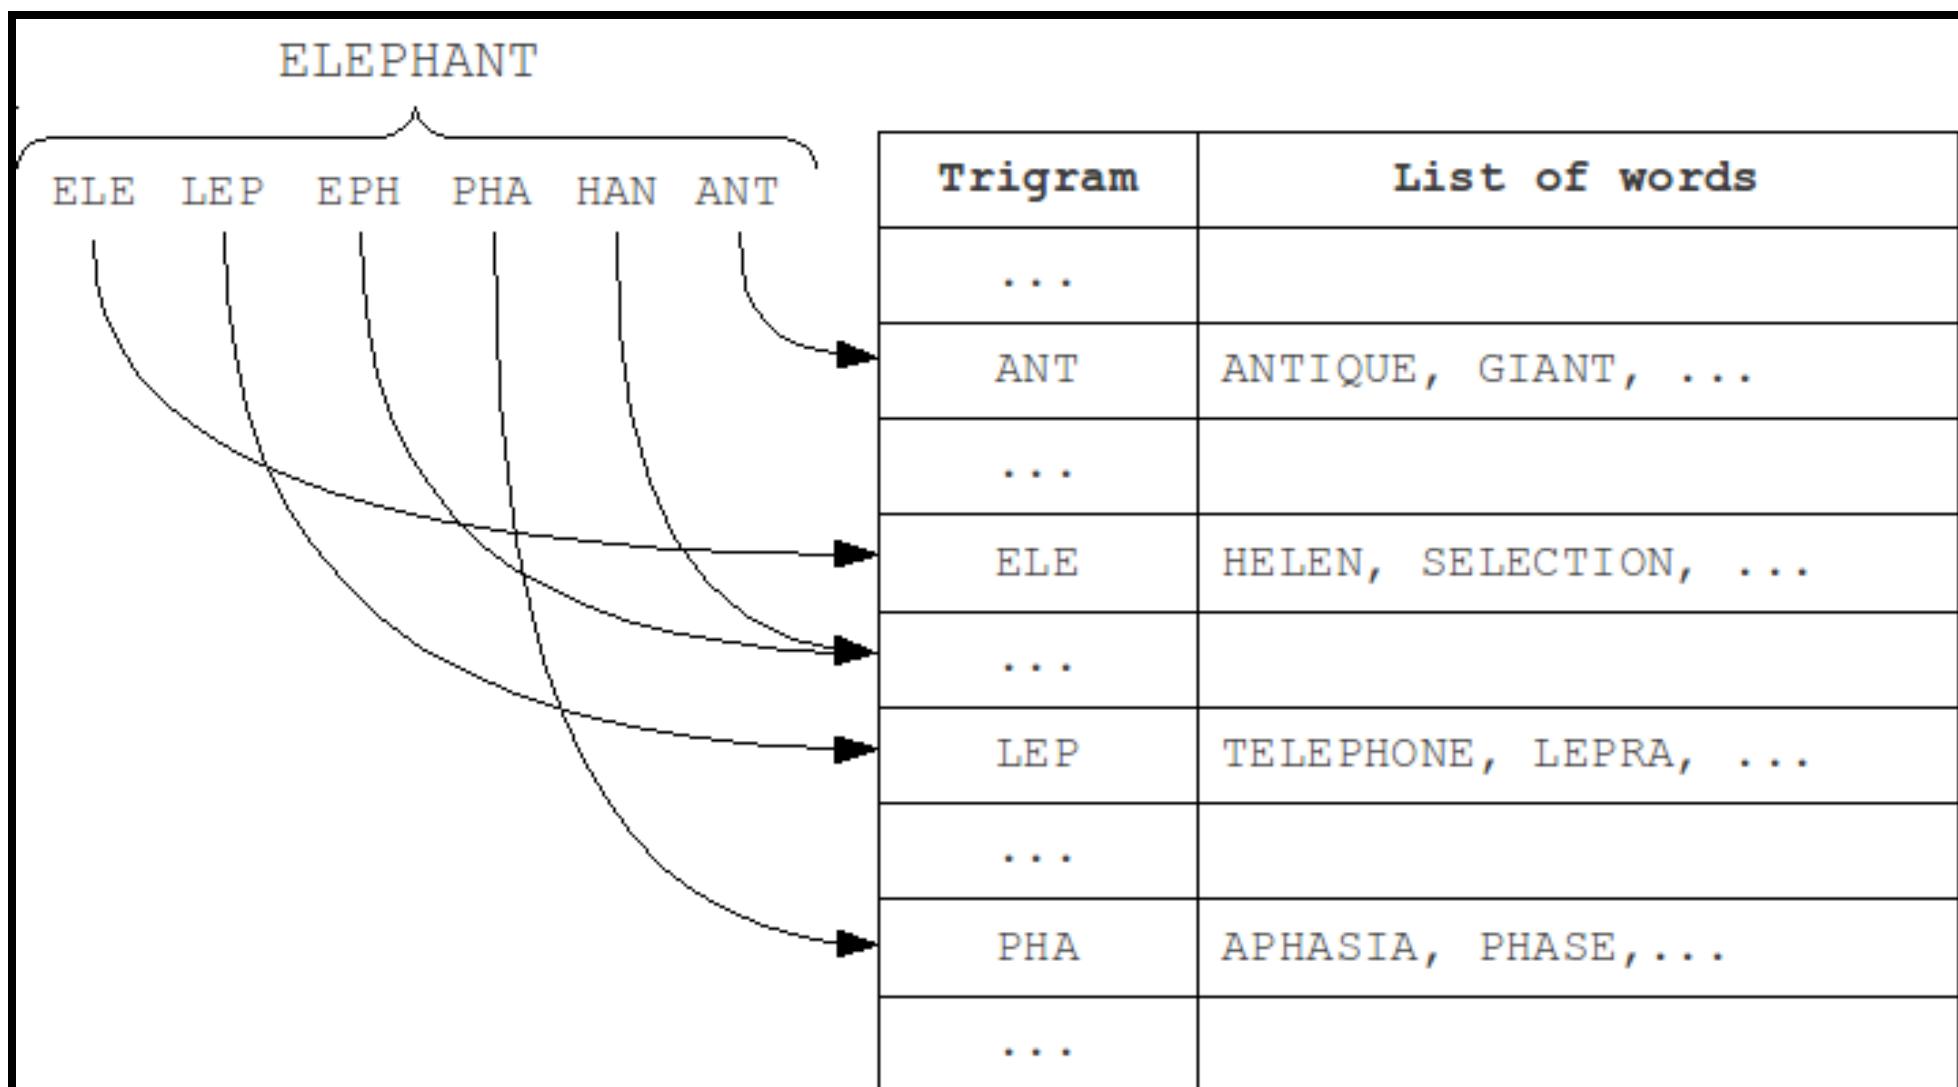
TP1: Distance Levenshtein

Utilisez les fonctions du package [fuzzystrgmatch](#) pour trouver :

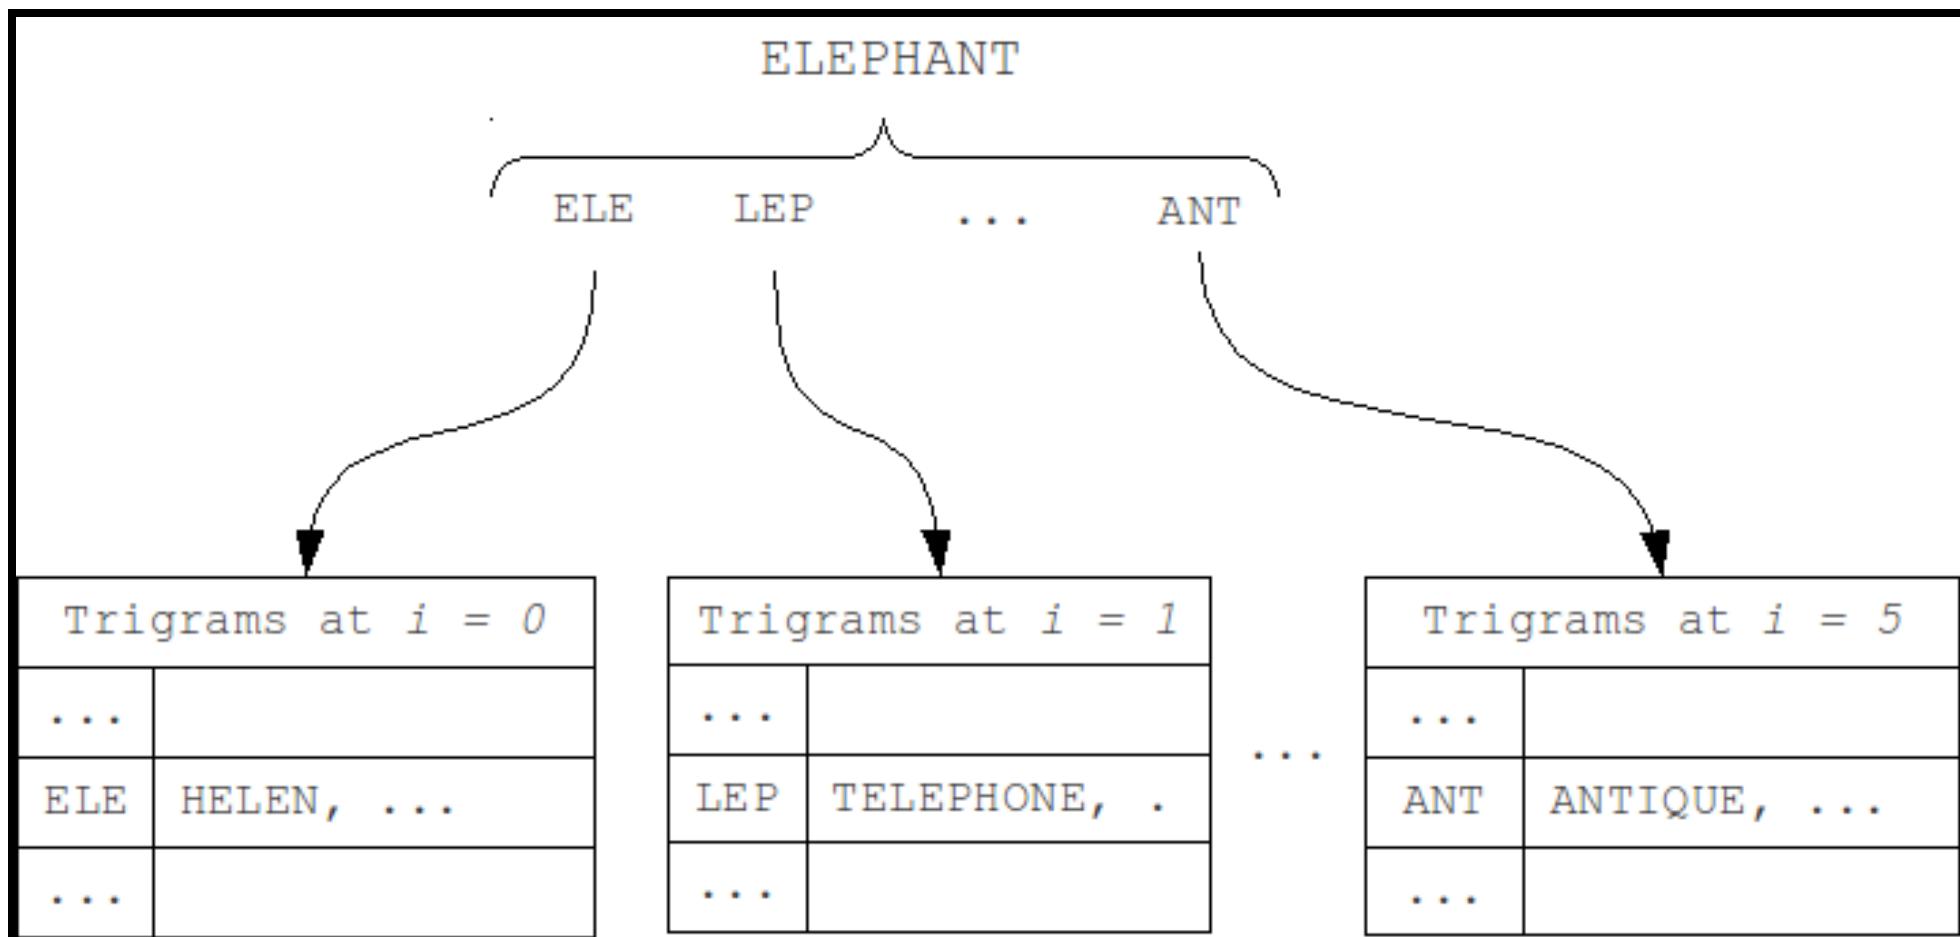
1. La distance levenshtein entre les mots ***execution*** et ***intention***
2. Tous les films qui sont a une distance *levenshtein* inférieure a 9 de la chaine suivante: ***a hard day nght***

[Documentation extension fuzzystrgmatch](#)

TP1: N-gram



TP1: N-gram, similarity search



TP1: N-gram, similarity search (%)

Écrivez les requêtes pour trouver :

1. Tous les tri-grammes du mot ***Avatar***
2. La similarité entre **VOTKA** et **VODKA**
3. Tous les films dont le titre est similaire a plus de 0.1% du titre ***Avatar***.

Documentation extension trgm

TP1: Full text search

Trouver les films qui contiennent les formes grammaticales des mots 'night' et 'day':

(ignorer les mots de liaison/ pluriel/etc..)

Algorithme:

1. ***extraire les racines des mots (lexèmes) → spécifiques au langage !***
2. comparer les vecteurs des lexèmes

TP1: Full text search

```
SELECT to_tsvector('A Hard Day''s Night'),  
       to_tsquery('english', 'night & day');  
  
      to_tsvector          |      to_tsquery  
-----+-----  
'day':3 'hard':2 'night':5 | 'night' & 'day'
```

- ***tsvector*** : lexèmes :position
- ***tsquery*** : lexèmes séparées par &
- ***spécifique au language !***

Full text search

```
SELECT title  
FROM movies  
WHERE to_tsvector(title) @@ to_tsquery('english', 'night & day');
```

```
SELECT title  
FROM movies  
WHERE title @@ 'night & day';
```

A Hard Day's Night
Six Days Seven Nights
Long Day's Journey Into Night

TP1: Recherche phonétique

- plusieurs fonctions pour la codification phonétique des mots

```
SELECT name, dmetaphone(name), dmetaphone_alt(name),
       metaphone(name, 8), soundex(name)
  FROM actors;

      name      | dmetaphone | dmetaphone_alt | metaphone | soundex
-----+-----+-----+-----+-----+
  50 Cent      | SNT        | SNT          | SNT      | C530
Aaron Eckhart | ARNK       | ARNK         | ARNKHRT  | A652
Agatha Hurle  | AK0R       | AKTR         | AK0HRL   | A236
```

Documentation fuzzystrmatch...

TP1: Recherche phonétique

1. Trouver les films qui ont des acteurs dont les noms se prononcent pareil.
2. Trouver les acteurs avec un nom similaire a ***Robin Williams***, triés par similarité (combiner %, metaphone et levenshtein):

actor_id	name
4093	Robin Williams
2442	John Williams
4479	Steven Williams
4090	Robin Shou

TP1: Search

- *Recherche exacte / pattern matching*
- *Distance de Levenshtein* → typos simples
- *N-gram/similarité* → trouver les erreurs modérées
- *Full text match @@* → similarité grammaticale
- *Métagone* → similarité phonétique

TP1: Recherche multi-dimensionnelle

```
CREATE TABLE movies (
    movie_id SERIAL PRIMARY KEY,
    title text,
    genre cube (1)
);

INSERT INTO movies (movie_id,title,genre) VALUES
(1, 'Star Wars',
'(0,7,0,0,0,0,0,0,7,0,0,0,0,10,0,0,0)' ) (2),
)
```

- on utilise le type cube <1> pour mapper les notes sur un vecteur n-dimensionnel de valeurs (= score du film <2>)

TP1: Recherche multi-dimensionnelle

- les noms pour les dimensions sont définis dans la table genres

```
CREATE TABLE genres (
    name text UNIQUE,
    position integer
);

INSERT INTO genres (name,position) VALUES
('Action',1),
('Adventure',2),
('Animation',3),
...
('Sport',16),
('Thriller',17),
('Western',18);
```

TP1: Recherche multi-dimensionnelle

Utiliser le module **cube** pour recommander des films ***similaires*** (du même genre)

- Afficher les notes du film ***Star Wars***
- Quelle est la note du film ***Star Wars*** dans la catégorie 'Animation'
- Afficher les films avec les meilleures notes dans la catégorie SciFi

TP1: Recherche multi-dimensionnelle

- Afficher les films similaires (`cube_distance`) à *Star Wars* (vecteur = $(0, 7, 0, 0, 0, 0, 0, 0, 0, 7, 0, 0, 0, 0, 10, 0, 0, 0)$) du plus similaire au moins similaire

title	dist
Star Wars	0
Star Wars: Episode V - The Empire Strikes Back	2
Avatar	5
Explorers	5.74456264653803
Krull	6.48074069840786
E.T. The Extra-Terrestrial	7.61577310586391

- Écrivez une requête pour trouver les films qui sont à moins de 5 points de différence sur chaque dimension (utiliser `cube_enlarge` et `@>`).

Ressources:

Bigdata - Nathan Marz book

NoSQL Distilled - Martin Fowler

7 databases

BigTable paper

Postgres cheatsheat

MovieLens dataset



TP1: Installation en détail

Fedora Ubuntu

```
yum install postgresql postgresql-server postgresql-contrib (1)
postgresql-setup initdb (2)

systemctl start postgresql.service (3)

yum install pgadmin3 (4)

CREATE EXTENSION tablefunc; (5)
CREATE EXTENSION dict_xsyn;
CREATE EXTENSION fuzzystrmatch;
CREATE EXTENSION pg_trgm;
CREATE EXTENSION cube;
```

1 Instalation du client/serveur/extensions suplementaires

2 Initialisation de la base

3 Demarage du serveur

4 Front-end requetage

5 Installation des extensions **Verifier les extensions installees**

Create index

```
CREATE INDEX [ nom ] ON table [ USING method ]
              ( { colonne | ( expression ) } [ classeop ] ... )
```

- *method*: btree/hash/gin/gist
- *classeop* : operator class that can use the index

[Documentation](#)